



Inspire...Educate...Transform.
**Statistical Techniques in
Data Science**
Time Series Forecasting

Dr. Sridhar Pappu
Executive VP – Academics, INSOFE

January 15, 2017

Why Time Series

- Causal independent variables are
 - Unknown to us
 - Not available
 - Might not fit the data well
 - Difficult to forecast

NASDAQ: CTSH - 26-Feb 4:00 PM GMT-5

57.04 USD ↓0.20 (0.35%)



2024



Typical Time Series

$$\hat{y}_{t+1} = f(y_t, y_{t-1}, y_{t-2} \dots) + f(x_1, x_2, x_3 \dots)$$

f can be linear or nonlinear



IMPORTANT CONCEPTS



Time Series Descriptive Statistics

- In descriptive statistics covered earlier (central tendencies, measures of variability, skewness, kurtosis, distributions, correlations, etc.), the order of observations in the data was of no consequence.
- In time series descriptive statistics, order of observations is of primary importance and so autocorrelations, etc. play a vital role in identifying the models and their characteristics.

Autocorrelation (ACF) and Partial ACF (PACF)

- ACF: n^{th} lag of ACF is the correlation between a day and n days before that.
- PACF: The same as ACF with all intermediate correlations removed. It is the k_{th} coefficient of the ordinary least squares regression.

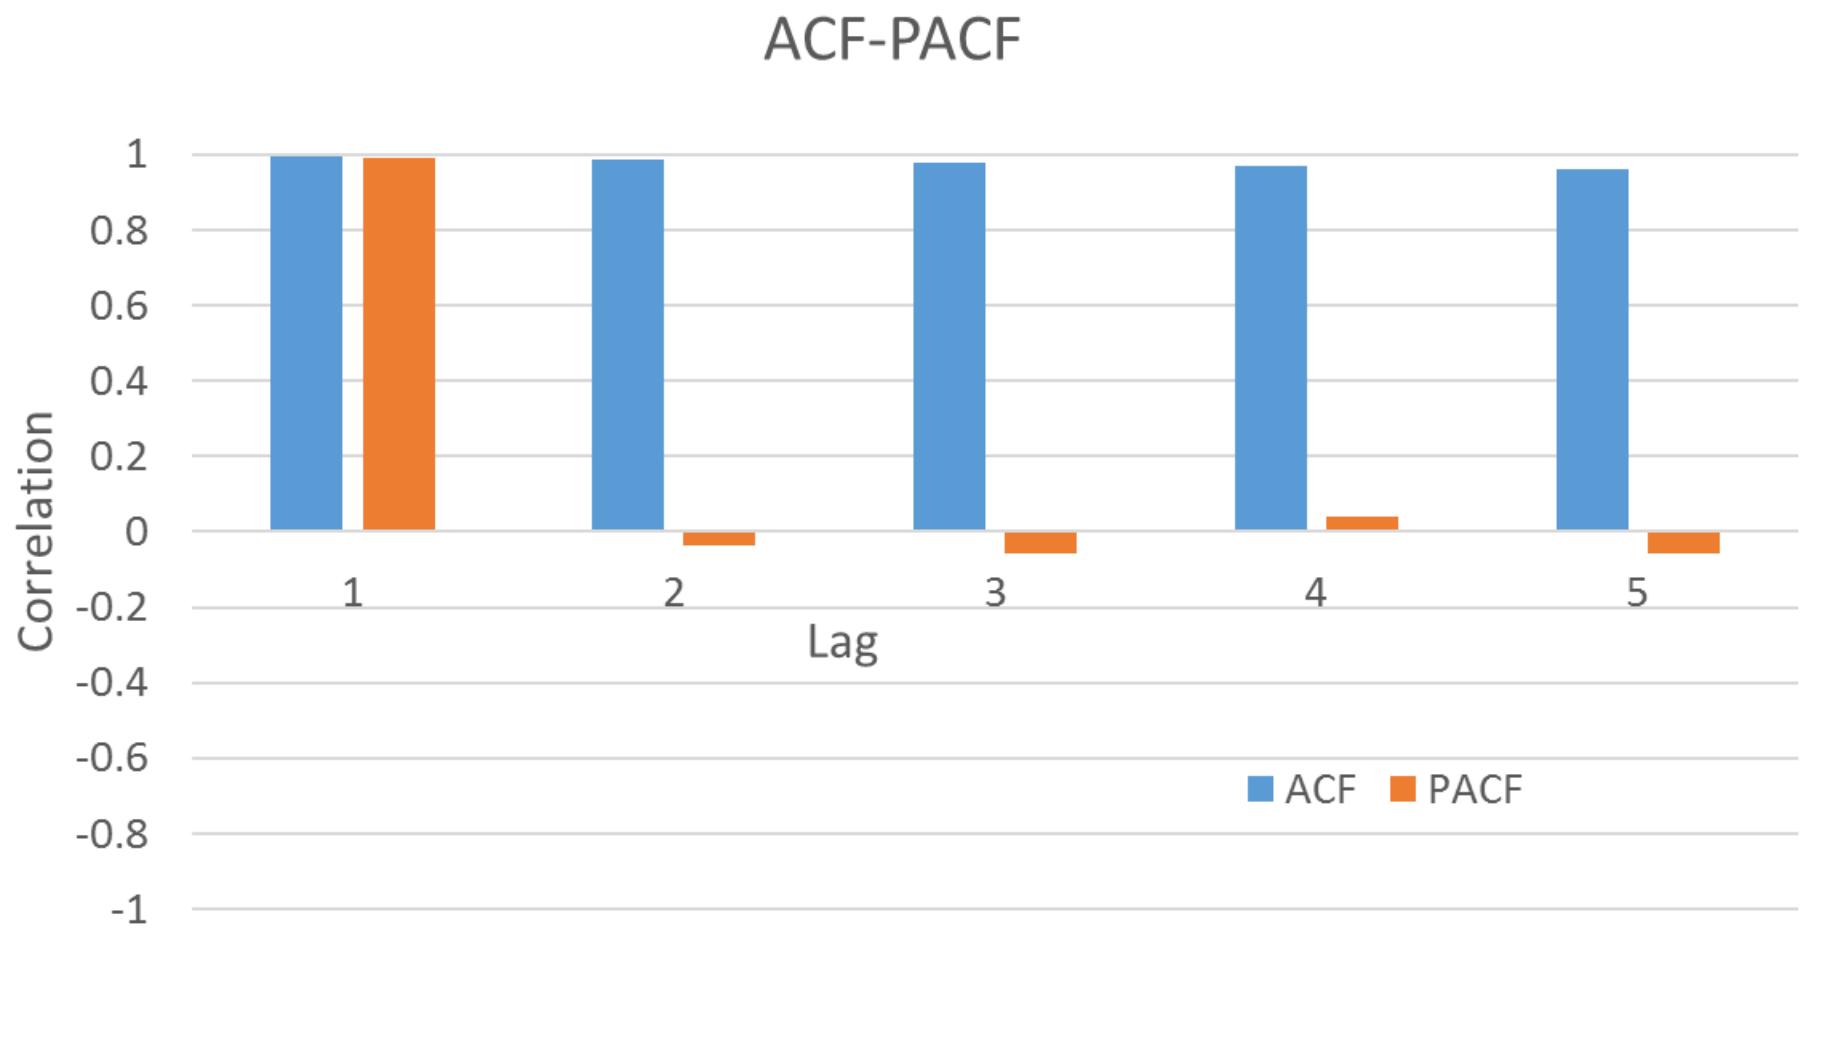
$$[y_t] = \beta_0 + \sum_{i=1}^k \beta_i [y_{t-i}] \text{ where}$$

$[y_t]$ is the input time series, k is the lag order and β_i is the i_{th} coefficient of the linear multiple regression.

EXCEL ACTIVITY

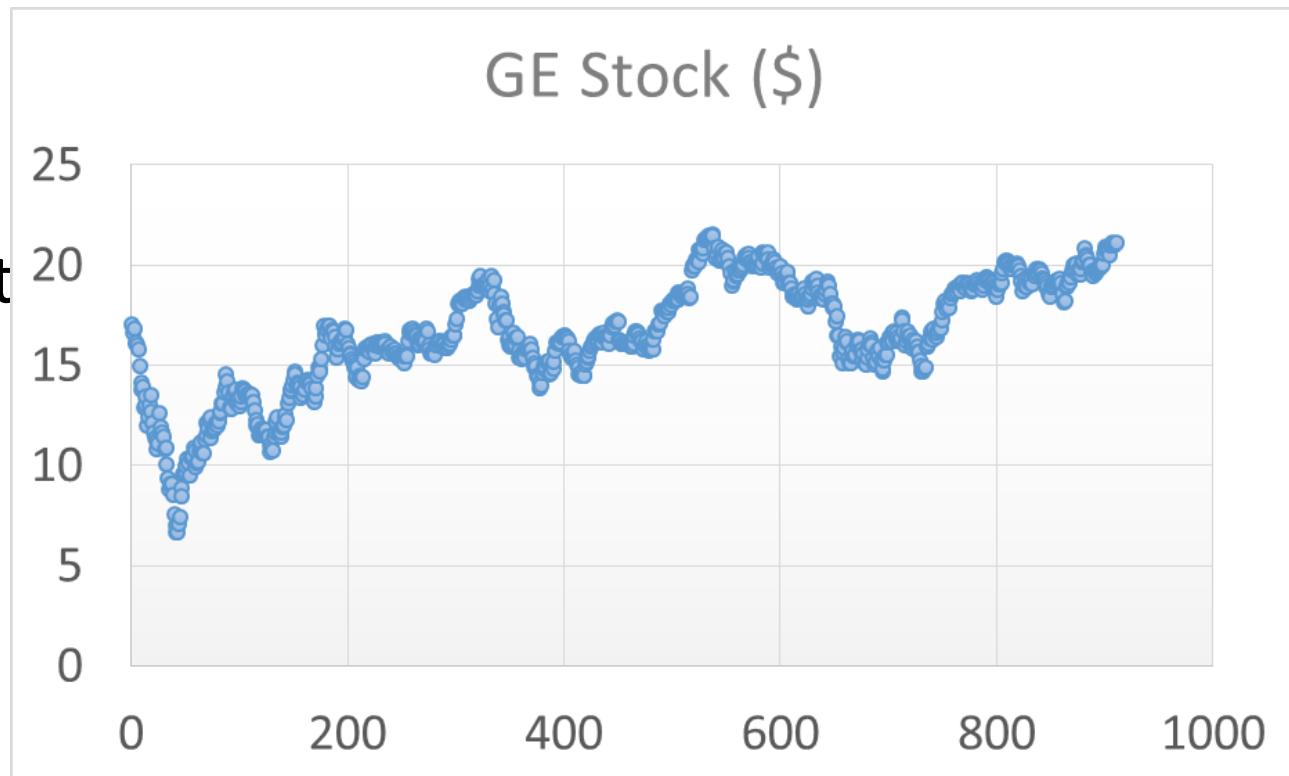


Autocorrelation (ACF) and Partial ACF (PACF)

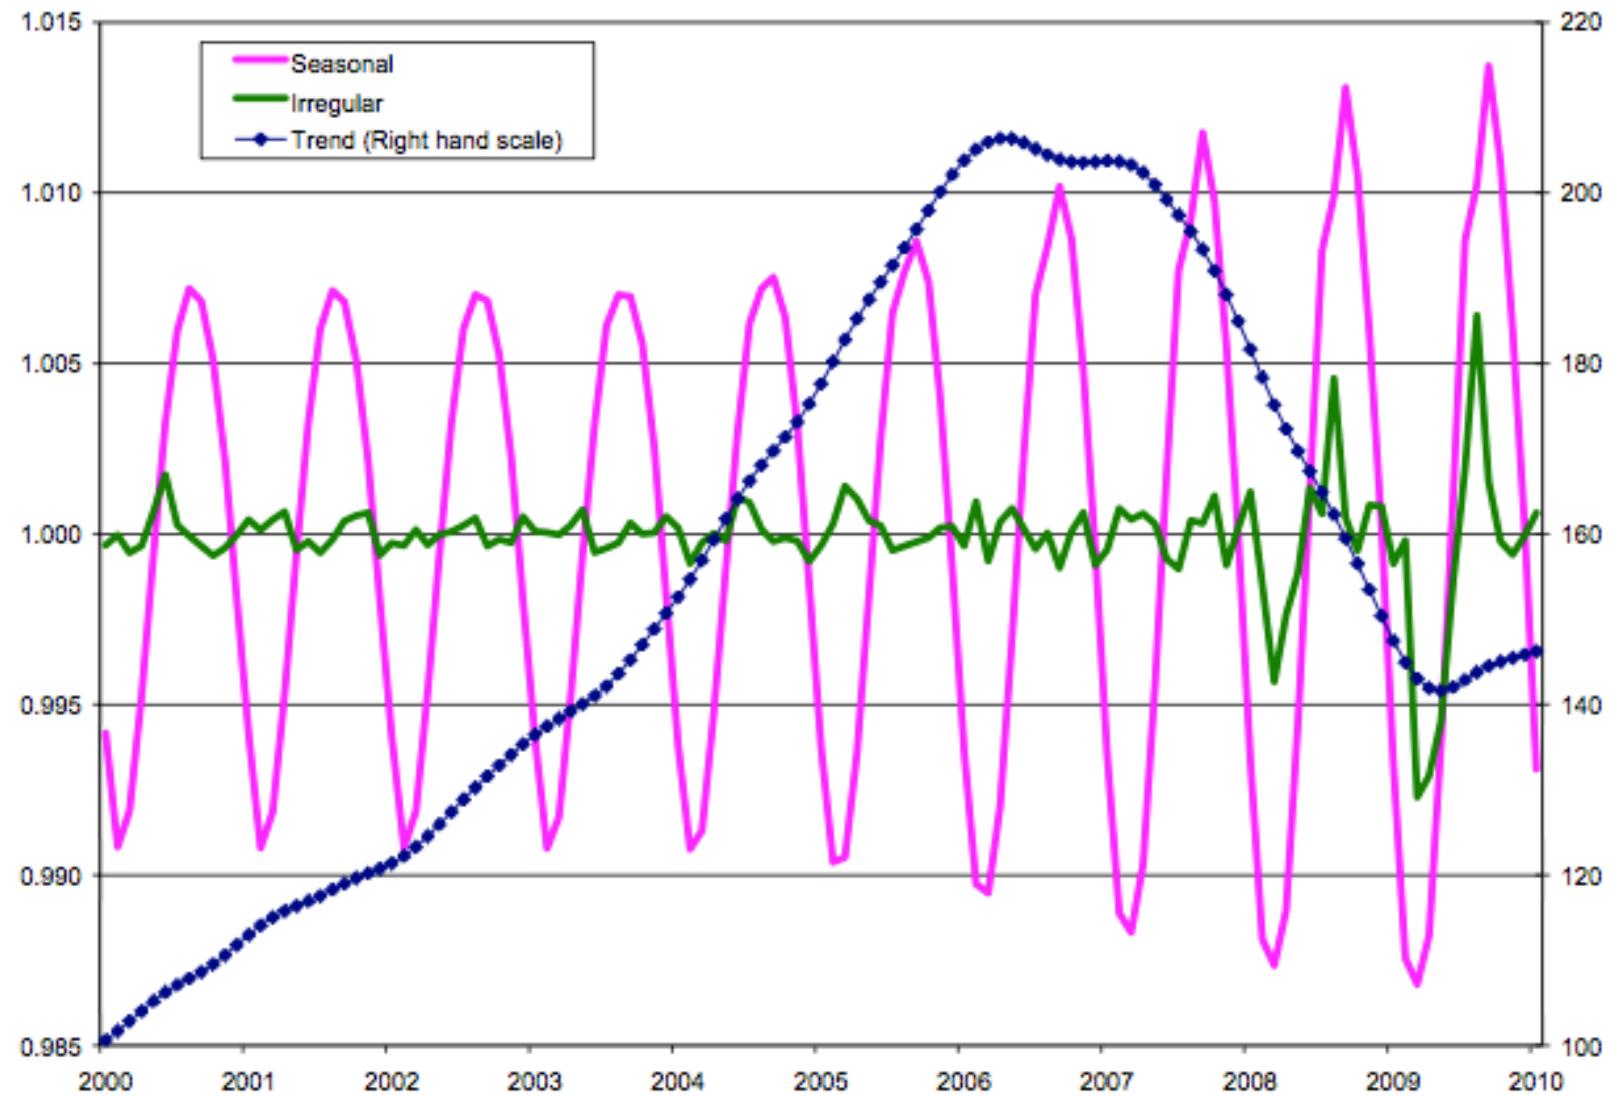


Components of Time Series

- Trend
- Seasonality
- Random component



Trend, Seasonality and Randomness



US Air Carrier Traffic – Revenue Passenger Miles ('000)

RPM

```
> milestimeseries <- ts(miles, frequency = 12, start = c(1996,1))
> milestimeseries
```

	Jan	Feb	Mar	Apr	May	Jun	Jul
1996	41972194	42054796	50443045	47112397	49118248	52880510	
1997	45850623	42838949	53620994	49282817	51191842	54707221	
1998	46514139	43769273	53361926	51968480	53515798	56460422	
1999	47988560	45241211	56555731	53920855	54674958	59213000	
2000	49045412	49306303	60443541	58286680	60533783	64903295	
2001	52634354	49532578	61575055	59151645	59662416	64353323	
2002	46224031	44615129	56897729	52542164	55116060	59745343	
2003	51197175	47040806	56766580	51857453	54335598	60272900	
2004	53979786	53179693	64035864	62340117	62530704	68866398	
2005	59629608	55795165	70595861	65145552	68268899	72952959	
2006	61035027	56729212	70799794	68120559	69352606	74099239	
2007	63016013	57793832	72700241	69836156	71933109	76926452	
2008	64667106	61504426	74575531	68906882	72725750	76162105	
2009	58373786	53506580	66027341	65166300	65868254	71350227	
2010	59651061	53240066	68307090	64953250	68850904	74474550	
2011	61630362	55391206	70158268	67683558	71711448	76057910	
2012	61940180	58243763	71696039	68669228	71887523	76760759	
	Aug	Sep	Oct	Nov	Dec		
1996	57723208	47035464	49263120	43937074	48539606		
1997	59715433	49418190	51058879	47056048	49654209		

Aug

Sep

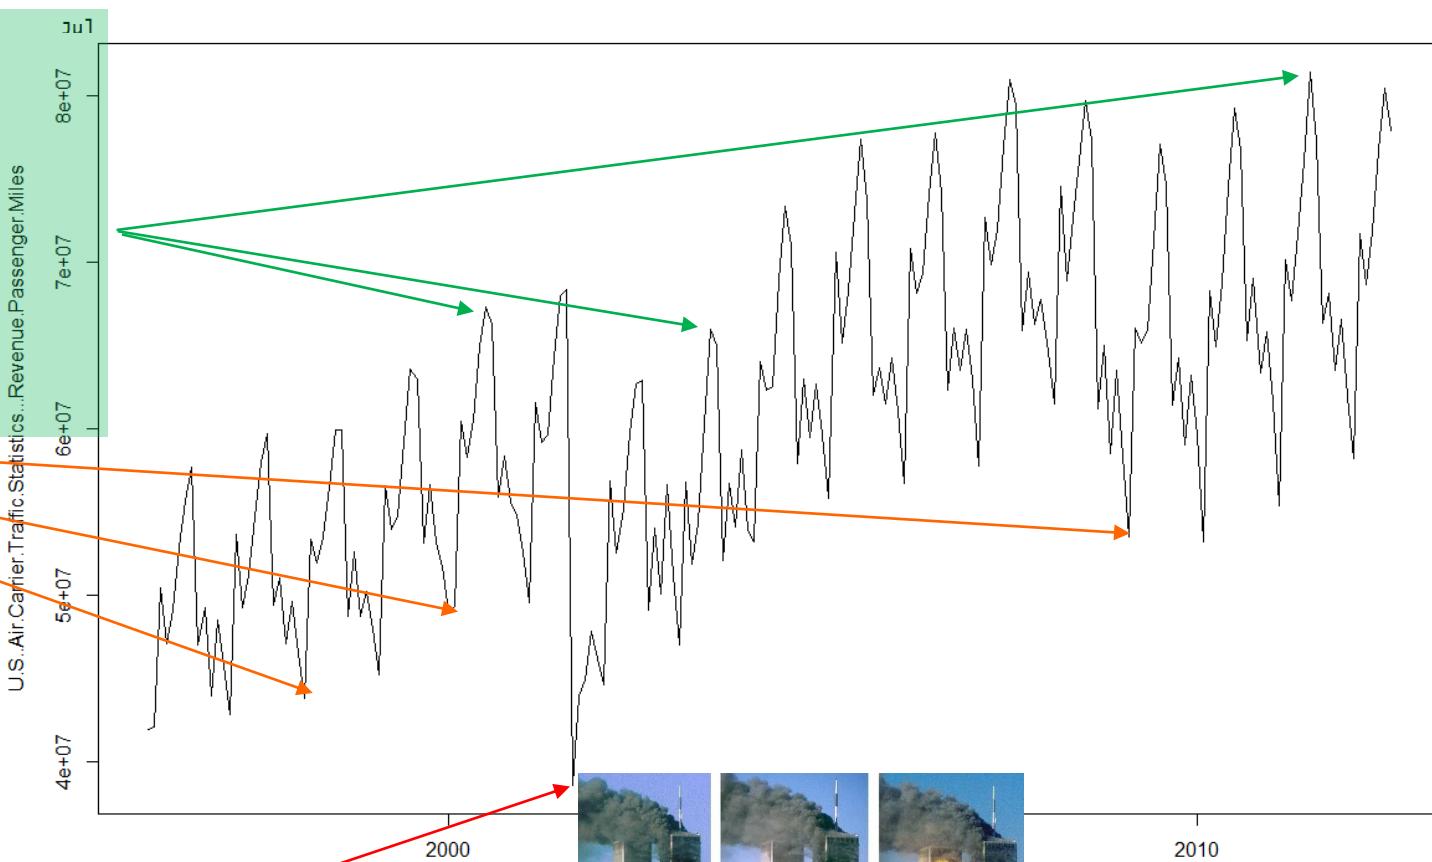
Oct

Nov

Dec

1996

1997



Data sources:

http://www.bts.gov/xml/air_traffic/src/index.xml

and <https://datamarket.com/data/set/281x/us-air-carrier-traffic-statistics-revenue-passenger-miles>

Last accessed: 31-Mar-2016

	Aug	Sep	Oct	Nov	Dec
1996	57723208	47035464	49263120	43937074	48539606
1997	59715433	49418190	51058879	47056048	49654209
1998	59927214	48751280	52578217	48734375	50208641
1999	63003663	53131972	56653901	53215500	51746821
2000	66256804	55900504	58373996	55590325	54822970
2001	68377080	38601868	43964788	44915764	47836501
2002	62944816	49096035	54019748	50106814	56656594
2003	64989766	52121480	56724551	54128776	58739845
2004	70961522	57881042	63021142	59453943	62680310



72020



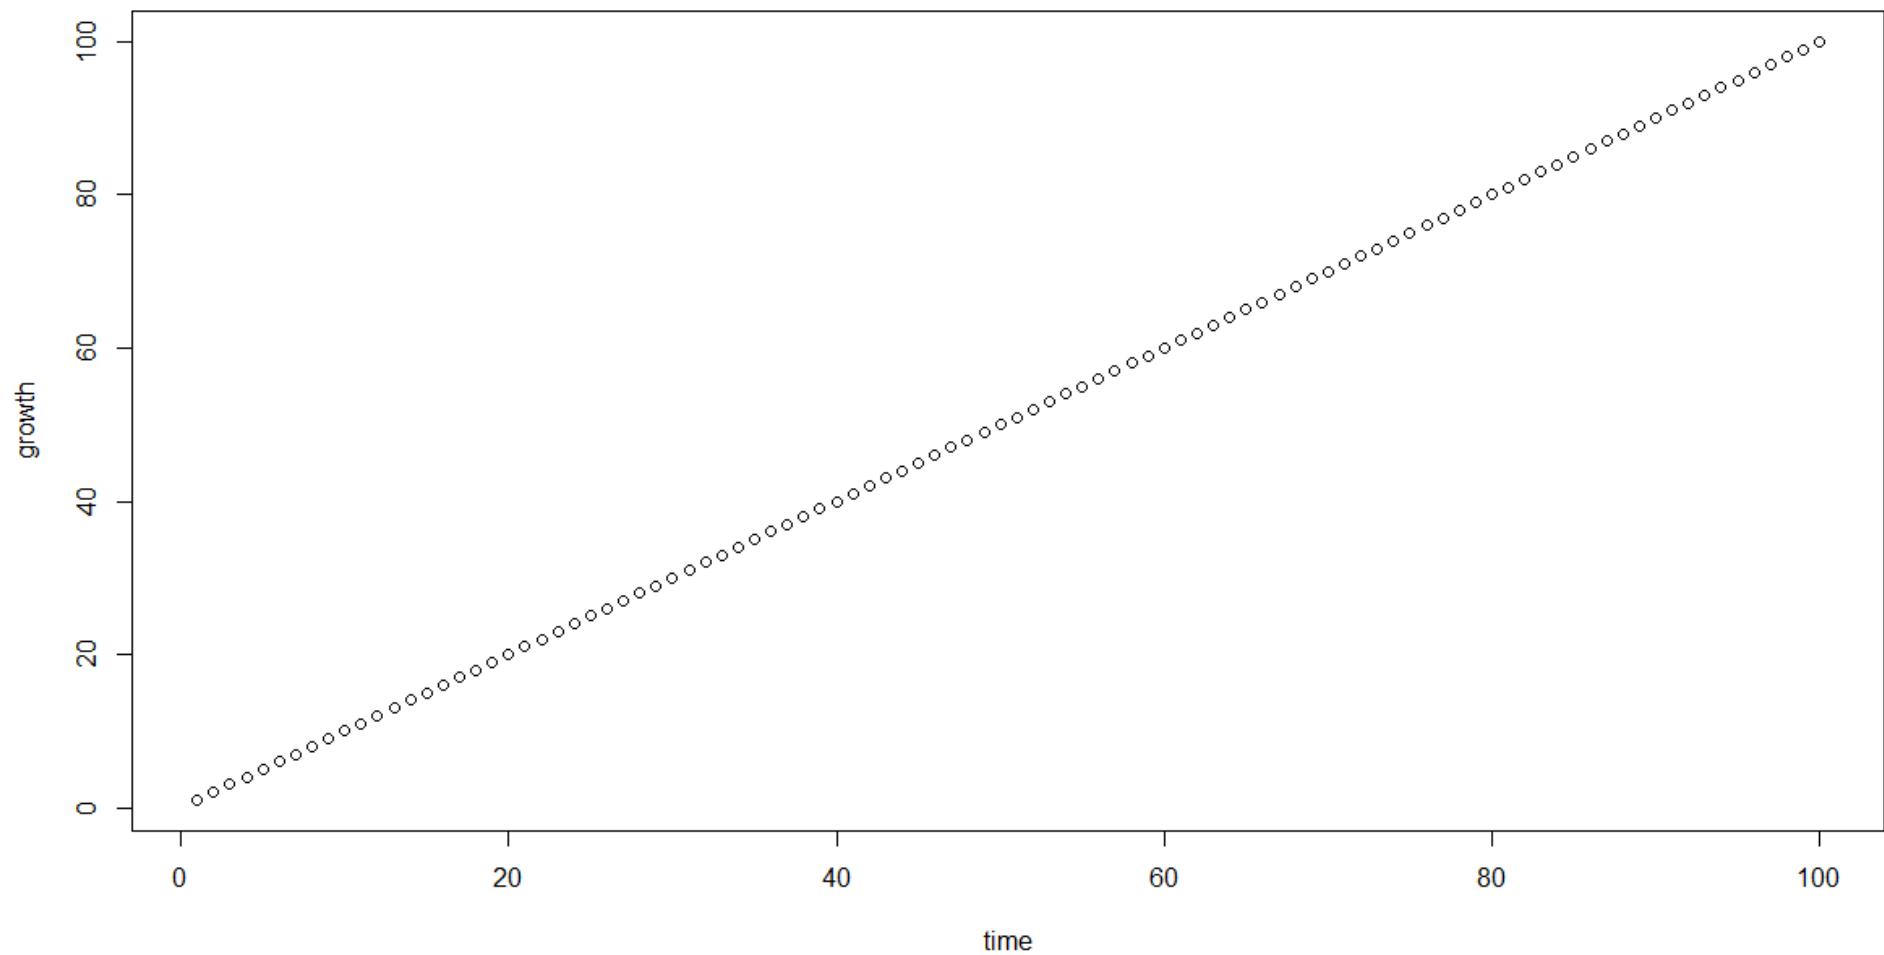
ACF and PACF – Idealized Trend, Seasonality and Randomness



CSE 7202c



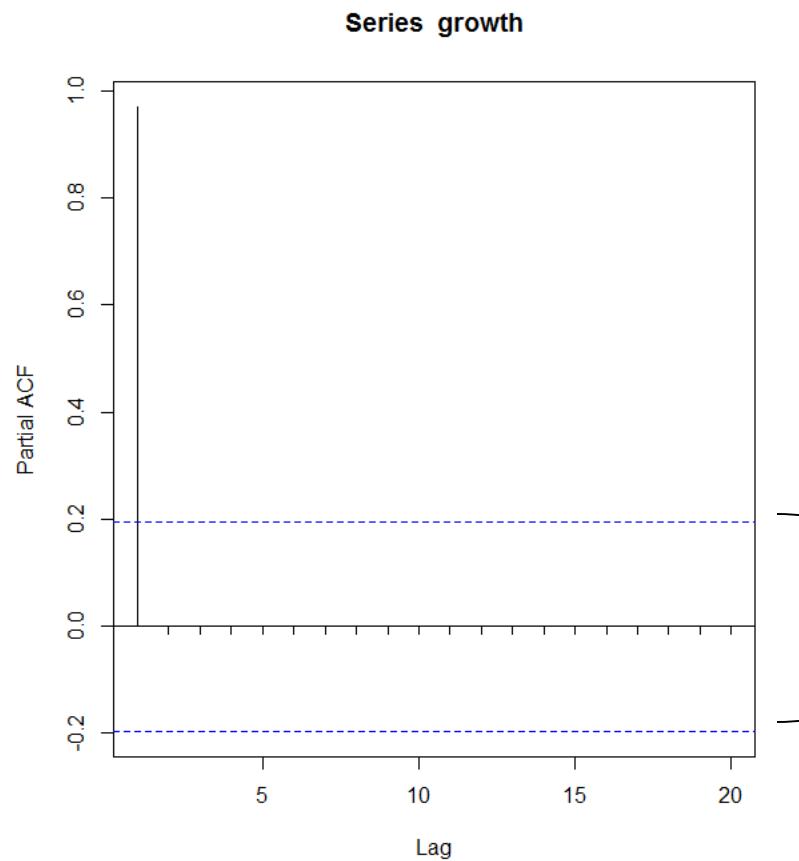
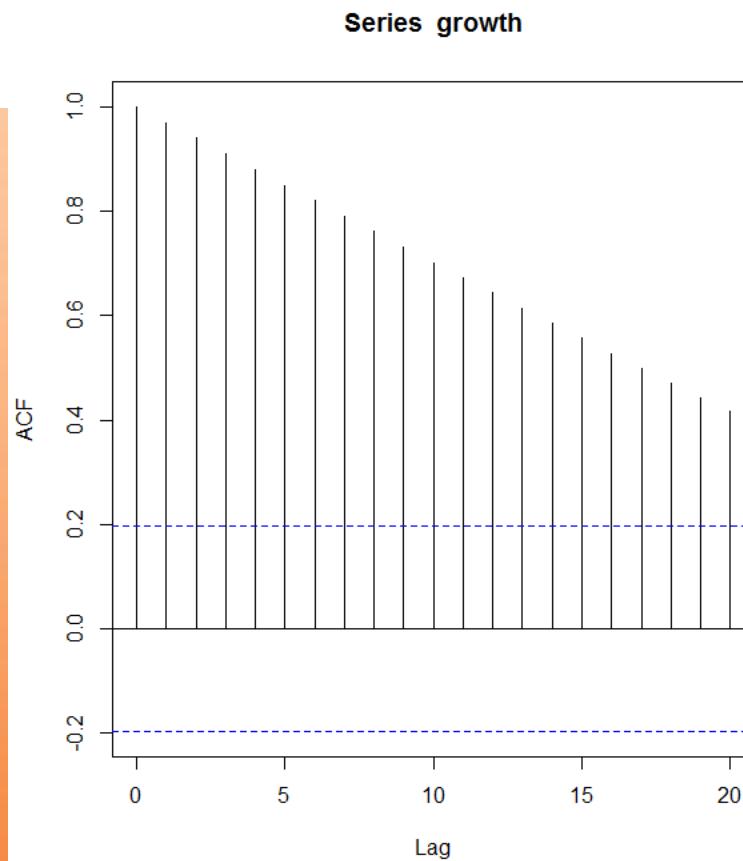
ACF and PACF – Idealized Trend



CSE 7202c



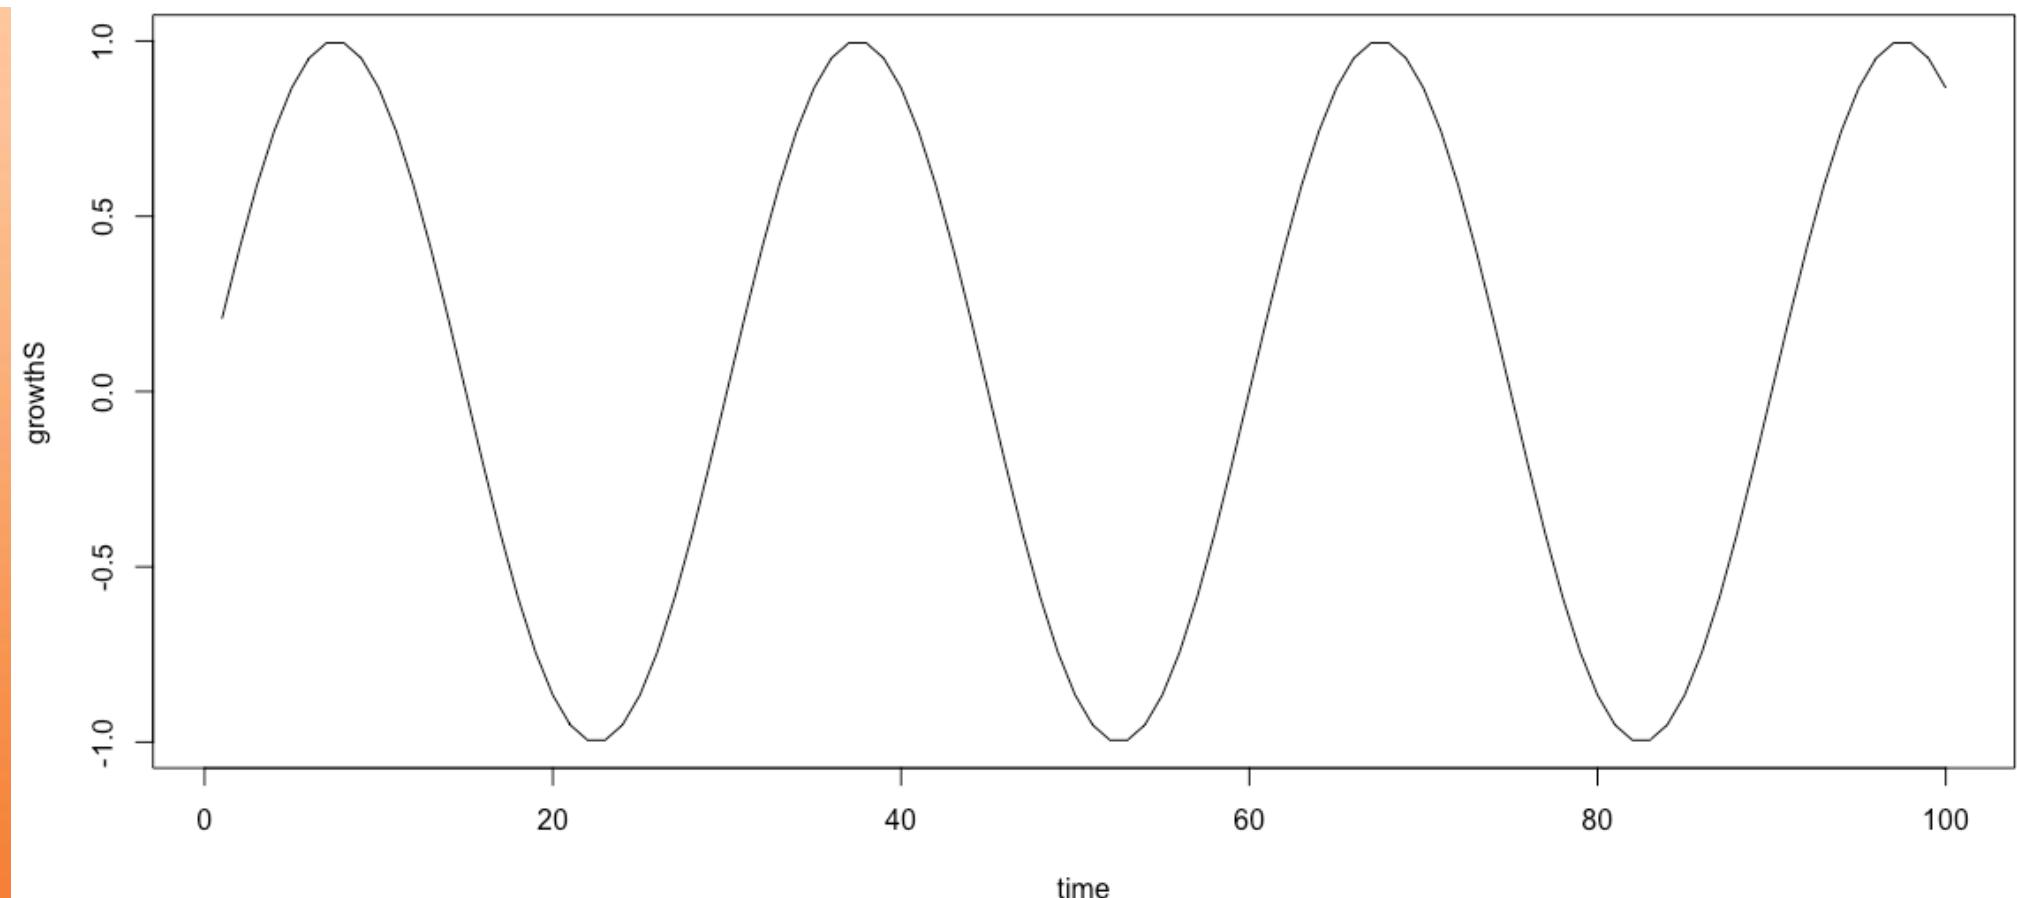
ACF and PACF – Idealized Trend



$$95\% \text{ CI: } 0 \pm \frac{1.96}{\sqrt{n}}$$

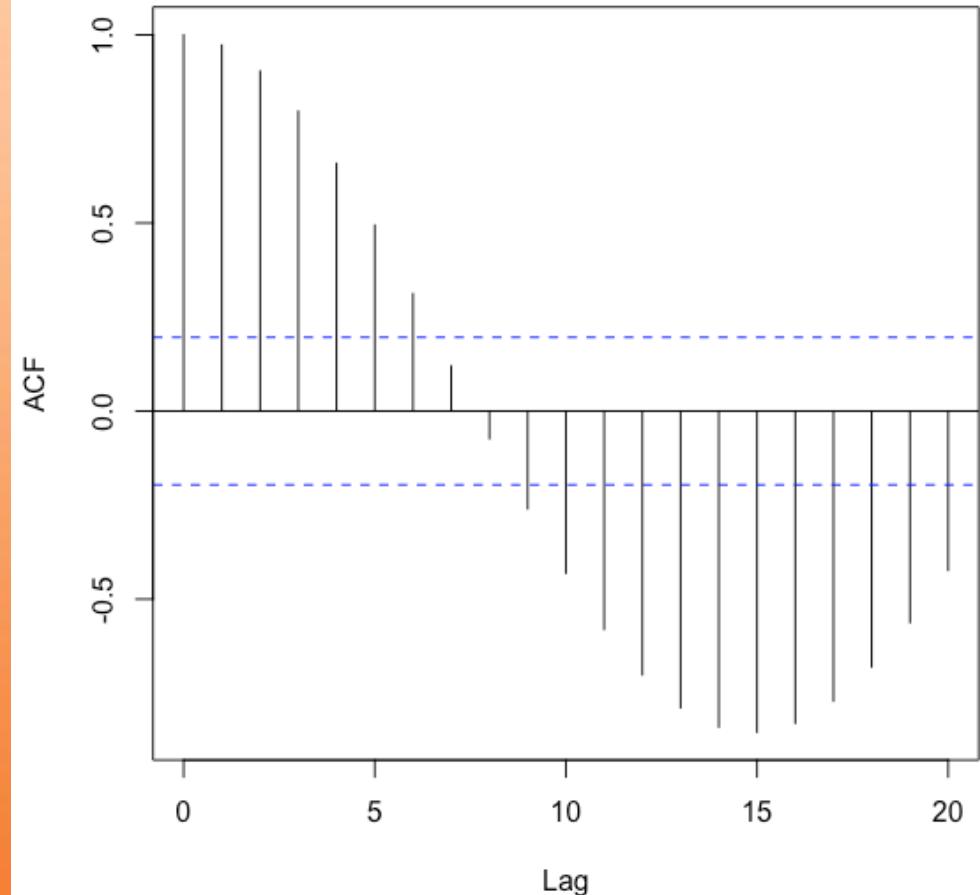
- ACF is a bar chart of correlation coefficients of the time series and its lags.
- PACF is a plot of the partial correlation coefficients of the time series and its lags.

ACF and PACF – Idealized Seasonality

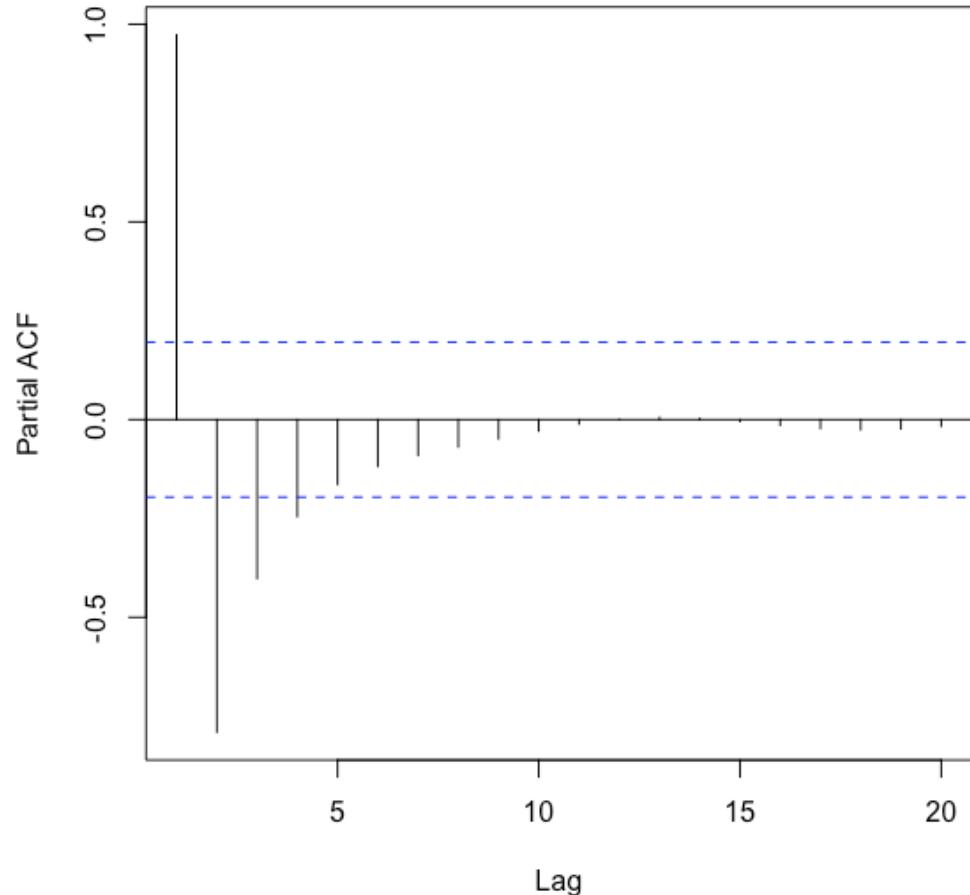


ACF and PACF – Idealized Seasonality

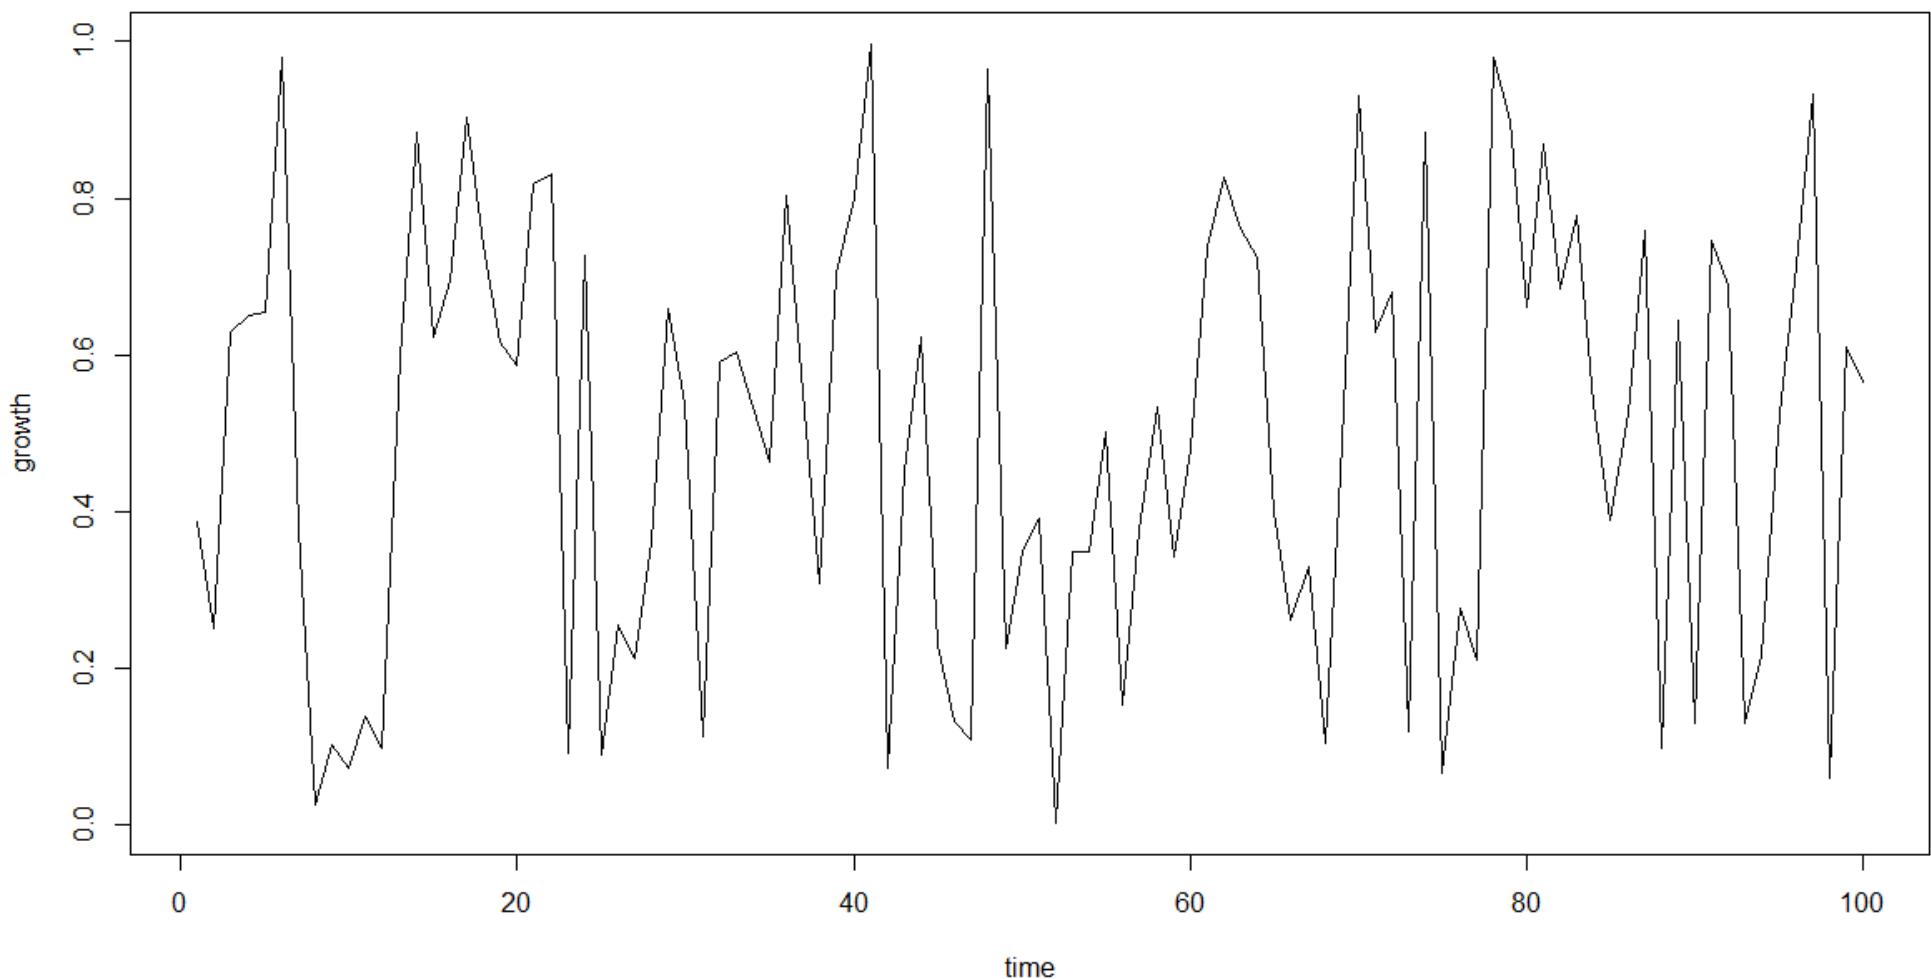
Series growth



Series growth



ACF and PACF – Idealized Randomness

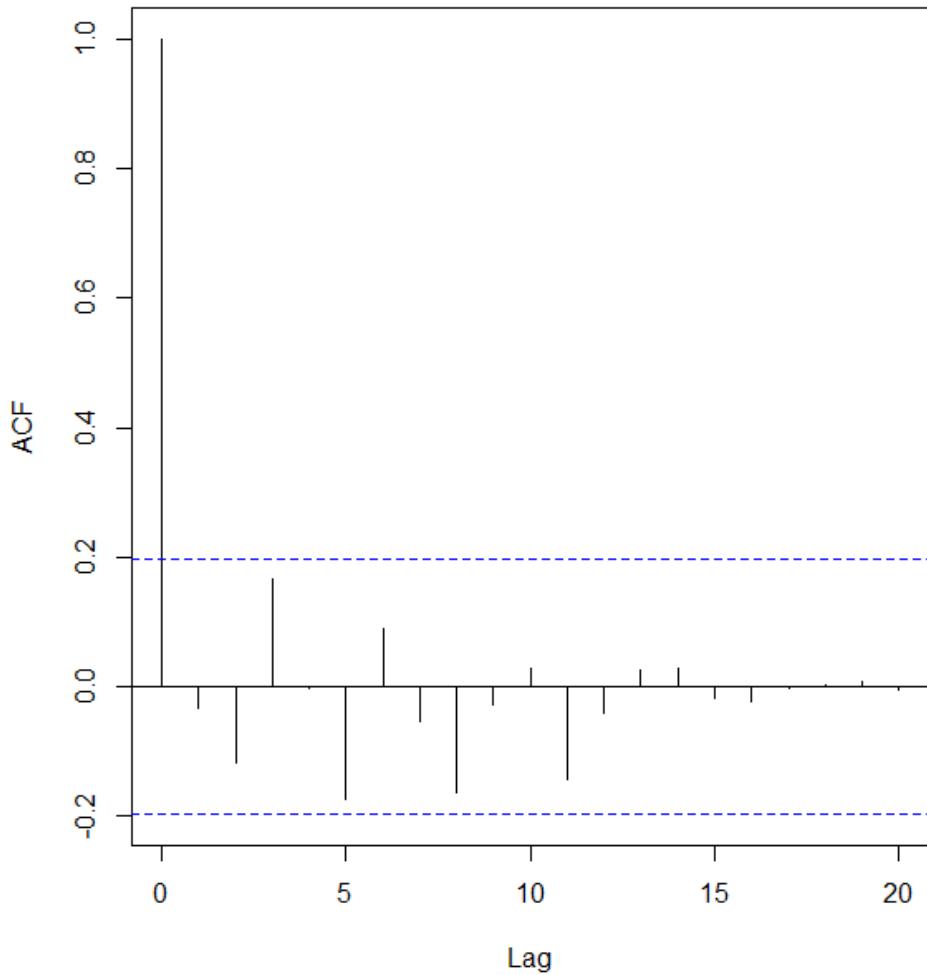


CSE 7202c

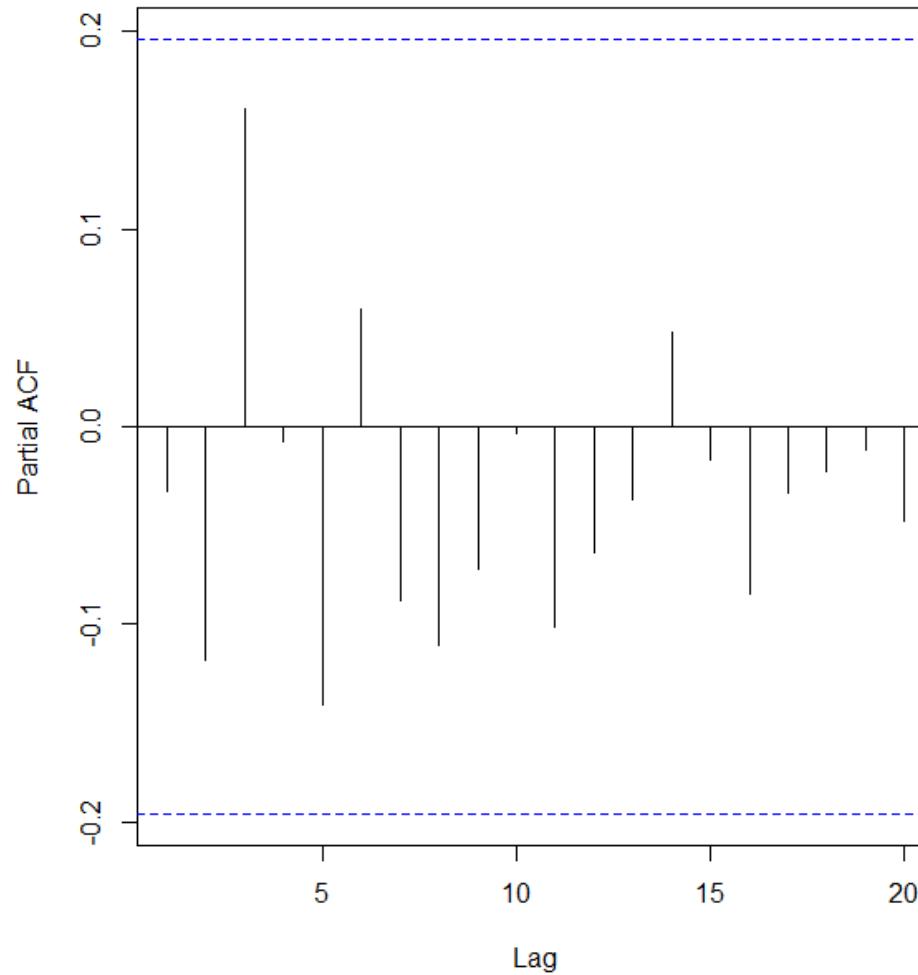


ACF and PACF – Idealized Randomness

Series growth



Series growth



ACF and PACF – Idealized Trend, Seasonality and Randomness

- Ideal Trend: Decreasing ACF and 1 or 2 lags of PACF
- Ideal Seasonality: Cyclical in ACF and a few lags of PACF with some positive and some negative
- Ideal Random: A spike may or may not be present; even if present, magnitude will be small

CSE 7202c



ACF and PACF (Real-world): Decomposing Time Series into the 3 Components

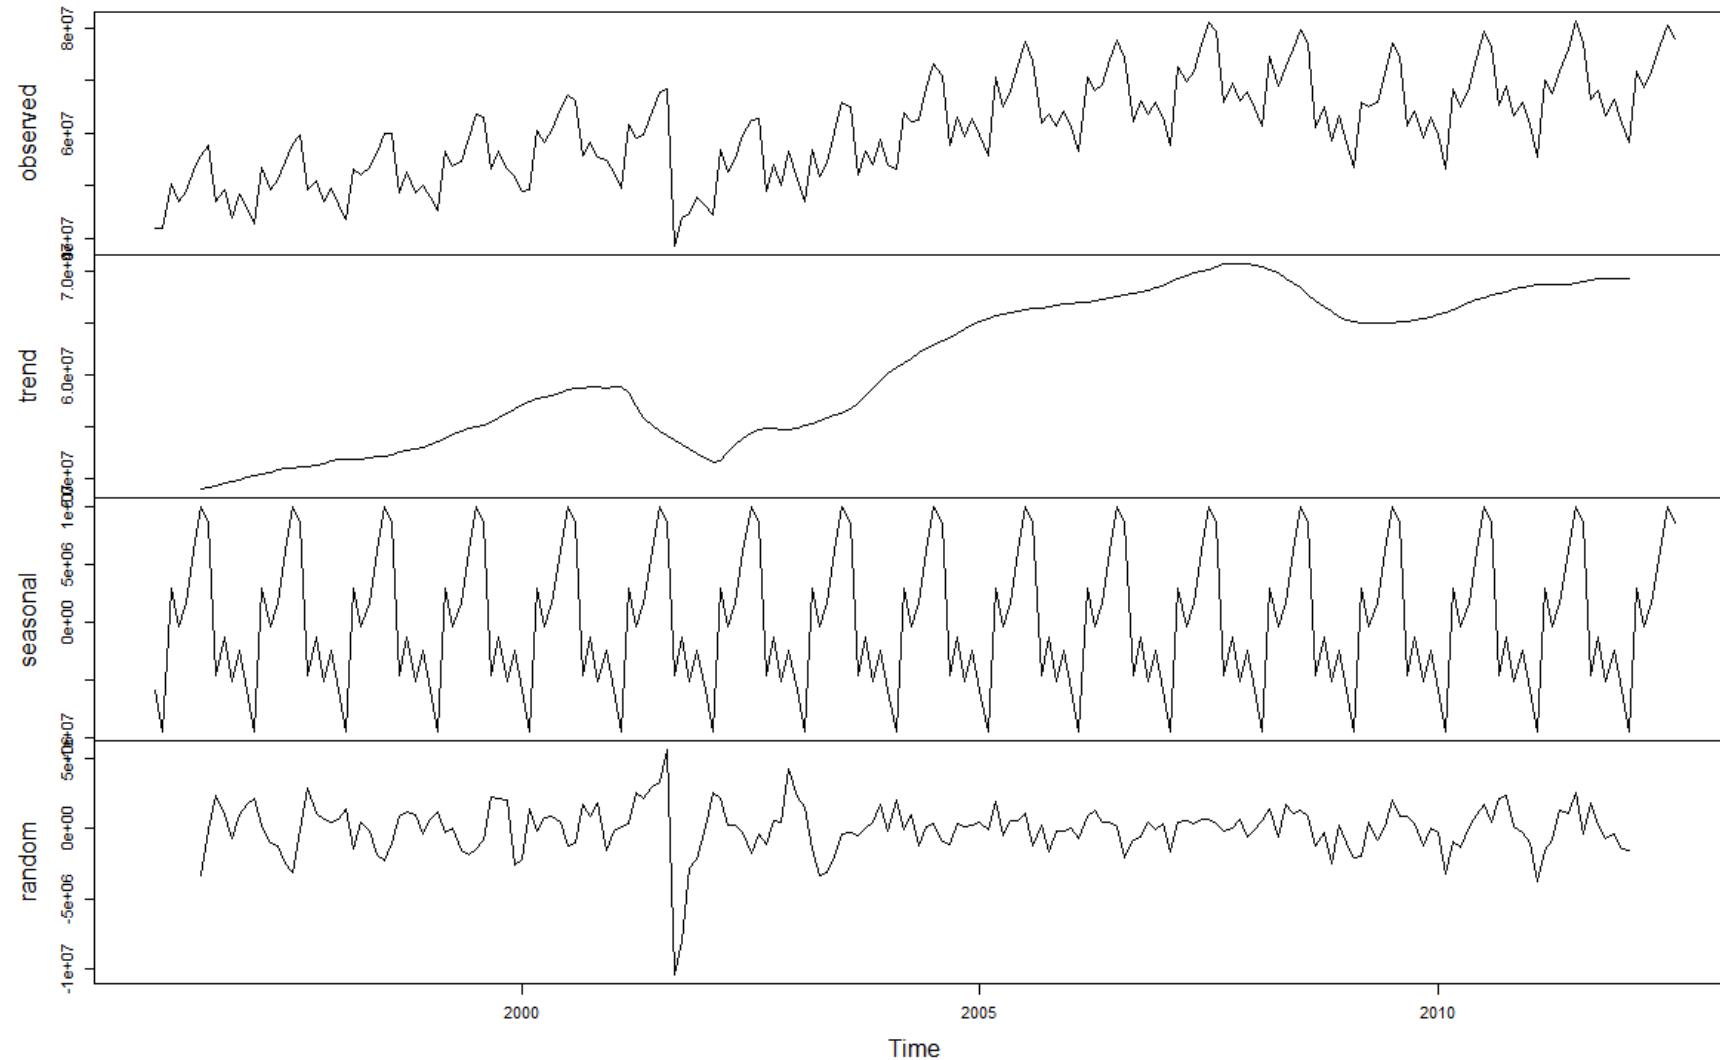


CSE 7202c



ACF and PACF (Real-world): Decomposing Time Series into the 3 Components – Revenue Passenger Miles (RPM)

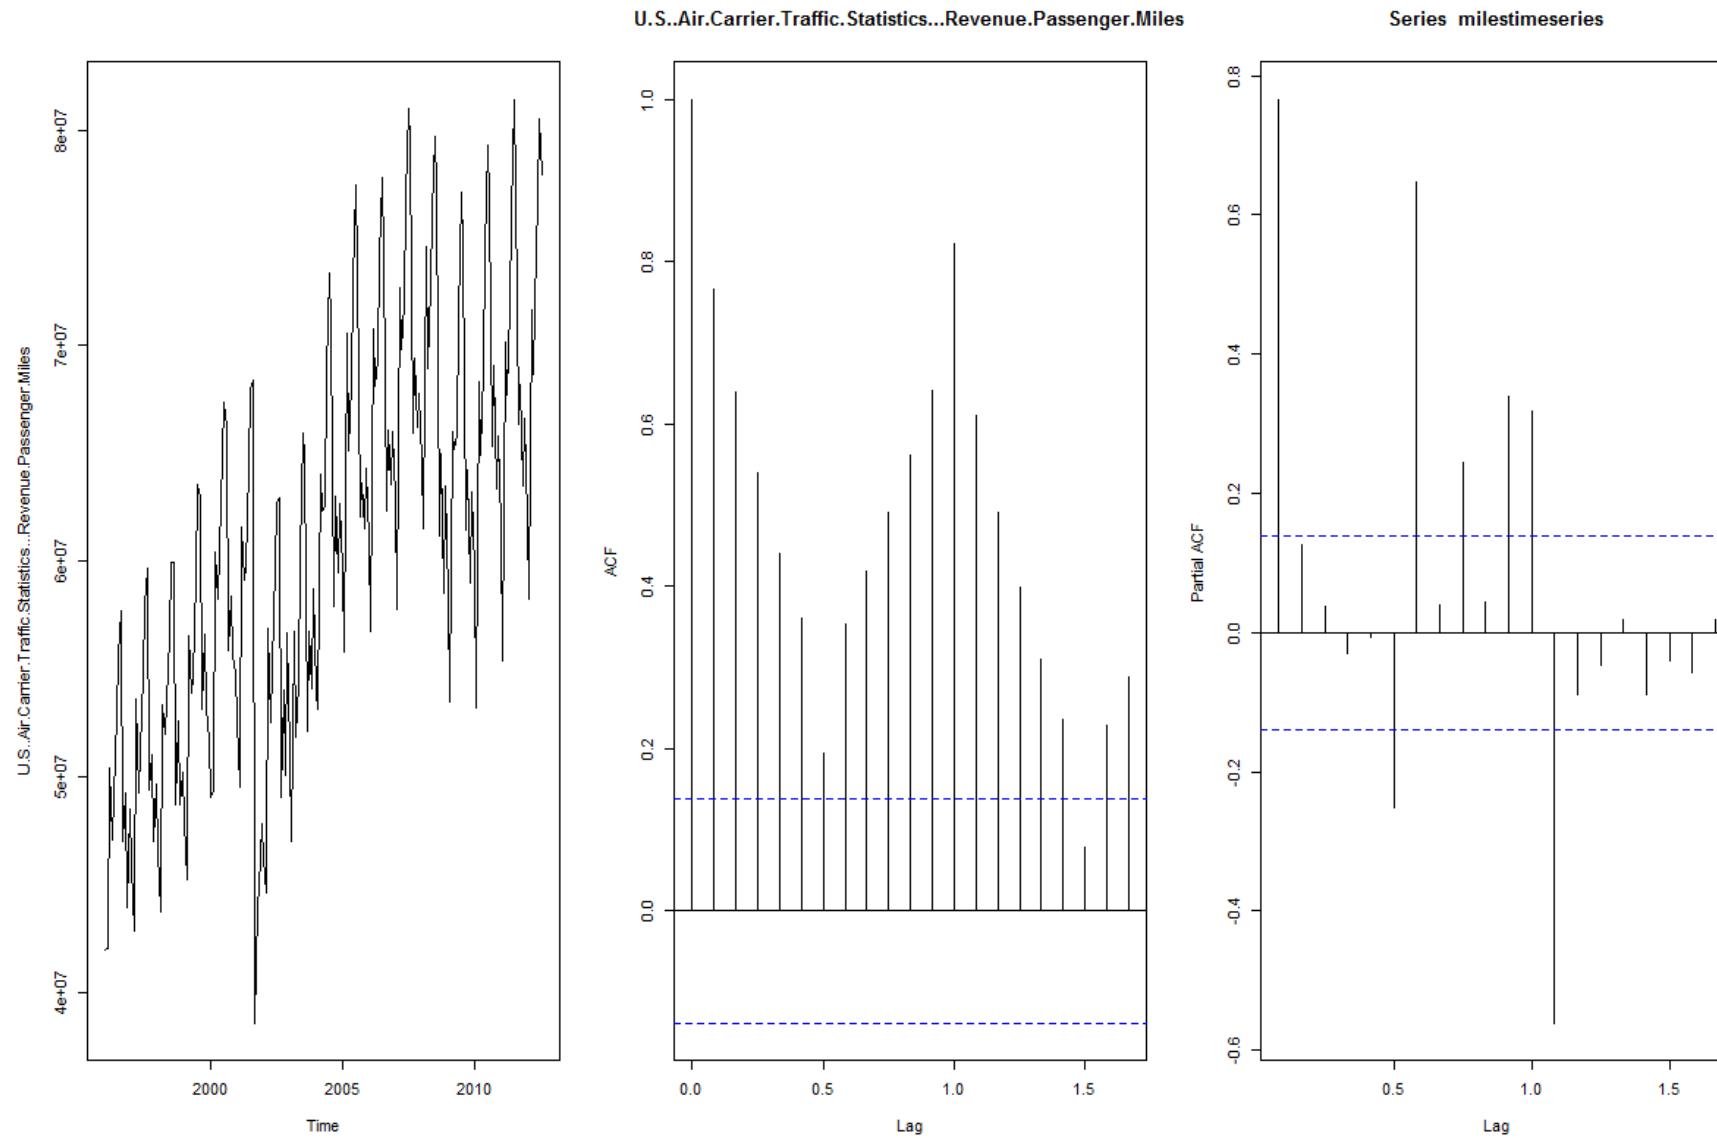
Decomposition of additive time series



CSE 7202c



ACF and PACF (Real-world): Decomposing Time Series into the 3 Components – RPM



CSE 72026



Stationary and Non-Stationary

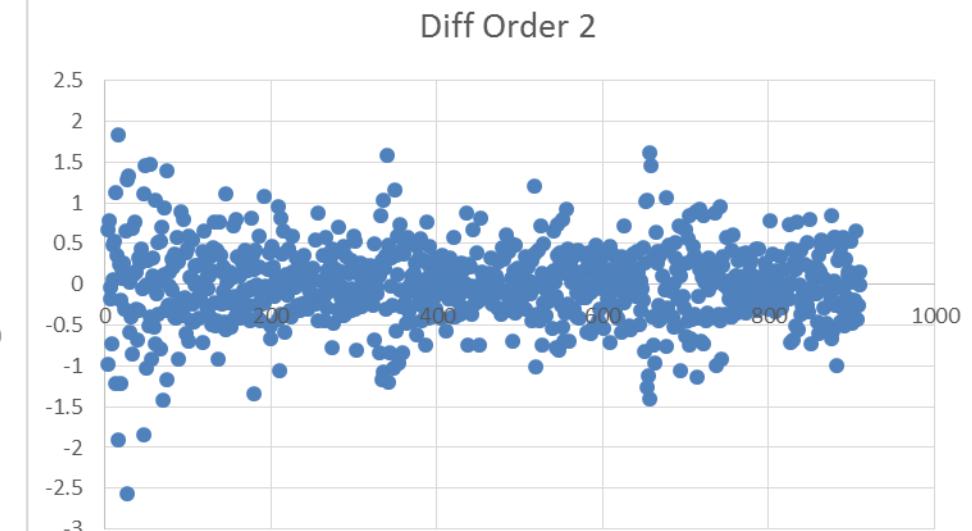
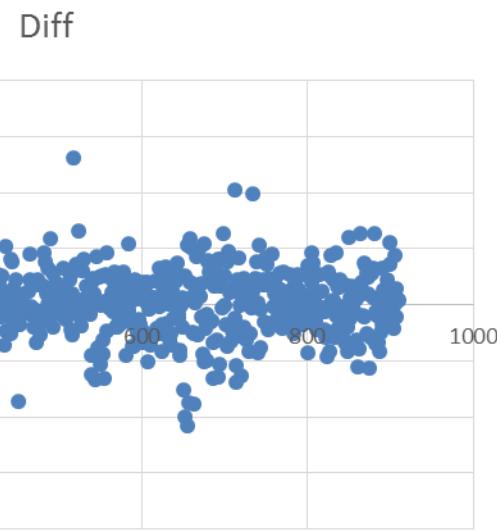
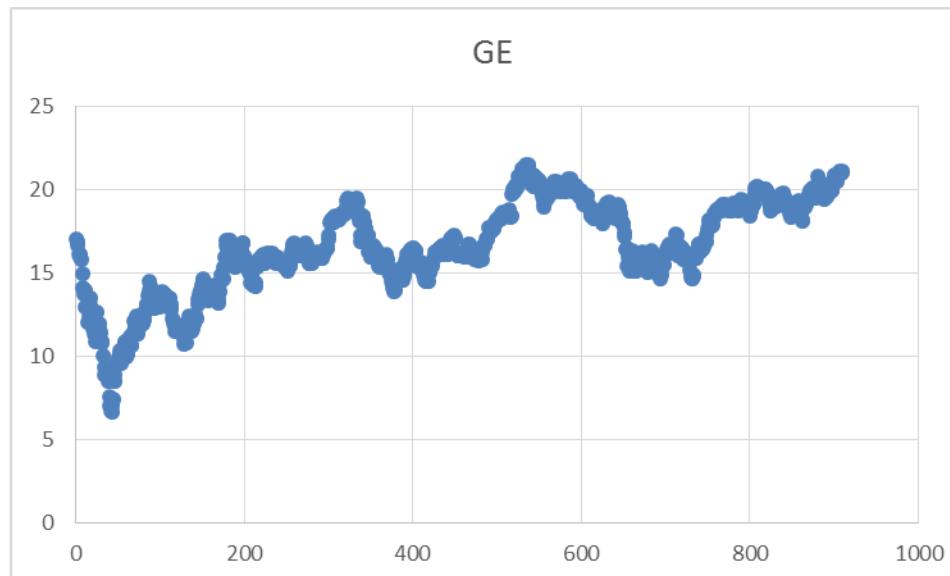
- Stationary data has constant statistical properties – mean, variance, autocorrelation, etc. – over time
- If the data is stationary, forecasting is easier!
- Differencing to convert non-stationary to stationary

EXCEL ACTIVITY

CSE 7202C



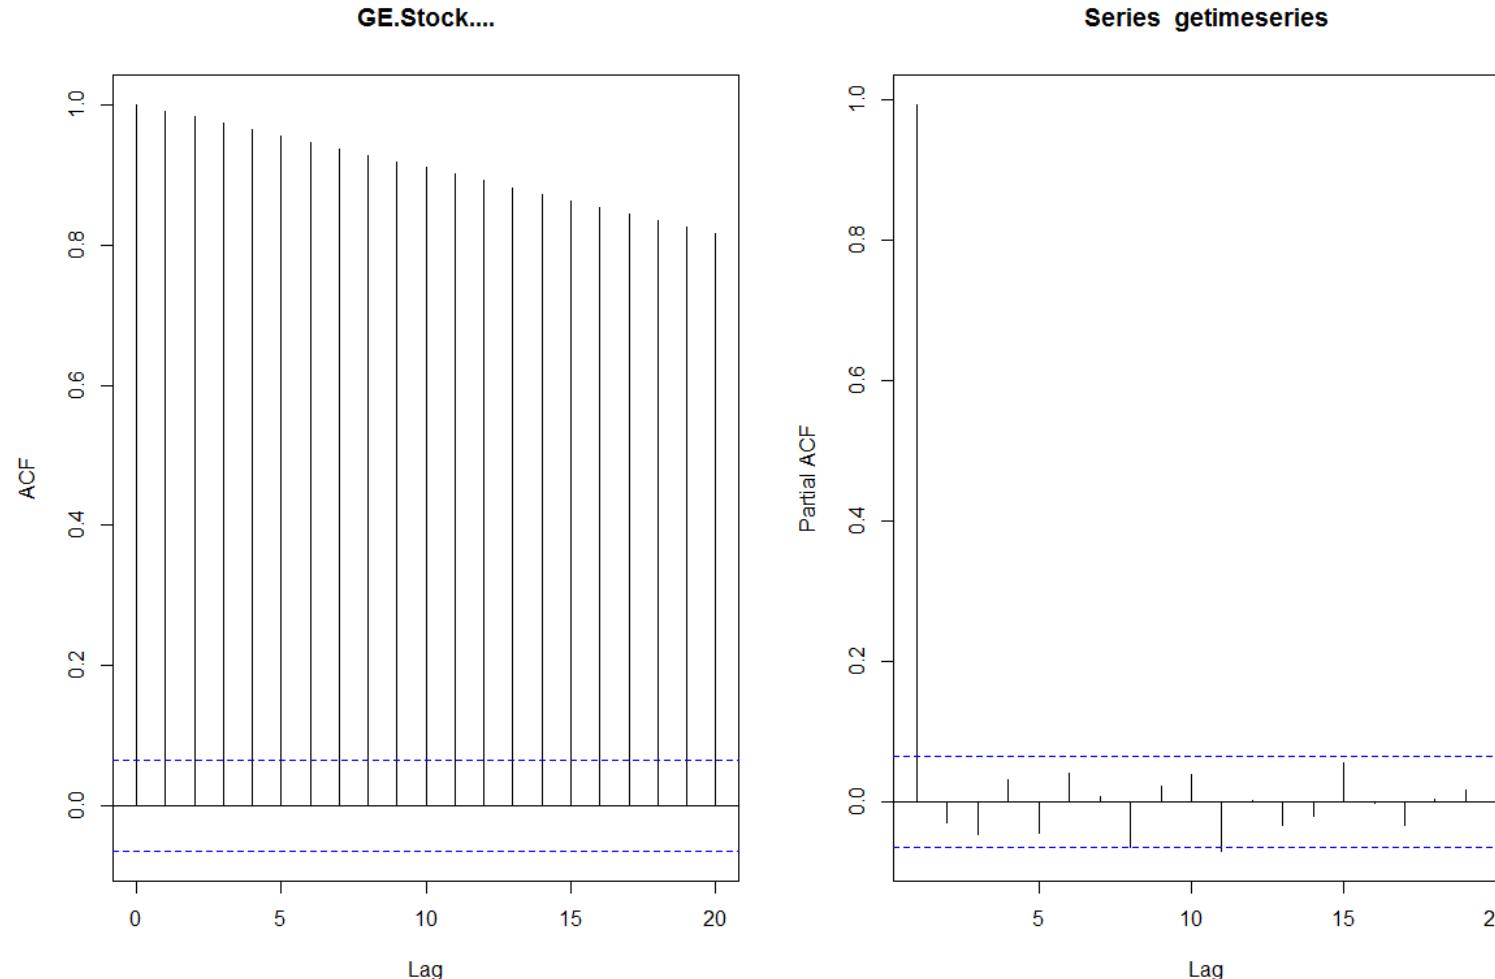
Removing Trend from Data



CSE 7202c



ACF and PACF of Stationary and Non-Stationary

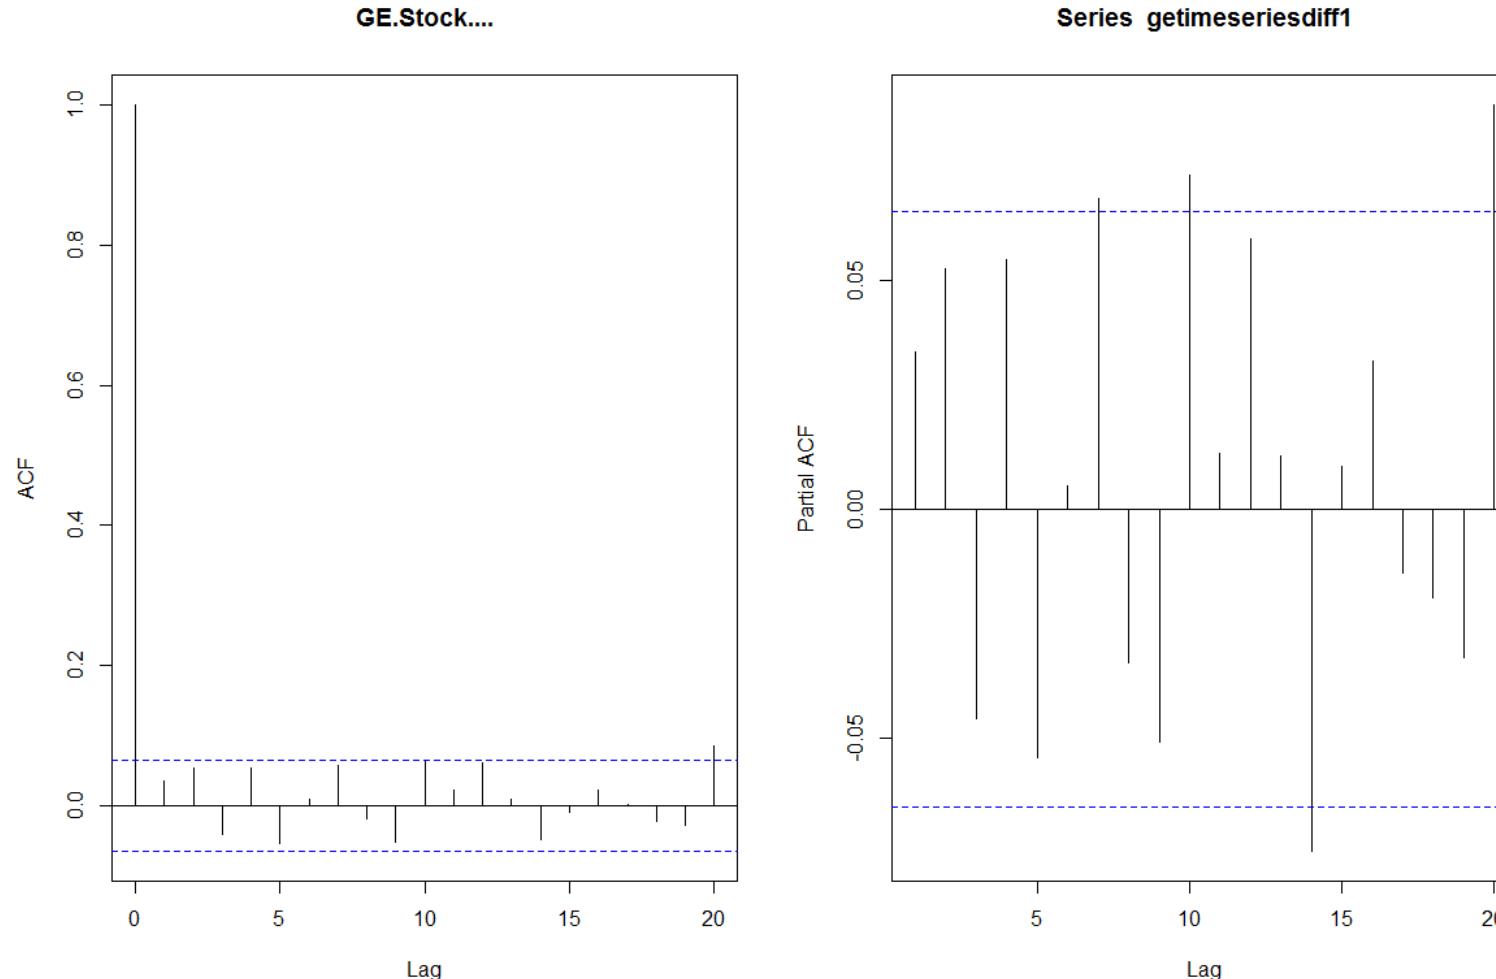


Price of GE stock is highly correlated with the previous day's value.

CESE 72026



ACF and PACF of Stationary and Non-Stationary



Daily changes in GE stock price are essentially random.

ACF and PACF of Stationary and Non-Stationary

- Non-stationary series have an ACF that remains significant for half a dozen or more lags, rather than quickly declining to zero.
- You must difference such a series until it is stationary before you can identify the process.

CSE 7202c



Time Series

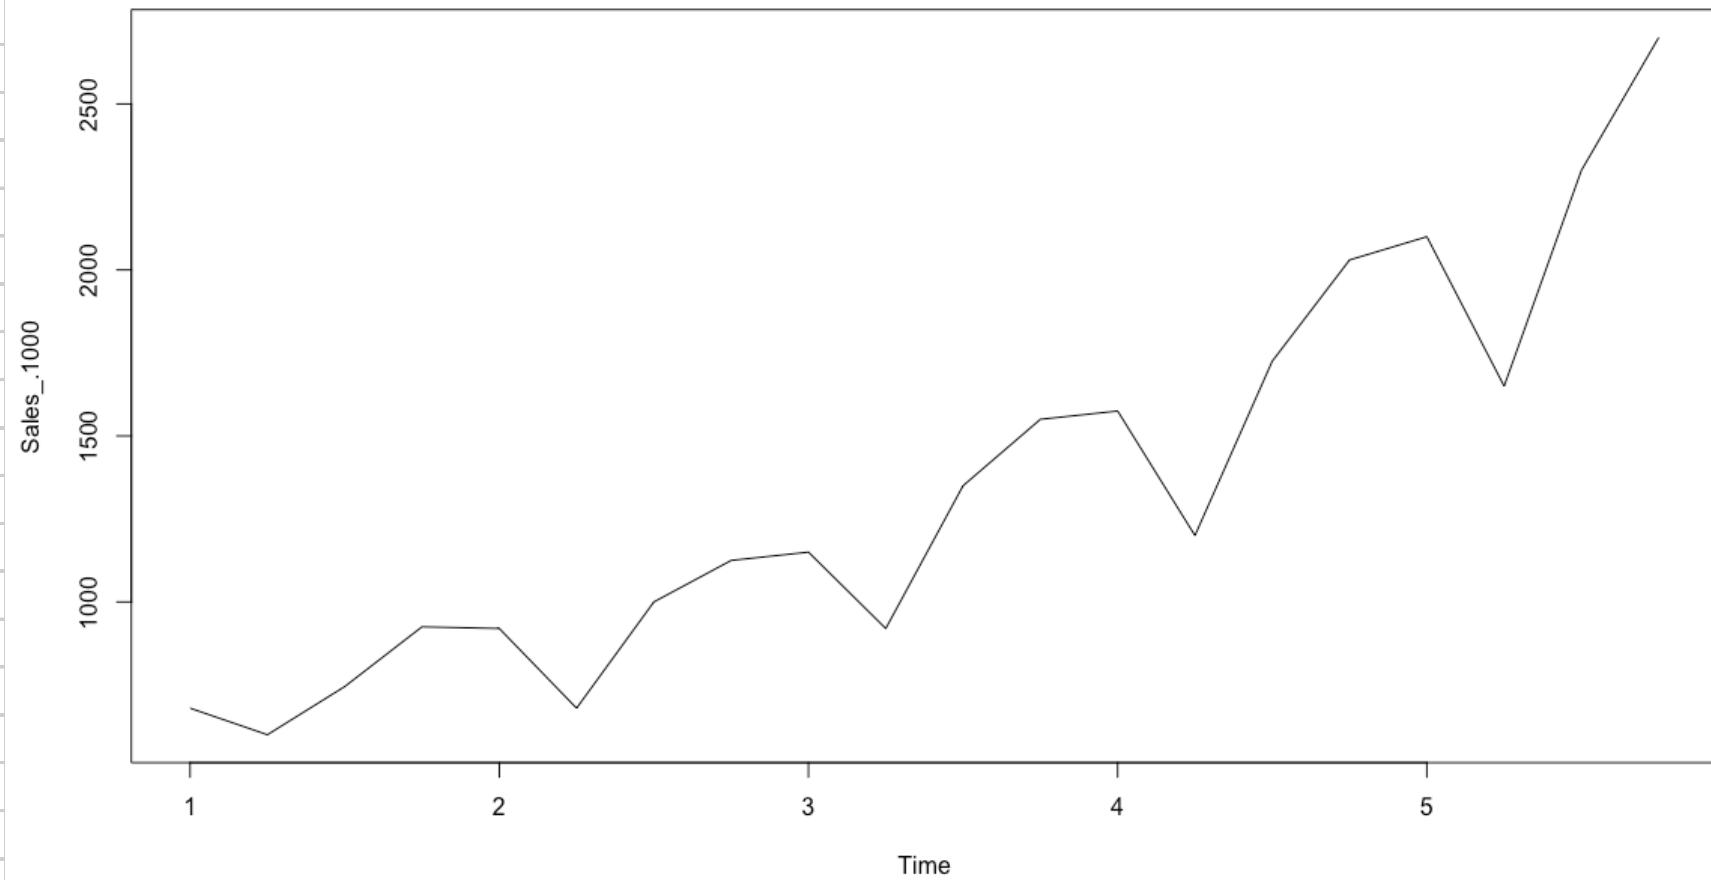
CURVE FITTING / REGRESSION ON TIME METHODS



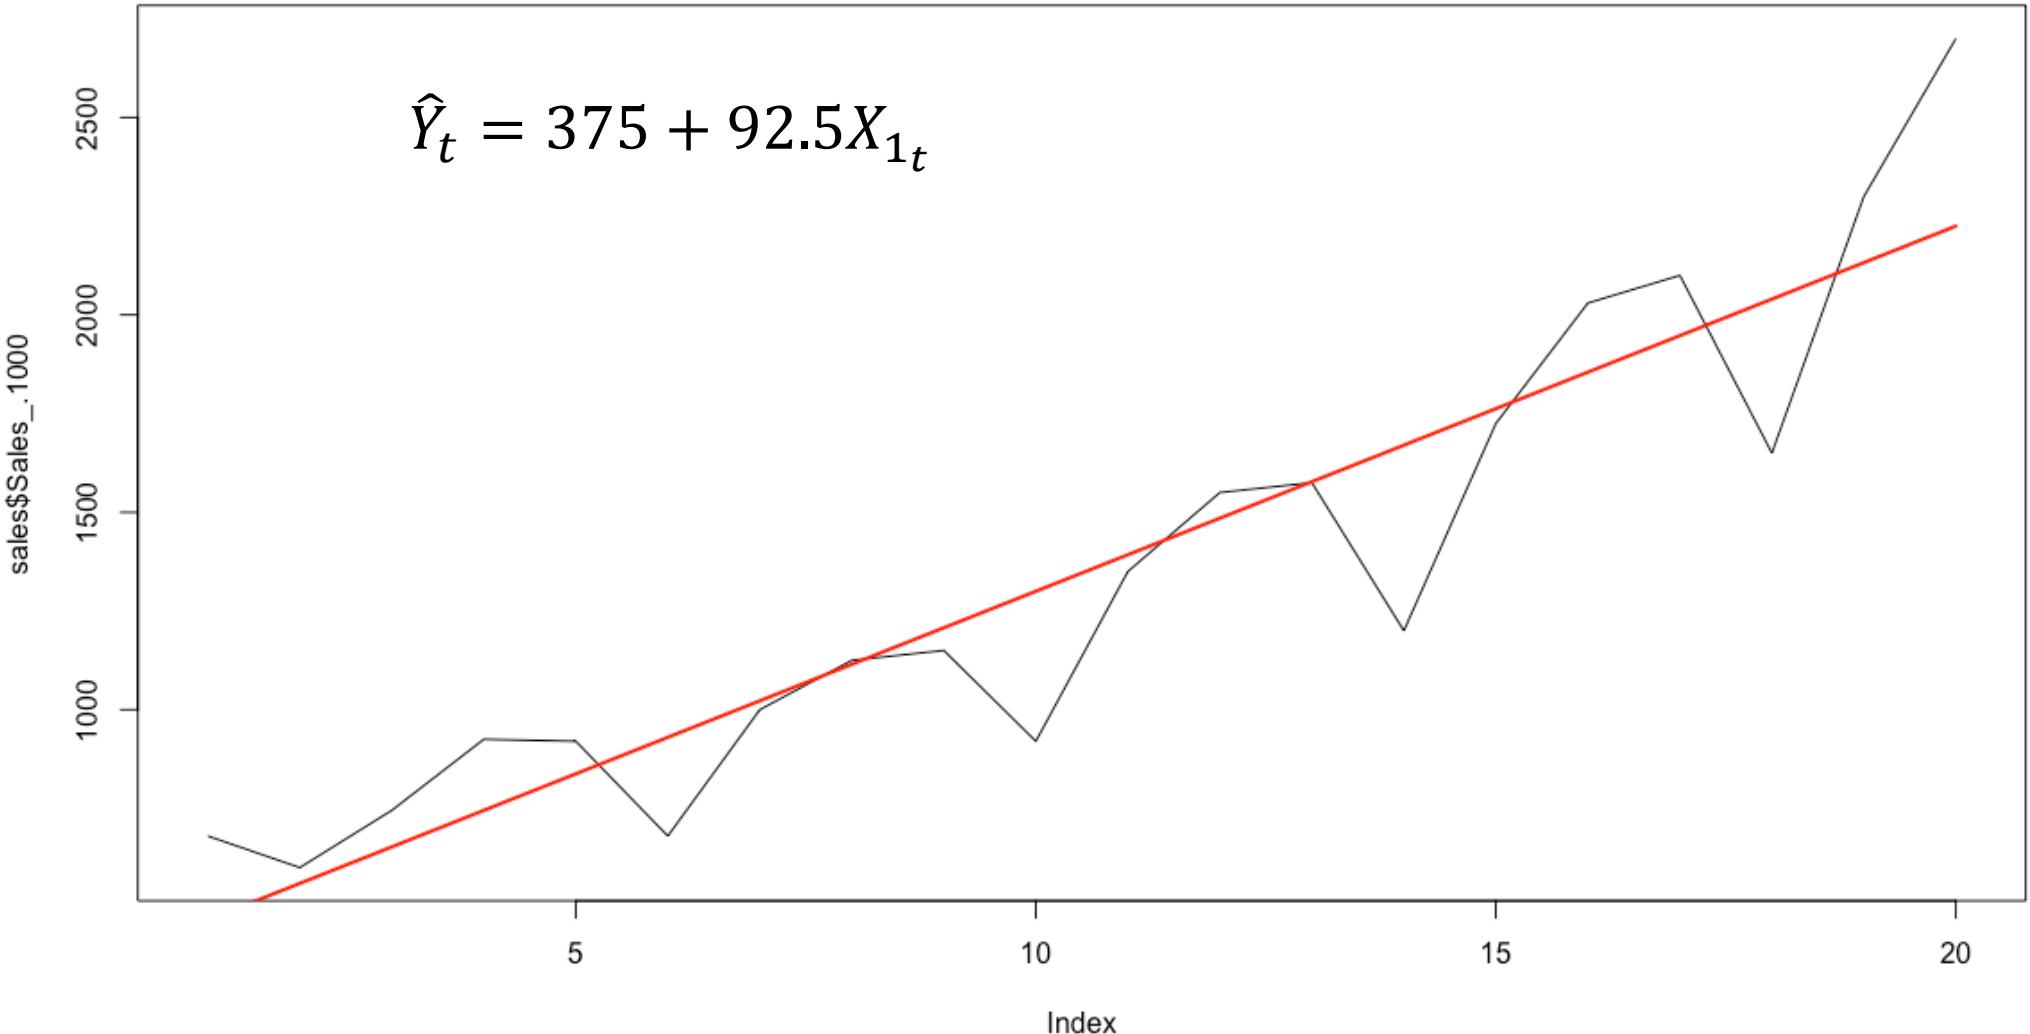
Regression on Time

- Use when trend is the most pronounced
- ACF decays exponentially and PACF has very few spikes

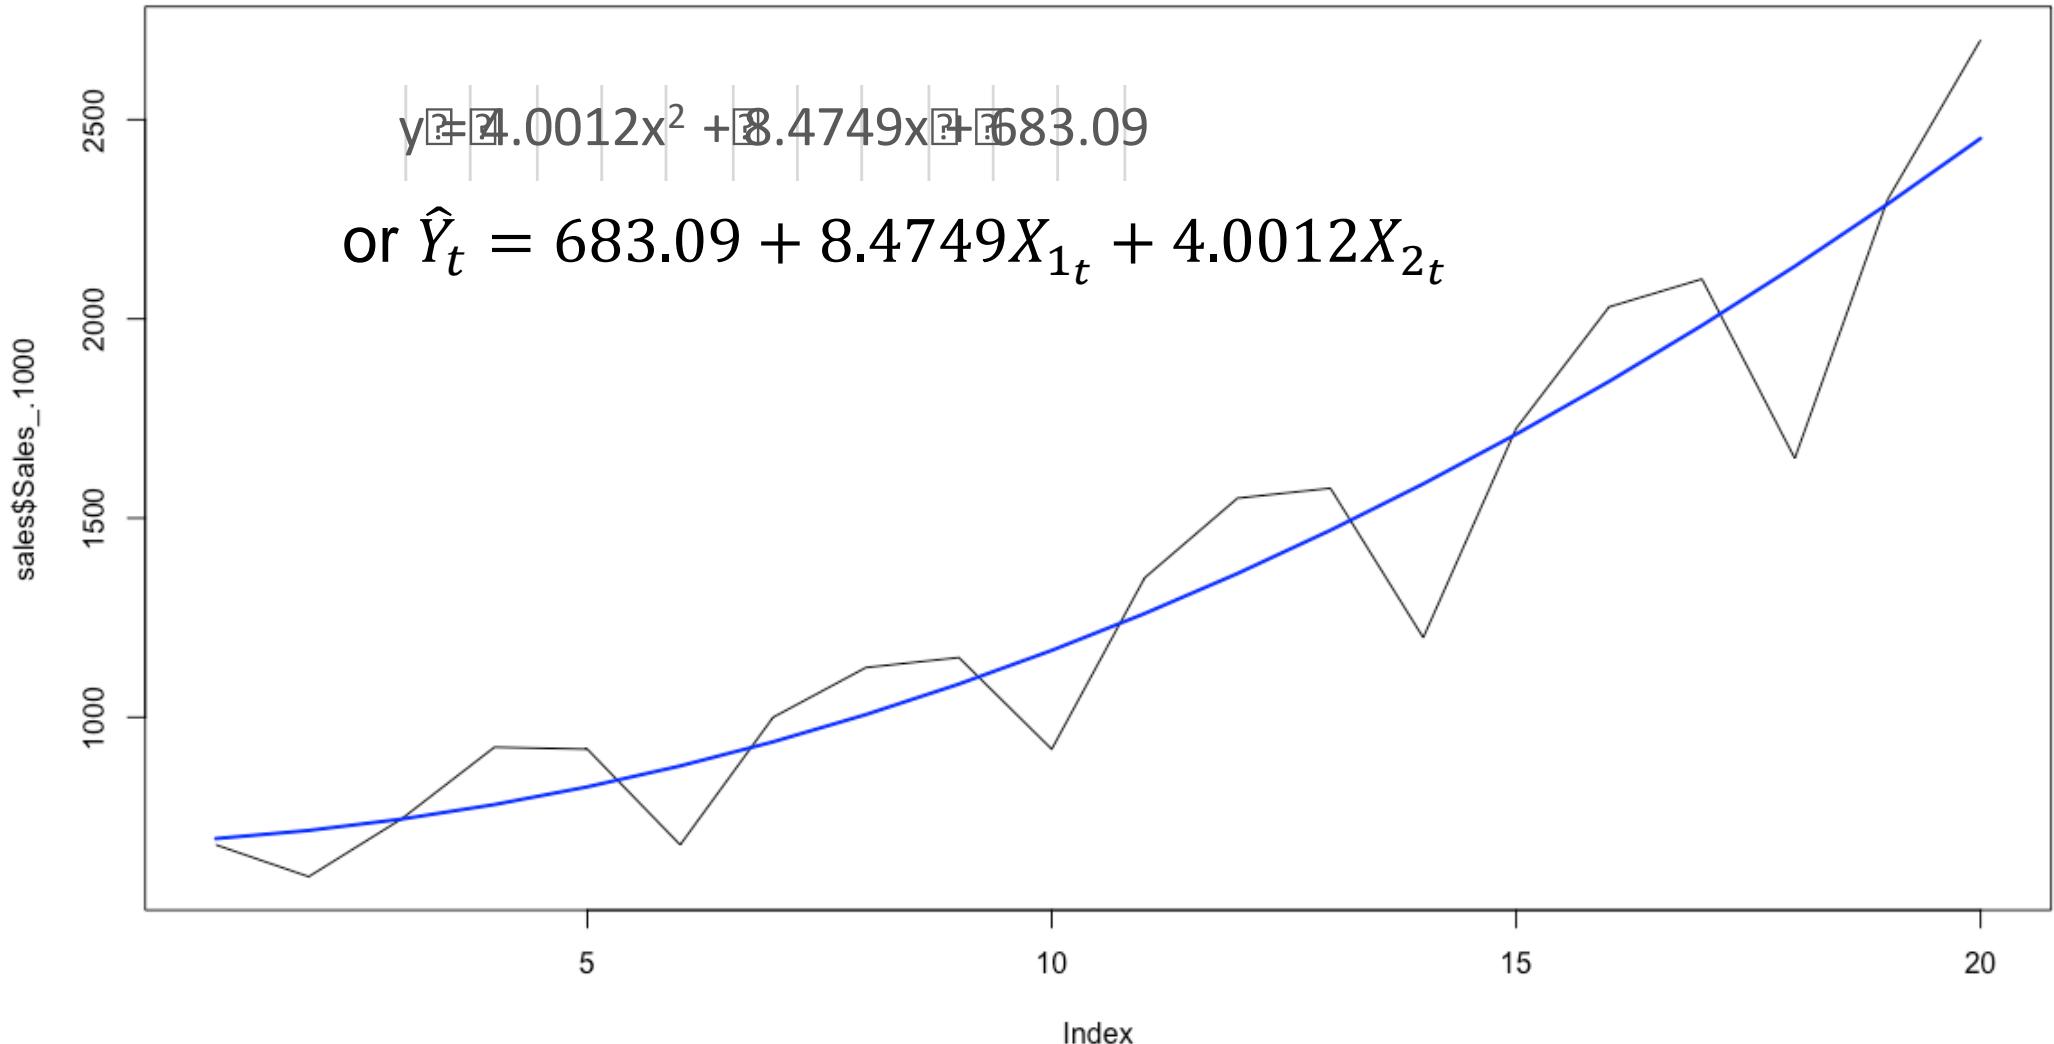
Quarter	Sales_\$1000
1	680
2	600
3	745
4	925
5	920
6	680
7	1000
8	1125
9	1150
10	920
11	1350
12	1550
13	1575
14	1200
15	1725
16	2030
17	2100
18	1650
19	2300
20	2700



Regression Analysis



Quadratic Trend

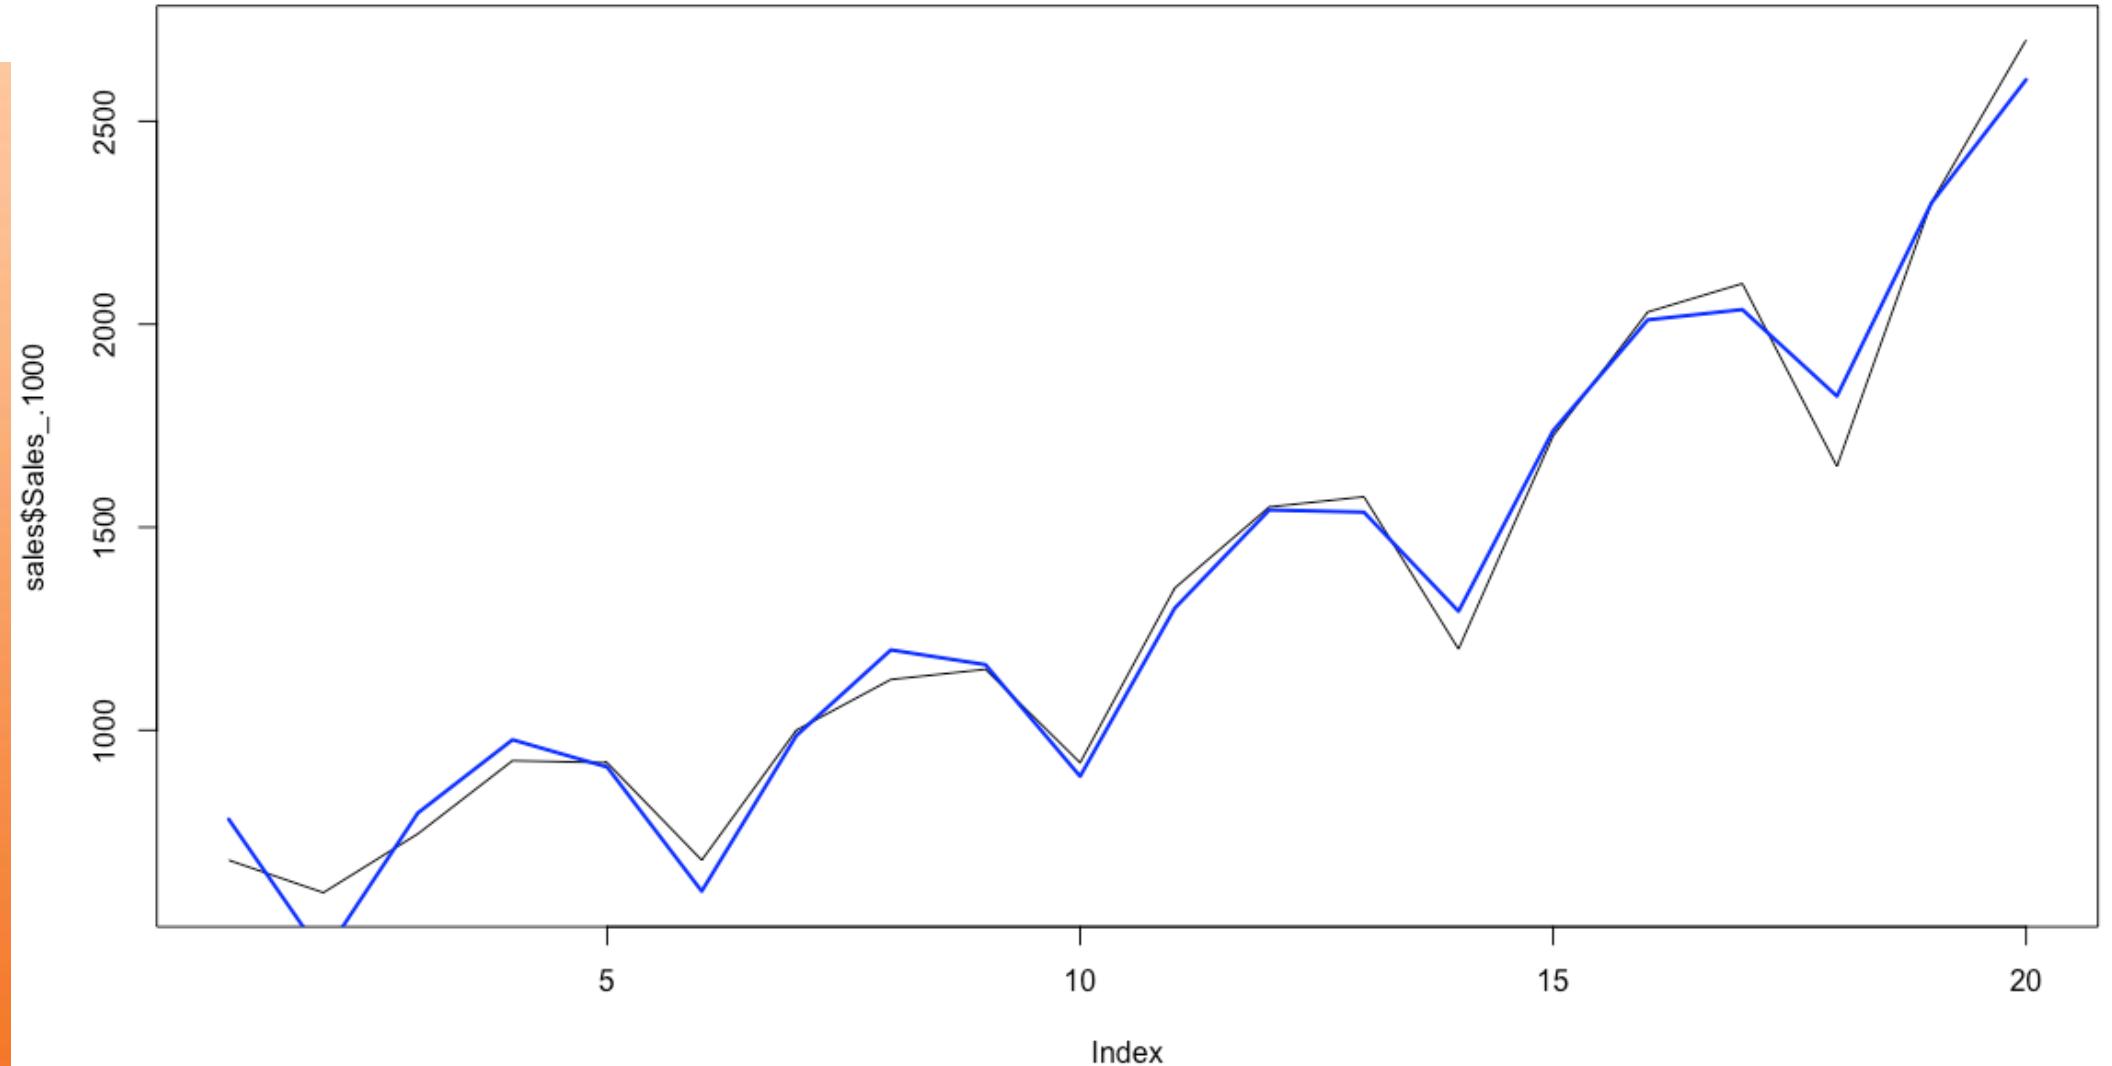


Seasonal Regression Models

Quarter	Value of		
	X _{3t}	X _{4t}	X _{5t}
1	1	0	0
2	0	1	0
3	0	0	1
4	0	0	0

$$\hat{Y}_t = \beta_0 + \beta_1 X_{1t} + \beta_2 X_{2t} + \beta_3 X_{3t} + \beta_4 X_{4t} + \beta_5 X_{5t} + \varepsilon_t$$

Seasonal Regression Models



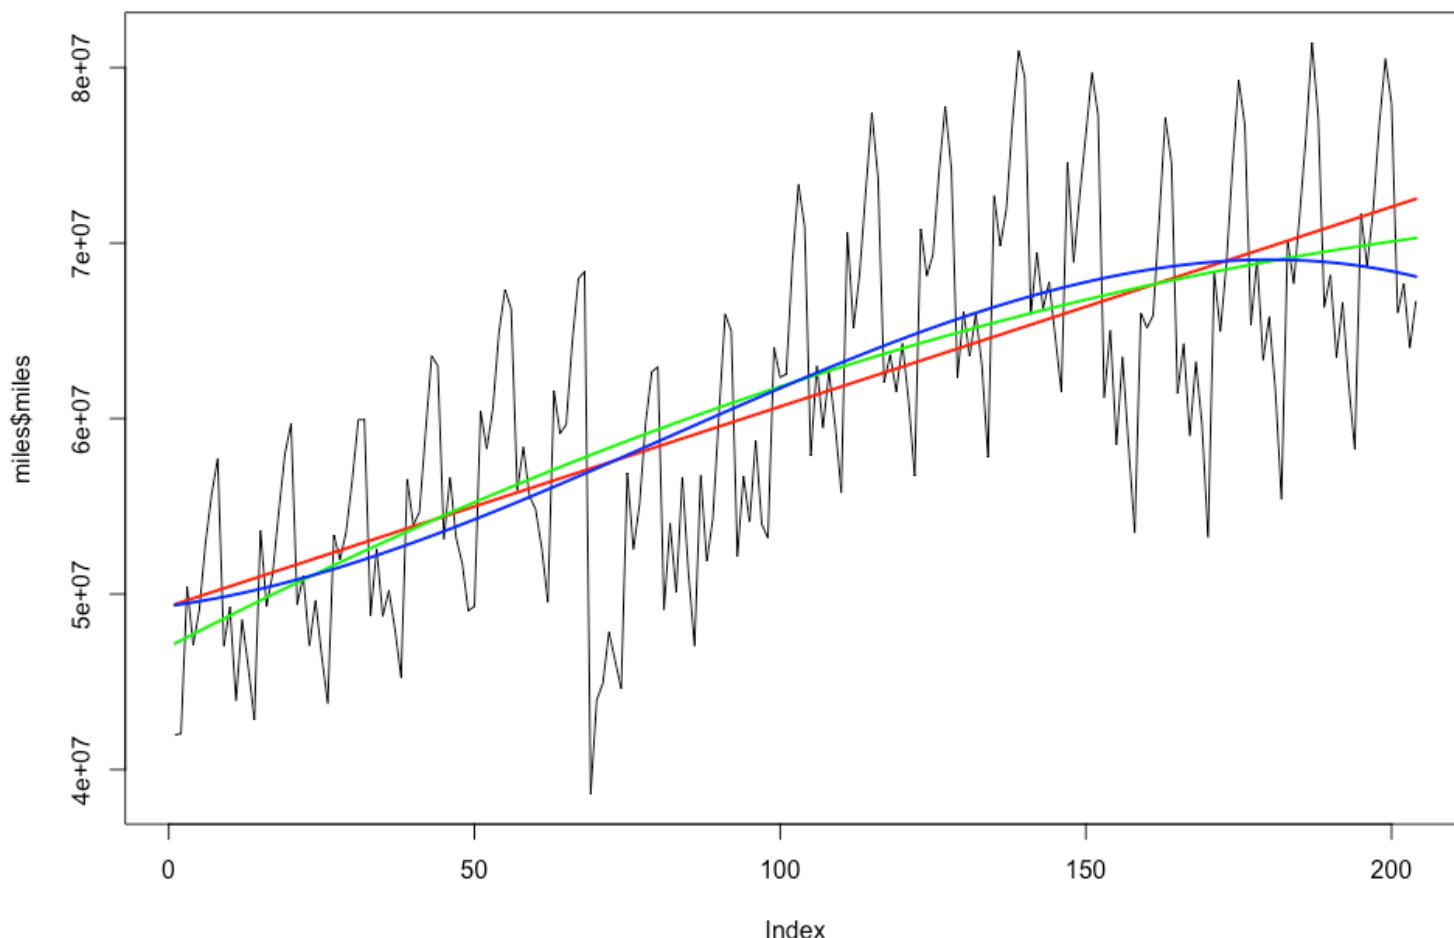
Seasonal Regression Models



CSE 7202c



Seasonal Regression Models - RPM

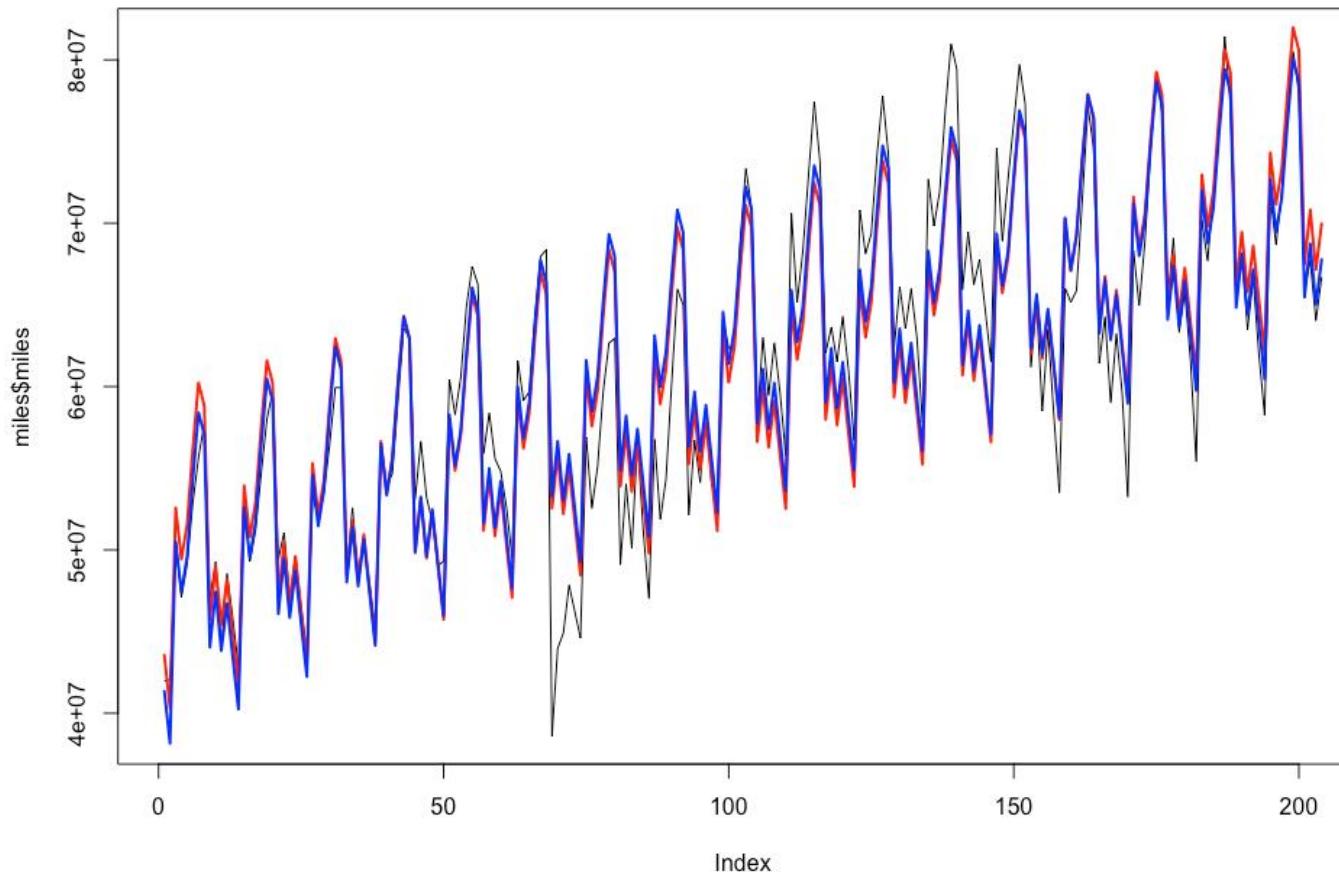


	miles	time	var3
1	41972194	1	
2	42054796	2	
3	50443045	3	
4	47112397	4	
5	49118248	5	
6	52880510	6	
7	55664750	7	
8	57723208	8	
9	47035464	9	
10	49263120	10	
11	43937074	11	
12	48539606	12	
13	45850623	13	
14	42838949	14	
15	53620994	15	

72026



Seasonal Regression Models - RPM



	miles	time	seasonal
1	41972194	1	1
2	42054796	2	2
3	50443045	3	3
4	47112397	4	4
5	49118248	5	5
6	52880510	6	6
7	55664750	7	7
8	57723208	8	8
9	47035464	9	9
10	49263120	10	10
11	43937074	11	11
12	48539606	12	12
13	45850623	13	1
14	42838949	14	2
15	53620994	15	3

EE
72026



Another Way of Incorporating Seasonality

- Take the trend prediction and actual prediction.
- Depending on additive or multiplicative model compute the deviation and map it as seasonality effect for each prediction.
- Take averages of the seasonality value. Use this to make future predictions.

Case

Year	Quarter	Time variable (this is created)	Revenues (in \$M)
2008	I	1	10.2
	II	2	12.4
	III	3	14.8
	IV	4	15
2009	I	5	11.2
	II	6	14.3
	III	7	18.4
	IV	8	18

Call:
lm(formula = y ~ x)

Residuals:

Min	1Q	Median	3Q	Max
-3.5595	-0.9384	0.4405	1.3265	1.9286

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	10.0393	1.5531	6.464	0.00065 ***
x	0.9440	0.3076	3.069	0.02196 *

Signif. codes: 0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.993 on 6 degrees of freedom

Multiple R-squared: 0.6109, Adjusted R-squared: 0.5461

F-statistic: 9.422 on 1 and 6 DF, p-value: 0.02196

What is the Regression equation?

$$y = 10.0393 + 0.9440x$$

Seasonality: Multiplicative

Time	Observed values TSI* (assuming no impact of cyclicality)	Predicted values (per the regression) T*	SI* = TSI/T
1	10.2	10.983	0.929
2	12.4	11.927	1.040
3	14.8	12.871	1.150
4	15.0	13.815	1.086
5	11.2	14.759	0.759
6	14.3	15.703	0.911
7	18.4	16.647	1.105
8	18.0	17.591	1.023

* T: Trend; S: Seasonal; I: Irregular

Quarterly Seasonality

Time	Average seasonality factor
Q1	$0.844 \left(= \frac{0.929+0.759}{2} \right)$
Q2	0.975
Q3	1.127
Q4	1.054

Time	Observed values TSI* (assuming no impact of cyclicity)	Predicted values (per the regression) T*	SI* = TSI/T
1	10.2	10.983	0.929
2	12.4	11.927	1.040
3	14.8	12.871	1.150
4	15.0	13.815	1.086
5	11.2	14.759	0.759
6	14.3	15.703	0.911
7	18.4	16.647	1.105
8	18.0	17.591	1.023

Computations

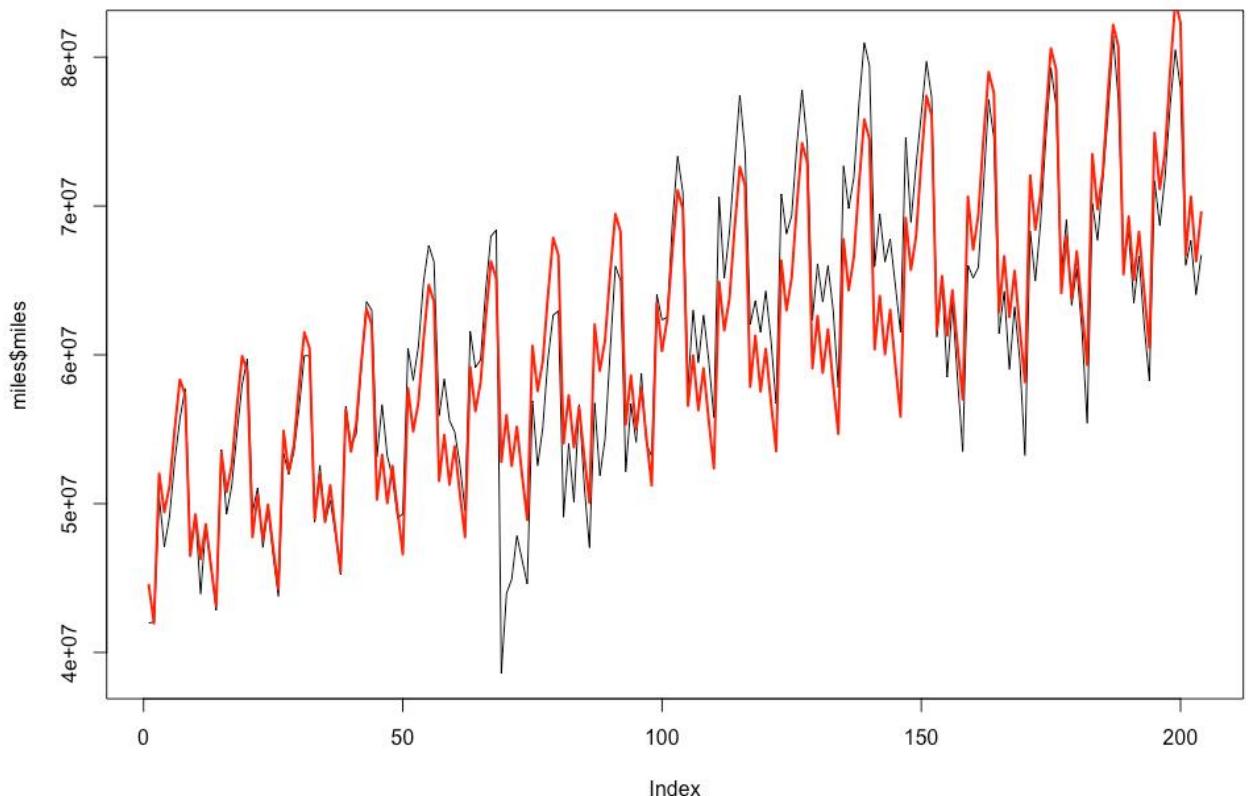
- Trend $Y_9 = 10.039 + 0.944(9) = 18.535$
- Corrected for seasonality and randomness: $18.535 * 0.844 = 15.643$



CSE 7202C

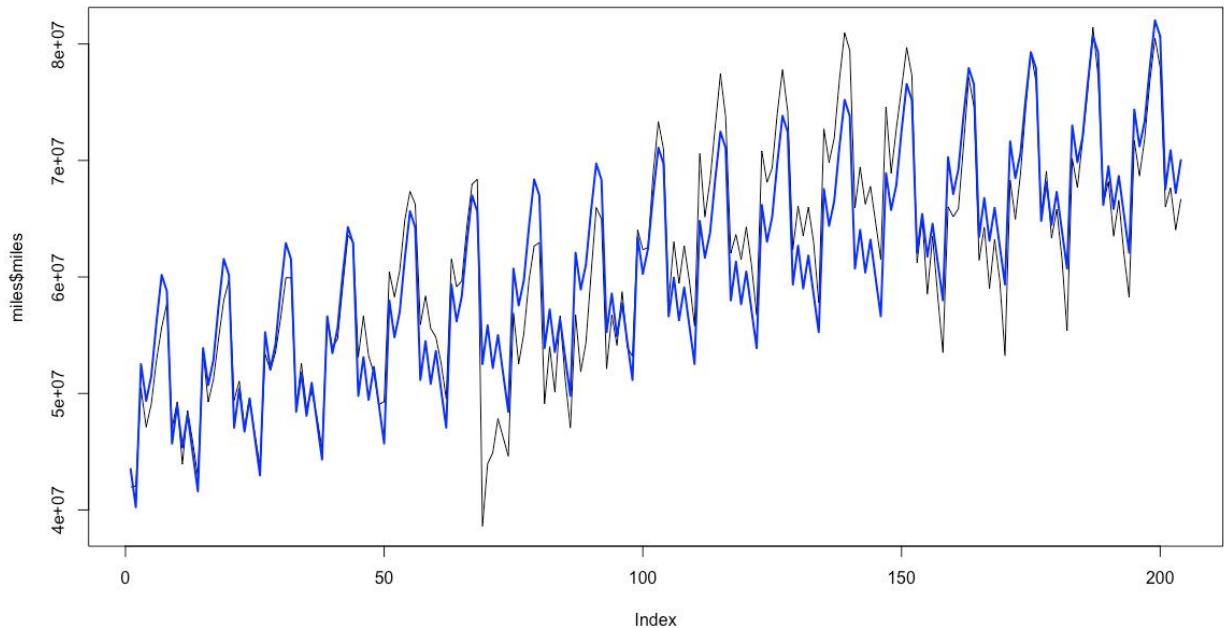


Seasonality: Multiplicative



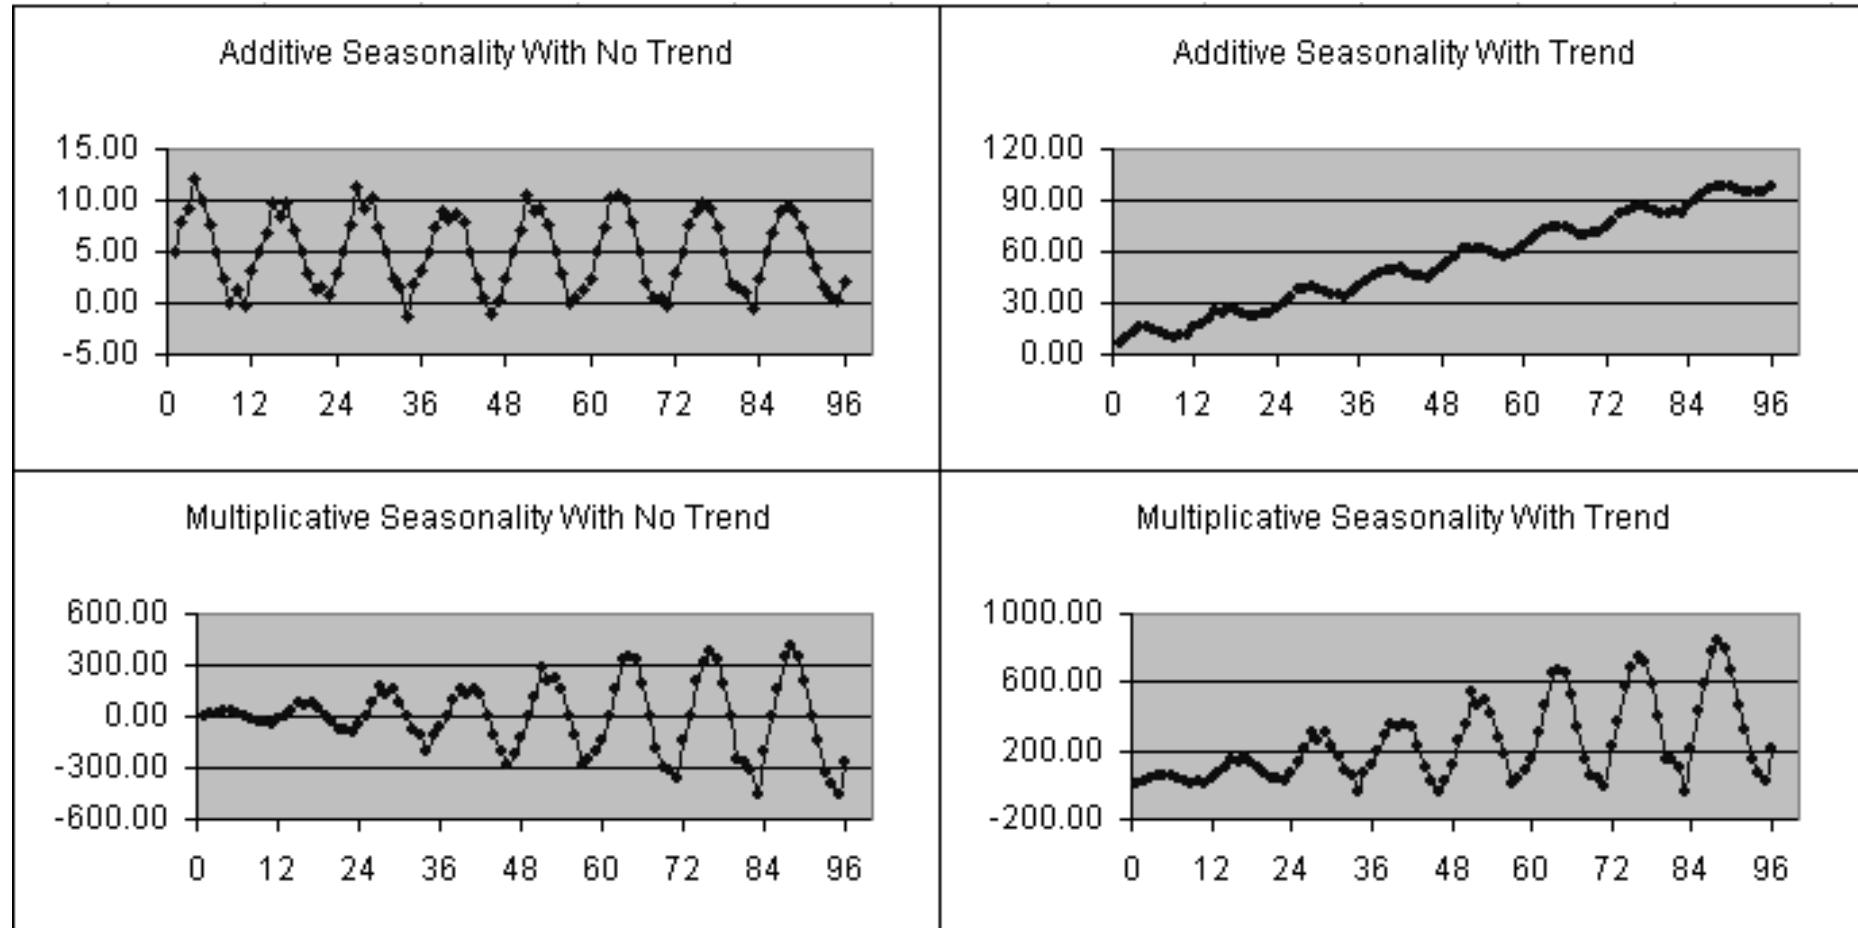
	miles	time	seasonal	mae
1	41972194	1	1	0.849386
2	42054796	2	2	0.8491019
3	50443045	3	3	1.016129
4	47112397	4	4	0.946865
5	49118248	5	5	0.9849257
6	52880510	6	6	1.057953
7	55664750	7	7	1.111125
8	57723208	8	8	1.149602
9	47035464	9	9	0.9346292
10	49263120	10	10	0.9766855
11	43937074	11	11	0.8691307
12	48539606	12	12	0.9580177
13	45850623	13	1	0.9029174
14	42838949	14	2	0.8417232
15	53620994	15	3	1.051224

Seasonality: Additive



	miles	time	seasonal	mae
1	41972194	1	1	-7442550
2	42054796	2	2	-7473763
3	50443045	3	3	800670.5
4	47112397	4	4	-2643793
5	49118248	5	5	-751757.1
6	52880510	6	6	2896690
7	55664750	7	7	5567114
8	57723208	8	8	7511757
9	47035464	9	9	-3289802
10	49263120	10	10	-1175962
11	43937074	11	11	-6615823
12	48539606	12	12	-2127106
13	45850623	13	1	-4929905
14	42838949	14	2	-8055394
15	53620994	15	3	2612836
.

Additive or Multiplicative



Issues with Regressing on Time

- If there is no trend or if seasonality and fluctuations are more important than trend, then the coefficients behave weirdly



Goodness of Fit

- MAE (Mean absolute error)

$$\frac{\sum |y_i - \hat{y}_i|}{n}$$

- MSE (Mean square error)

$$\frac{\sum (y_i - \hat{y}_i)^2}{n}$$

- RMSE (Root mean square error)

$$\sqrt{\frac{\sum (y_i - \hat{y}_i)^2}{n}}$$

- MAPE (Mean absolute percent error)

$$\frac{1}{n} \left(\frac{\sum |y_i - \hat{y}_i|}{y_i} \right) * 100$$



Time Series

SIMPLE FORECASTING / BENCHMARKING METHODS



Average Method

$$\hat{y}_{t+h} = \bar{y} = \frac{y_1 + \cdots + y_t}{t}$$

- Forecasts of all future values are equal to the **mean of the historical data**
- R – meanf(y(ts),h)

Naïve Method

$$\hat{y}_{t+h} = y_t$$

- Forecasts of all future values are equal to the last **observed** value
- A very useful method in many economic and financial time series
- R – naive(y(ts),h) or r wf(y(ts),h)

Seasonal Naïve Method

$$\hat{y}_{t+h} = y_{t+h-km}$$

Integer part

where m is the seasonal period and $k = \left\lfloor \frac{h-1}{m} \right\rfloor + 1$

- Forecasts of all future values for a particular period (e.g., July) are equal to the last **observed** value for the same period (e.g., last July)
- R – snaive(y(ts),h)

Drift Method

$$\hat{y}_{t+h} = y_t + \frac{h}{t-1} \sum_{t=2}^h (y_t - y_{t-1}) = y_t + h \left(\frac{y_t - y_1}{t-1} \right)$$

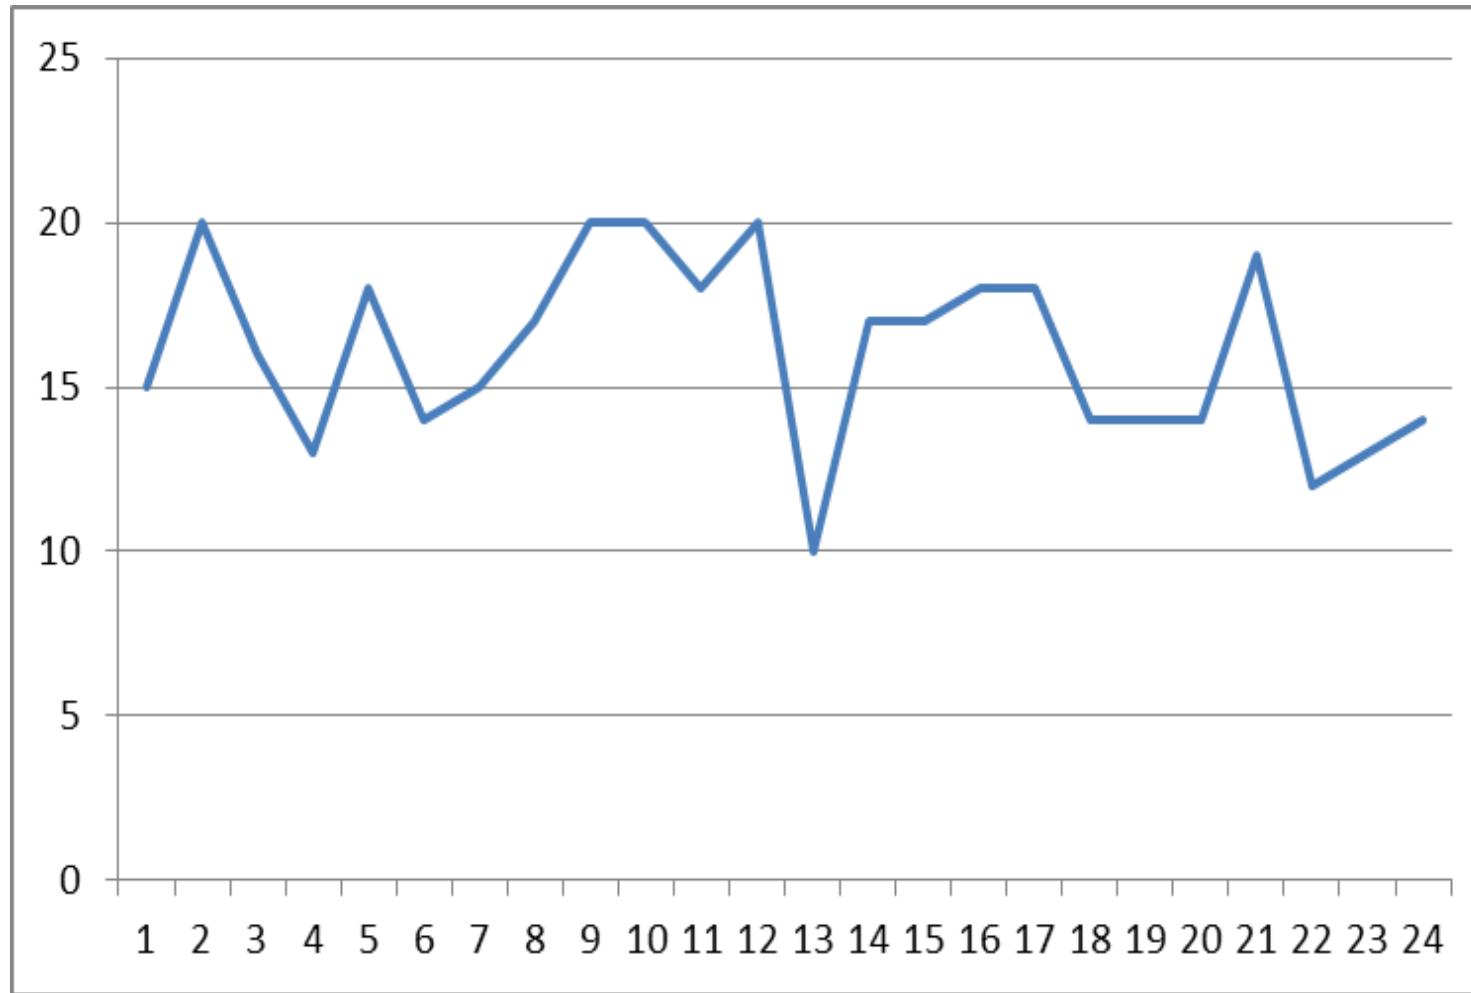
- Forecasts increase or decrease over time
- Amount of change over time (drift) is the average historical change (*In Naïve method, this change is zero*)
- R – rwf(y(ts),h, drift=TRUE)

Time Series

SMOOTHING METHODS

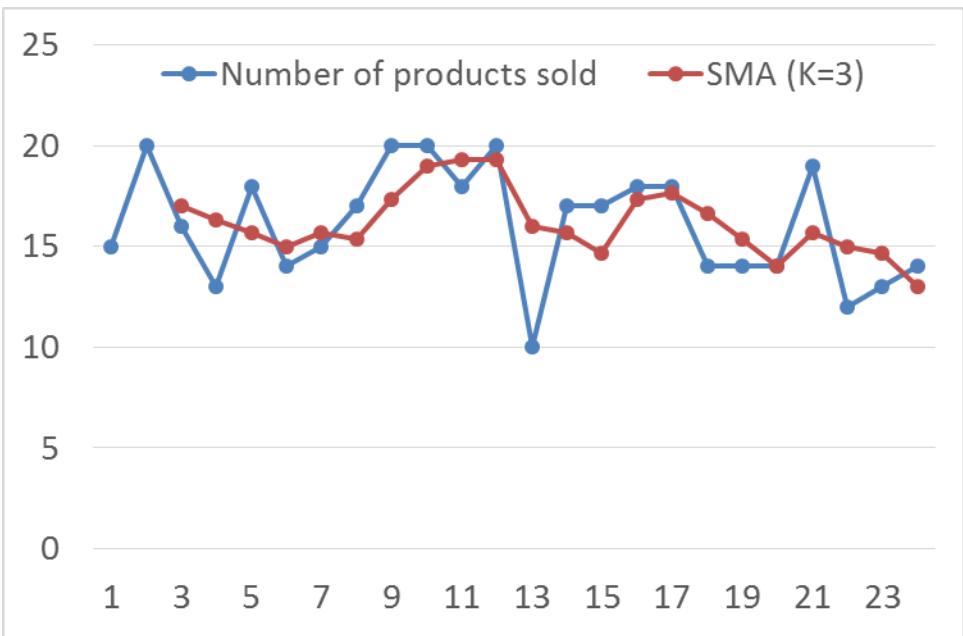
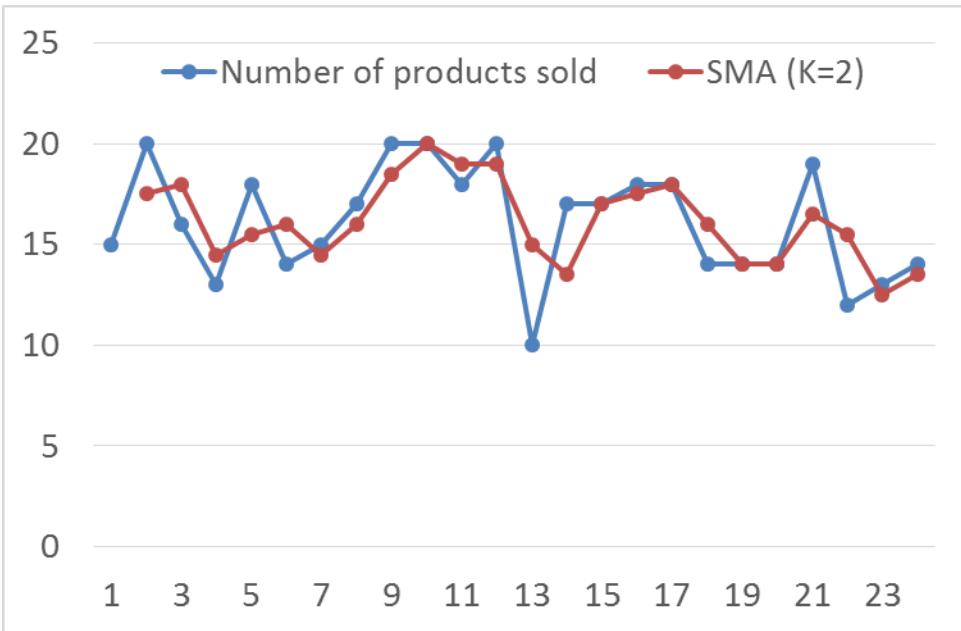
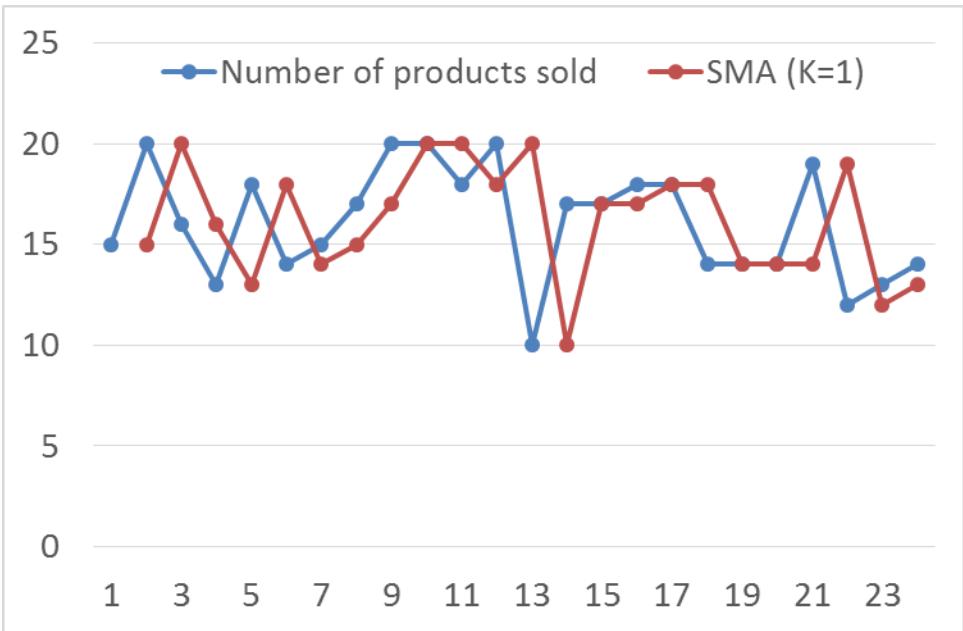


Stationary Model: Moving Averages



Stationary Model: Case 1 – Simple Moving Averages

	Number of products sold	SMA (K=1)	Error	SMA (K=2)	Error	SMA (K=3)	Error
1							
2	15						
3	20	15	5	17.5	2.5		
4	16	20	4	18	2	17	1
5	13	16	3	14.5	1.5	16.333333	3.33333
6	18	13	5	15.5	2.5	15.666667	2.33333
7	14	18	4	16	2	15	1
8	15	14	1	14.5	0.5	15.666667	0.66667
9	17	15	2	16	1	15.333333	1.66667
10	20	17	3	18.5	1.5	17.333333	2.66667
11	20	20	0	20	0	19	1
12	18	20	2	19	1	19.333333	1.33333
13	20	18	2	19	1	19.333333	0.66667
14	10	20	10	15	5	16	6
15	17	10	7	13.5	3.5	15.666667	1.33333
16	17	17	0	17	0	14.666667	2.33333
17	18	17	1	17.5	0.5	17.333333	0.66667
18	18	18	0	18	0	17.666667	0.33333
19	14	18	4	16	2	16.666667	2.66667
20	14	14	0	14	0	15.333333	1.33333
21	14	14	0	14	0	14	0
22	19	14	5	16.5	2.5	15.666667	3.33333
23	12	19	7	15.5	3.5	15	3
24	13	12	1	12.5	0.5	14.666667	1.66667
25	14	13	1	13.5	0.5	13	1
26			2.91304		1.45652		1.78788



Only decision point is K

Stationary Model: Case 2 – Weighted Moving Averages

$$\hat{Y}_{t+1} = w_1 Y_t + w_2 Y_{t-1} + \cdots + w_k Y_{t-k+1}$$

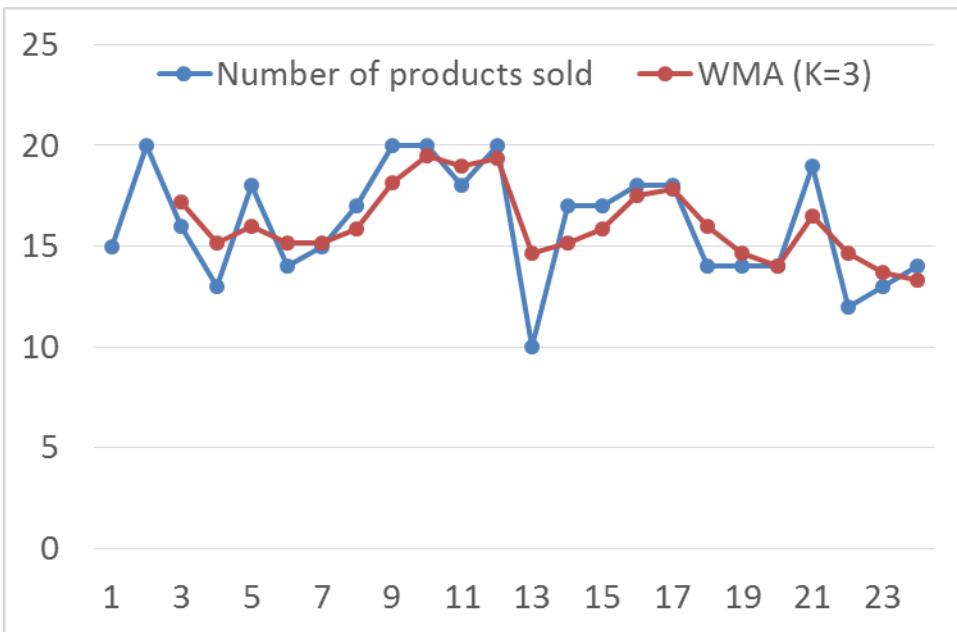
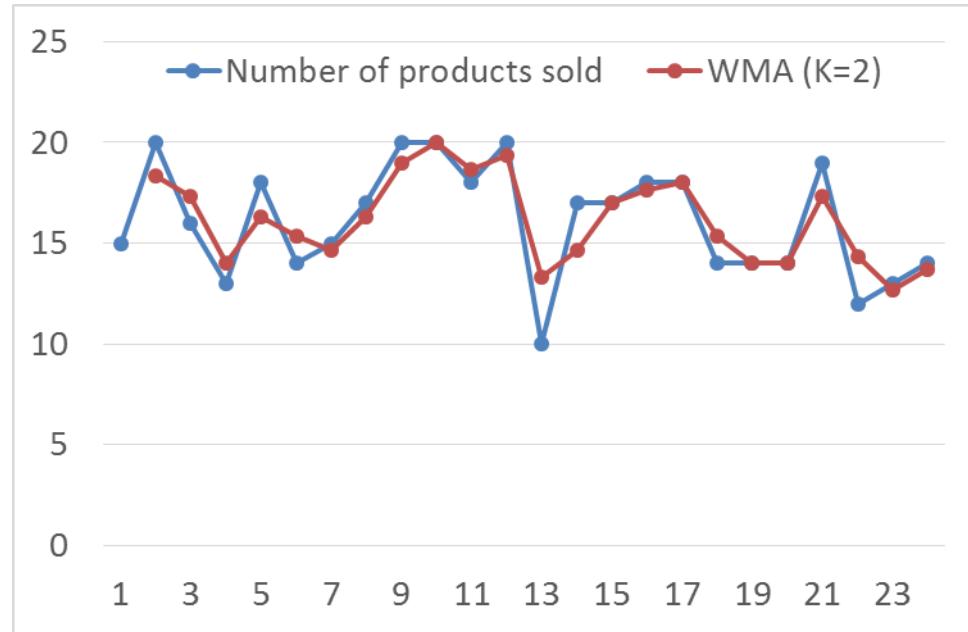
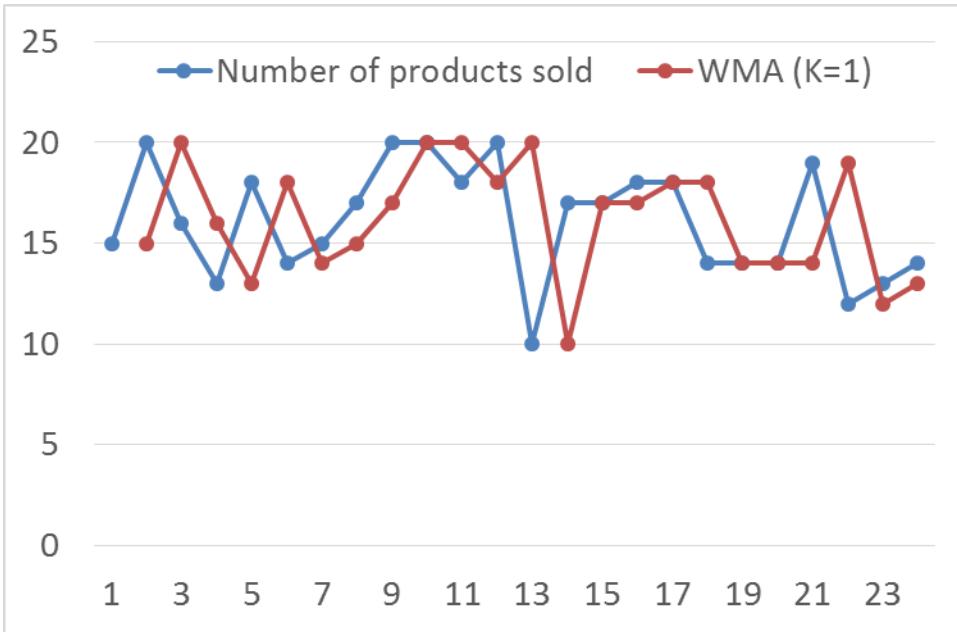
- Typically we choose a time period of moving average and weights are chosen such that the error is minimized

Stationary Model: Case 2 – Weighted Moving Averages

	Number of products sold	WMA (K=1)	Error	WMA (K=2)	Error	WMA (K=3)	Error
1							
2	15						
3	20	=A2*1	=ABS(B3-A3)	=(A3*2+A2*1)/3	=ABS(A3-D3)		
4	16	=A3*1	=ABS(B4-A4)	=(A4*2+A3*1)/3	=ABS(A4-D4)	=(A4*3+A3*2+A2*1)/6	=ABS(A4-F4)
5	13	=A4*1	=ABS(B5-A5)	=(A5*2+A4*1)/3	=ABS(A5-D5)	=(A5*3+A4*2+A3*1)/6	=ABS(A5-F5)
6	18	=A5*1	=ABS(B6-A6)	=(A6*2+A5*1)/3	=ABS(A6-D6)	=(A6*3+A5*2+A4*1)/6	=ABS(A6-F6)
7	14	=A6*1	=ABS(B7-A7)	=(A7*2+A6*1)/3	=ABS(A7-D7)	=(A7*3+A6*2+A5*1)/6	=ABS(A7-F7)
8	15	=A7*1	=ABS(B8-A8)	=(A8*2+A7*1)/3	=ABS(A8-D8)	=(A8*3+A7*2+A6*1)/6	=ABS(A8-F8)
9	17	=A8*1	=ABS(B9-A9)	=(A9*2+A8*1)/3	=ABS(A9-D9)	=(A9*3+A8*2+A7*1)/6	=ABS(A9-F9)
10	20	=A9*1	=ABS(B10-A10)	=(A10*2+A9*1)/3	=ABS(A10-D10)	=(A10*3+A9*2+A8*1)/6	=ABS(A10-F10)
11	20	=A10*1	=ABS(B11-A11)	=(A11*2+A10*1)/3	=ABS(A11-D11)	=(A11*3+A10*2+A9*1)/6	=ABS(A11-F11)
12	18	=A11*1	=ABS(B12-A12)	=(A12*2+A11*1)/3	=ABS(A12-D12)	=(A12*3+A11*2+A10*1)/6	=ABS(A12-F12)
13	20	=A12*1	=ABS(B13-A13)	=(A13*2+A12*1)/3	=ABS(A13-D13)	=(A13*3+A12*2+A11*1)/6	=ABS(A13-F13)
14	10	=A13*1	=ABS(B14-A14)	=(A14*2+A13*1)/3	=ABS(A14-D14)	=(A14*3+A13*2+A12*1)/6	=ABS(A14-F14)
15	17	=A14*1	=ABS(B15-A15)	=(A15*2+A14*1)/3	=ABS(A15-D15)	=(A15*3+A14*2+A13*1)/6	=ABS(A15-F15)
16	17	=A15*1	=ABS(B16-A16)	=(A16*2+A15*1)/3	=ABS(A16-D16)	=(A16*3+A15*2+A14*1)/6	=ABS(A16-F16)
17	18	=A16*1	=ABS(B17-A17)	=(A17*2+A16*1)/3	=ABS(A17-D17)	=(A17*3+A16*2+A15*1)/6	=ABS(A17-F17)
18	18	=A17*1	=ABS(B18-A18)	=(A18*2+A17*1)/3	=ABS(A18-D18)	=(A18*3+A17*2+A16*1)/6	=ABS(A18-F18)
19	14	=A18*1	=ABS(B19-A19)	=(A19*2+A18*1)/3	=ABS(A19-D19)	=(A19*3+A18*2+A17*1)/6	=ABS(A19-F19)
20	14	=A19*1	=ABS(B20-A20)	=(A20*2+A19*1)/3	=ABS(A20-D20)	=(A20*3+A19*2+A18*1)/6	=ABS(A20-F20)
21	14	=A20*1	=ABS(B21-A21)	=(A21*2+A20*1)/3	=ABS(A21-D21)	=(A21*3+A20*2+A19*1)/6	=ABS(A21-F21)
22	19	=A21*1	=ABS(B22-A22)	=(A22*2+A21*1)/3	=ABS(A22-D22)	=(A22*3+A21*2+A20*1)/6	=ABS(A22-F22)
23	12	=A22*1	=ABS(B23-A23)	=(A23*2+A22*1)/3	=ABS(A23-D23)	=(A23*3+A22*2+A21*1)/6	=ABS(A23-F23)
24	13	=A23*1	=ABS(B24-A24)	=(A24*2+A23*1)/3	=ABS(A24-D24)	=(A24*3+A23*2+A22*1)/6	=ABS(A24-F24)
25	14	=A24*1	=ABS(B25-A25)	=(A25*2+A24*1)/3	=ABS(A25-D25)	=(A25*3+A24*2+A23*1)/6	=ABS(A25-F25)
26			=AVERAGE(C3:C25)		=AVERAGE(E3:E25)		=AVERAGE(G3:G25)

Stationary Model: Case 2 – Weighted Moving Averages

1	Number of products sold	WMA (K=1)	Error	WMA (K=2)	Error	WMA (K=3)	Error
2	15						
3	20	15	5	18.3333333	1.66666667		
4	16	20	4	17.3333333	1.33333333	17.1666667	1.16666667
5	13	16	3	14	1	15.1666667	2.16666667
6	18	13	5	16.3333333	1.66666667	16	2
7	14	18	4	15.3333333	1.33333333	15.1666667	1.16666667
8	15	14	1	14.6666667	0.33333333	15.1666667	0.16666667
9	17	15	2	16.3333333	0.66666667	15.8333333	1.16666667
10	20	17	3	19	1	18.1666667	1.83333333
11	20	20	0	20	0	19.5	0.5
12	18	20	2	18.6666667	0.66666667	19	1
13	20	18	2	19.3333333	0.66666667	19.3333333	0.66666667
14	10	20	10	13.3333333	3.33333333	14.6666667	4.66666667
15	17	10	7	14.6666667	2.33333333	15.1666667	1.83333333
16	17	17	0	17	0	15.8333333	1.16666667
17	18	17	1	17.6666667	0.33333333	17.5	0.5
18	18	18	0	18	0	17.8333333	0.16666667
19	14	18	4	15.3333333	1.33333333	16	2
20	14	14	0	14	0	14.6666667	0.66666667
21	14	14	0	14	0	14	0
22	19	14	5	17.3333333	1.66666667	16.5	2.5
23	12	19	7	14.3333333	2.33333333	14.6666667	2.66666667
24	13	12	1	12.6666667	0.33333333	13.6666667	0.66666667
25	14	13	1	13.6666667	0.33333333	13.3333333	0.66666667
26			2.91304348		0.97101449		1.33333333

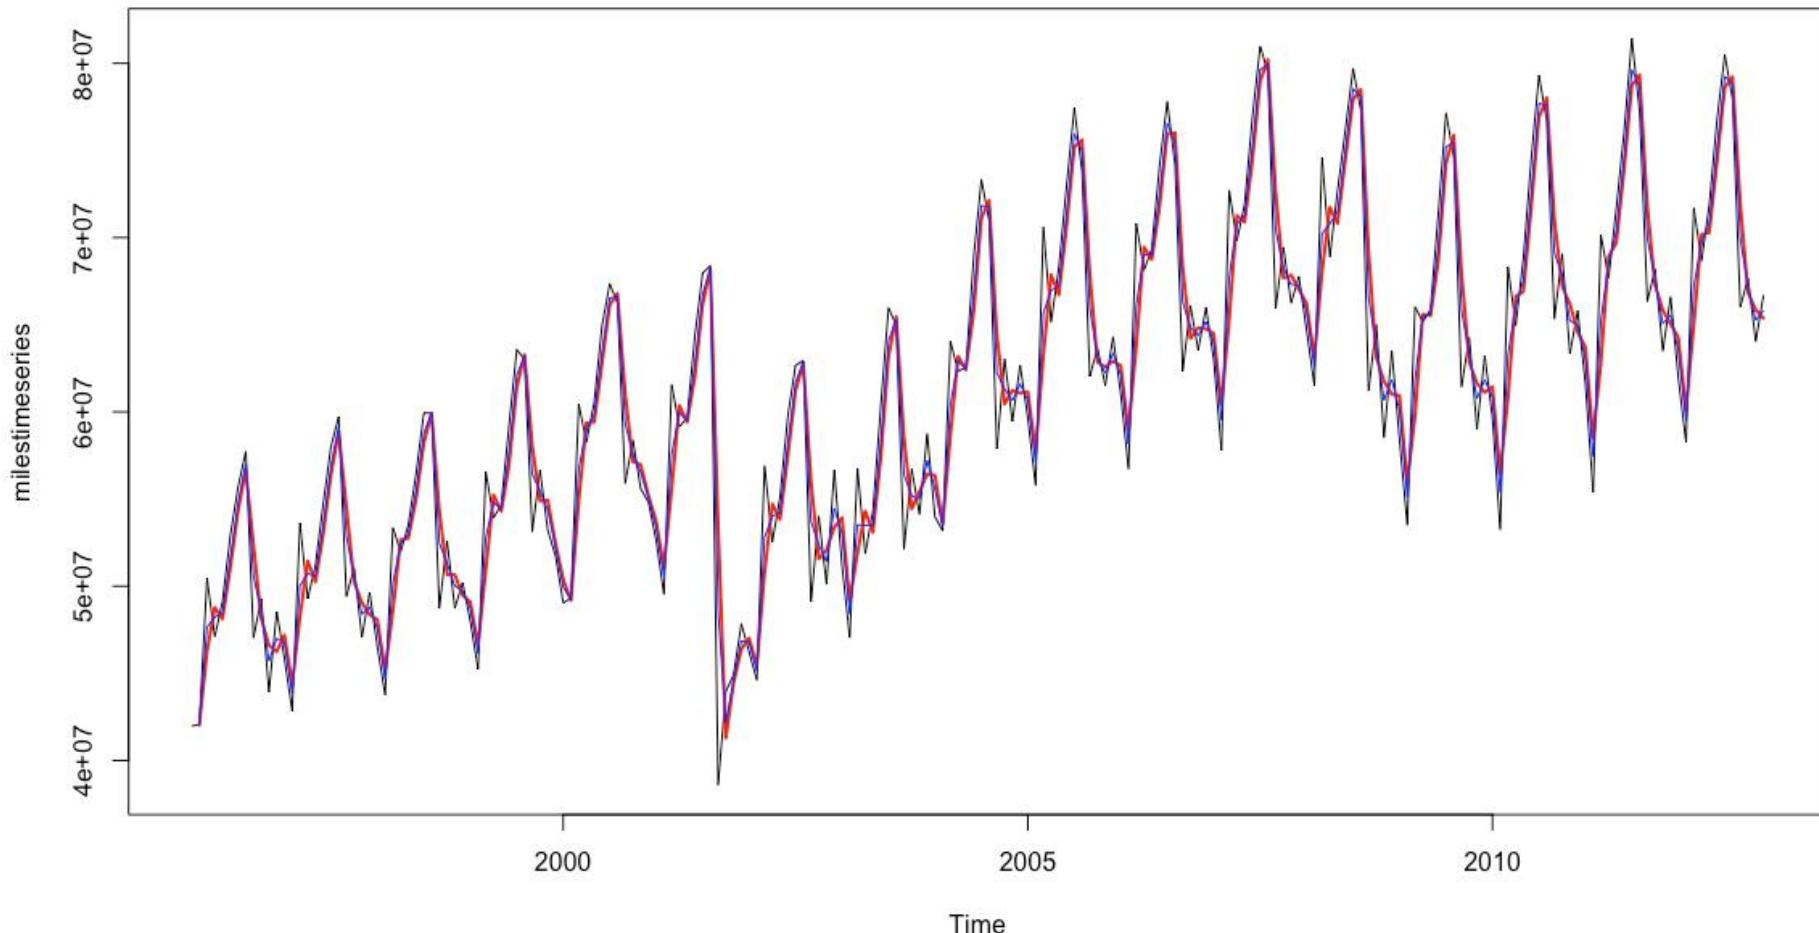




CSE 7202c



SMA and WMA – Revenue Passenger Miles



> MAPE-SMA 4.093731 > MAPE-WMA 2.729154

~~Stationary Model: Case 3 – Exponential Weighted Moving Averages or Exponential Smoothing~~

Averaging over long periods dampens fluctuations, removing not only the noise but also trend and seasonality.

Moving averages over short recent periods maintains trend and seasonality but determining an optimum number for periods is tricky, even when using metrics like MAE. If averaged over too few periods, irregularities continue to remain and if averaged over long periods, dampening again becomes a problem.

Exponential smoothing **retains all older periods** while giving a greater weight to more recent periods (hence not a MOVING average).

Caution: It doesn't make any one method superior for all situations.



Stationary Model: Case 3 – Exponential ~~Weighted Moving Averages~~ or Exponential Smoothing

$$\hat{Y}_{t+1} = \hat{Y}_t + \alpha(Y_t - \hat{Y}_t)$$

Above equation indicates that the predicted value for time period $t+1$ (\hat{Y}_{t+1}) is equal to the predicted value for the previous period (\hat{Y}_t) plus an adjustment for the error made in predicting the previous period's value ($\alpha(Y_t - \hat{Y}_t)$).

The parameter α can assume any value between 0 and 1 ($0 \leq \alpha \leq 1$).

Exponential Smoothing in Other Ways

$\hat{Y}_{t+1} = \hat{Y}_t + \alpha(Y_t - \hat{Y}_t)$ can be rewritten variously as

$$\begin{aligned}\hat{Y}_{t+1} &= \alpha Y_t + (1 - \alpha) \hat{Y}_t \\ &= Y_t - (1 - \alpha)(Y_t - \hat{Y}_t) \\ &\Rightarrow \hat{Y}_{t+1} = \alpha Y_t + \alpha(1 - \alpha)Y_{t-1} + \alpha(1 - \alpha)^2 Y_{t-2} + \dots + \alpha(1 - \alpha)^n Y_{t-n} + \dots\end{aligned}$$

Various Ways of Understanding Exponential Smoothing

- Forecast
 - Interpolation between previous *forecast* and previous *observation*
$$= \alpha Y_t + (1 - \alpha) \hat{Y}_t$$
 - Previous *forecast* plus fraction of previous error
$$= \hat{Y}_t + \alpha(Y_t - \hat{Y}_t)$$
 - Previous *observation* minus fraction 1- of previous error
$$= Y_t - (1 - \alpha)(Y_t - \hat{Y}_t)$$
 - *Exponentially weighted (i.e., discounted) moving average*
$$= \alpha Y_t + \alpha(1 - \alpha)Y_{t-1} + \alpha(1 - \alpha)^2 Y_{t-2} + \cdots + \alpha(1 - \alpha)^n Y_{t-n} + \cdots$$

Exponential Smoothing

- Y at $t+1$
$$\hat{Y}_{t+1} = \hat{Y}_t + \alpha(Y_t - \hat{Y}_t)$$
- Y at $t+2$
- All future predictions are same! This is in accordance with **stationary** assumption.

Exponential Smoothing

	A	B	C	D	E	F	G	H	I
1	Number of products sold	EMA (K=1)	Error	EMA (K=2)	Error	EMA (K=3)	Error	EMA (K=4)	Error
2	15								
3	20	=A2*1	=ABS(B3-A3)	=AVERAGE(A2:A3)					
4	16	=A3*\$K\$2+B3*\$L\$2	=ABS(B4-A4)	=A4*\$K\$3+D3*\$L\$3	=ABS(A4-D4)	=AVERAGE(A2:A4)	=ABS(A4-F4)		
5	13	=A4*\$K\$2+B4*\$L\$2	=ABS(B5-A5)	=A5*\$K\$3+D4*\$L\$3	=ABS(A5-D5)	=A5*\$K\$4+F4*\$L\$4	=ABS(A5-F5)	=AVERAGE(A2:A5)	=ABS(H5-A5)
6	18	=A5*\$K\$2+B5*\$L\$2	=ABS(B6-A6)	=A6*\$K\$3+D5*\$L\$3	=ABS(A6-D6)	=A6*\$K\$4+F5*\$L\$4	=ABS(A6-F6)	=A6*\$K\$5+H5*\$L\$5	=ABS(H6-A6)
7	14	=A6*\$K\$2+B6*\$L\$2	=ABS(B7-A7)	=A7*\$K\$3+D6*\$L\$3	=ABS(A7-D7)	=A7*\$K\$4+F6*\$L\$4	=ABS(A7-F7)	=A7*\$K\$5+H6*\$L\$5	=ABS(H7-A7)
8	15	=A7*\$K\$2+B7*\$L\$2	=ABS(B8-A8)	=A8*\$K\$3+D7*\$L\$3	=ABS(A8-D8)	=A8*\$K\$4+F7*\$L\$4	=ABS(A8-F8)	=A8*\$K\$5+H7*\$L\$5	=ABS(H8-A8)
9	17	=A8*\$K\$2+B8*\$L\$2	=ABS(B9-A9)	=A9*\$K\$3+D8*\$L\$3	=ABS(A9-D9)	=A9*\$K\$4+F8*\$L\$4	=ABS(A9-F9)	=A9*\$K\$5+H8*\$L\$5	=ABS(H9-A9)
10	20	=A9*\$K\$2+B9*\$L\$2	=ABS(B10-A10)	=A10*\$K\$3+D9*\$L\$3	=ABS(A10-D10)	=A10*\$K\$4+F9*\$L\$4	=ABS(A10-F10)	=A10*\$K\$5+H9*\$L\$5	=ABS(H10-A10)
11	20	=A10*\$K\$2+B10*\$L\$2	=ABS(B11-A11)	=A11*\$K\$3+D10*\$L\$3	=ABS(A11-D11)	=A11*\$K\$4+F10*\$L\$4	=ABS(A11-F11)	=A11*\$K\$5+H10*\$L\$5	=ABS(H11-A11)
12	18	=A11*\$K\$2+B11*\$L\$2	=ABS(B12-A12)	=A12*\$K\$3+D11*\$L\$3	=ABS(A12-D12)	=A12*\$K\$4+F11*\$L\$4	=ABS(A12-F12)	=A12*\$K\$5+H11*\$L\$5	=ABS(H12-A12)
13	20	=A12*\$K\$2+B12*\$L\$2	=ABS(B13-A13)	=A13*\$K\$3+D12*\$L\$3	=ABS(A13-D13)	=A13*\$K\$4+F12*\$L\$4	=ABS(A13-F13)	=A13*\$K\$5+H12*\$L\$5	=ABS(H13-A13)
14	10	=A13*\$K\$2+B13*\$L\$2	=ABS(B14-A14)	=A14*\$K\$3+D13*\$L\$3	=ABS(A14-D14)	=A14*\$K\$4+F13*\$L\$4	=ABS(A14-F14)	=A14*\$K\$5+H13*\$L\$5	=ABS(H14-A14)
15	17	=A14*\$K\$2+B14*\$L\$2	=ABS(B15-A15)	=A15*\$K\$3+D14*\$L\$3	=ABS(A15-D15)	=A15*\$K\$4+F14*\$L\$4	=ABS(A15-F15)	=A15*\$K\$5+H14*\$L\$5	=ABS(H15-A15)
16	17	=A15*\$K\$2+B15*\$L\$2	=ABS(B16-A16)	=A16*\$K\$3+D15*\$L\$3	=ABS(A16-D16)	=A16*\$K\$4+F15*\$L\$4	=ABS(A16-F16)	=A16*\$K\$5+H15*\$L\$5	=ABS(H16-A16)
17	18	=A16*\$K\$2+B16*\$L\$2	=ABS(B17-A17)	=A17*\$K\$3+D16*\$L\$3	=ABS(A17-D17)	=A17*\$K\$4+F16*\$L\$4	=ABS(A17-F17)	=A17*\$K\$5+H16*\$L\$5	=ABS(H17-A17)
18	18	=A17*\$K\$2+B17*\$L\$2	=ABS(B18-A18)	=A18*\$K\$3+D17*\$L\$3	=ABS(A18-D18)	=A18*\$K\$4+F17*\$L\$4	=ABS(A18-F18)	=A18*\$K\$5+H17*\$L\$5	=ABS(H18-A18)
19	14	=A18*\$K\$2+B18*\$L\$2	=ABS(B19-A19)	=A19*\$K\$3+D18*\$L\$3	=ABS(A19-D19)	=A19*\$K\$4+F18*\$L\$4	=ABS(A19-F19)	=A19*\$K\$5+H18*\$L\$5	=ABS(H19-A19)
20	14	=A19*\$K\$2+B19*\$L\$2	=ABS(B20-A20)	=A20*\$K\$3+D19*\$L\$3	=ABS(A20-D20)	=A20*\$K\$4+F19*\$L\$4	=ABS(A20-F20)	=A20*\$K\$5+H19*\$L\$5	=ABS(H20-A20)
21	14	=A20*\$K\$2+B20*\$L\$2	=ABS(B21-A21)	=A21*\$K\$3+D20*\$L\$3	=ABS(A21-D21)	=A21*\$K\$4+F20*\$L\$4	=ABS(A21-F21)	=A21*\$K\$5+H20*\$L\$5	=ABS(H21-A21)
22	19	=A21*\$K\$2+B21*\$L\$2	=ABS(B22-A22)	=A22*\$K\$3+D21*\$L\$3	=ABS(A22-D22)	=A22*\$K\$4+F21*\$L\$4	=ABS(A22-F22)	=A22*\$K\$5+H21*\$L\$5	=ABS(H22-A22)
23	12	=A22*\$K\$2+B22*\$L\$2	=ABS(B23-A23)	=A23*\$K\$3+D22*\$L\$3	=ABS(A23-D23)	=A23*\$K\$4+F22*\$L\$4	=ABS(A23-F23)	=A23*\$K\$5+H22*\$L\$5	=ABS(H23-A23)
24	13	=A23*\$K\$2+B23*\$L\$2	=ABS(B24-A24)	=A24*\$K\$3+D23*\$L\$3	=ABS(A24-D24)	=A24*\$K\$4+F23*\$L\$4	=ABS(A24-F24)	=A24*\$K\$5+H23*\$L\$5	=ABS(H24-A24)
25	14	=A24*\$K\$2+B24*\$L\$2	=ABS(B25-A25)	=A25*\$K\$3+D24*\$L\$3	=ABS(A25-D25)	=A25*\$K\$4+F24*\$L\$4	=ABS(A25-F25)	=A25*\$K\$5+H24*\$L\$5	=ABS(H25-A25)
26			=AVERAGE(C3:C25)		=AVERAGE(E3:E25)		=AVERAGE(G3:G25)		=AVERAGE(I3:I25)

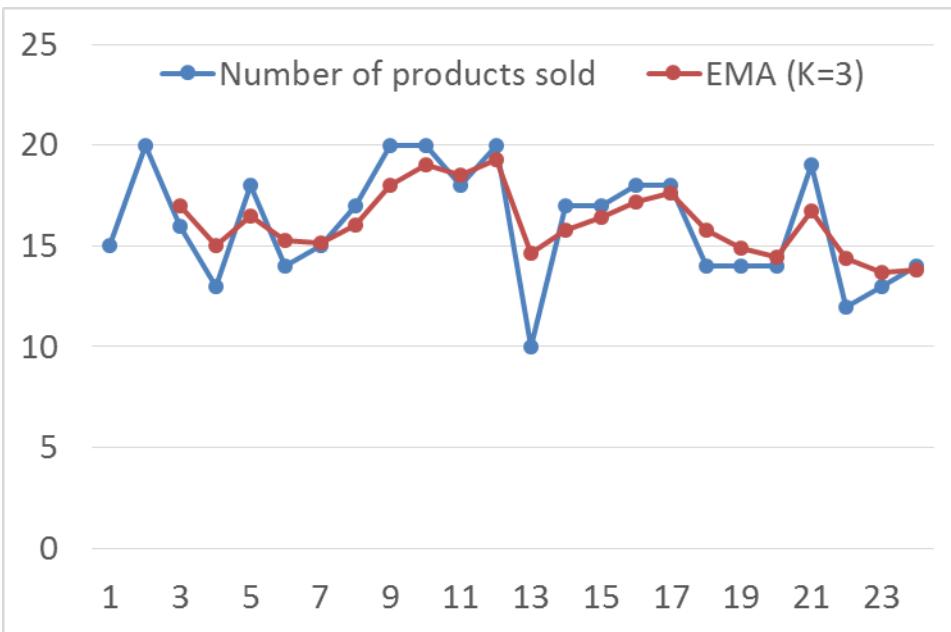
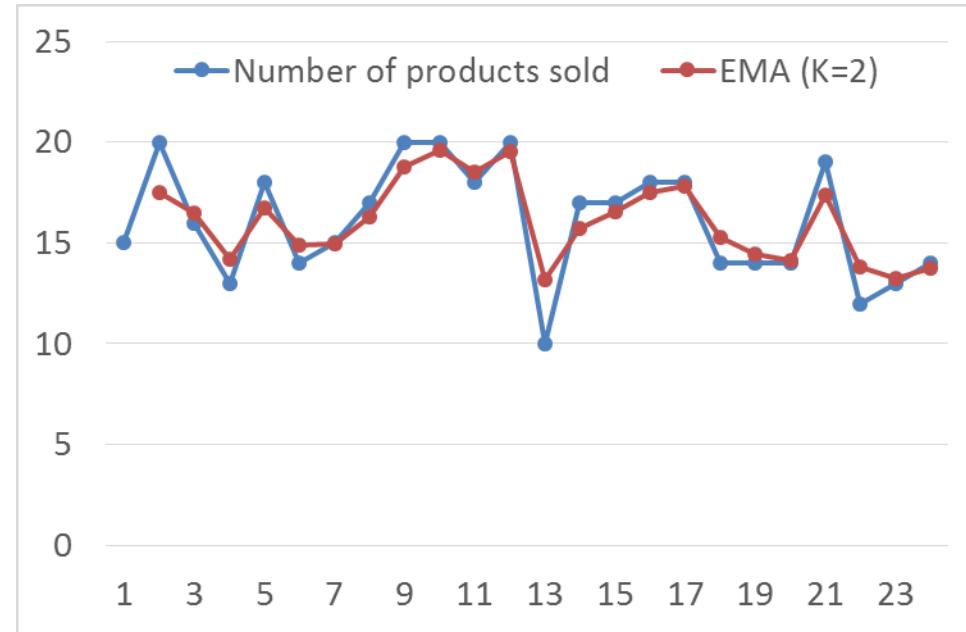
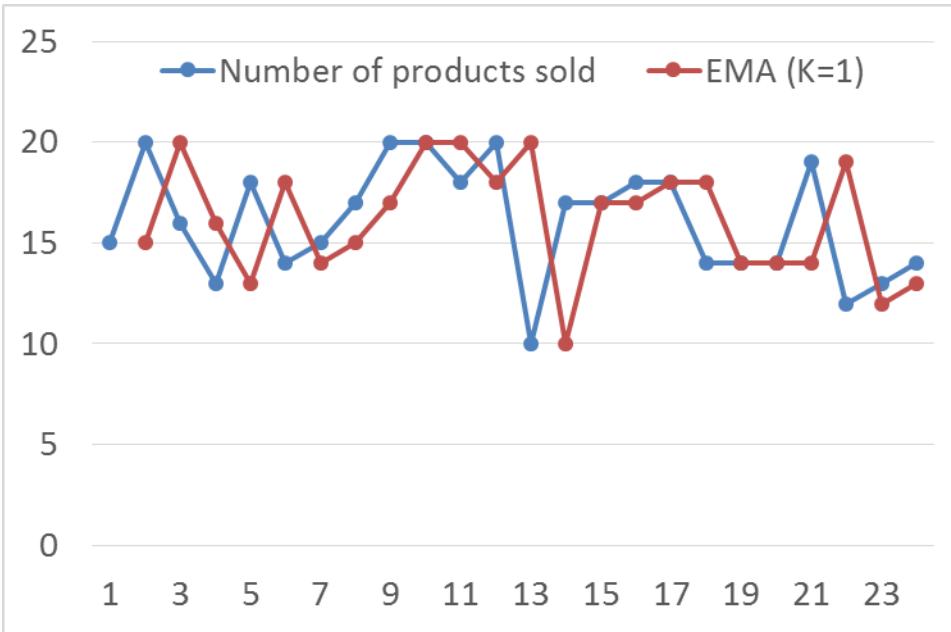
J	K	L
K	2/(K+1)	1-[2/(K+1)]
1	1	=1-K2
2	=2/(J3+1)	=1-K3
3	=2/(J4+1)	=1-K4
4	=2/(J5+1)	=1-K5

$$\hat{Y}_{t+1} = \alpha Y_t + (1 - \alpha) \hat{Y}_t$$



Exponential Smoothing

1	Number of products sold	EMA (K=1)	Error	EMA (K=2)	Error	EMA (K=3)	Error	EMA (K=4)	Error
2	15								
3	20	15	5	17.5					
4	16	20	4	16.5	0.5	17	1		
5	13	16	3	14.1666667	1.16666667	15	2	16	3
6	18	13	5	16.7222222	1.27777778	16.5	1.5	16.8	1.2
7	14	18	4	14.9074074	0.90740741	15.25	1.25	15.68	1.68
8	15	14	1	14.9691358	0.0308642	15.125	0.125	15.408	0.408
9	17	15	2	16.3230453	0.67695473	16.0625	0.9375	16.0448	0.9552
10	20	17	3	18.7743484	1.22565158	18.03125	1.96875	17.62688	2.37312
11	20	20	0	19.5914495	0.40855053	19.015625	0.984375	18.576128	1.423872
12	18	20	2	18.5304832	0.53048316	18.5078125	0.5078125	18.3456768	0.3456768
13	20	18	2	19.5101611	0.48983895	19.2539063	0.74609375	19.0074061	0.99259392
14	10	20	10	13.1700537	3.17005368	14.6269531	4.62695313	15.4044436	5.40444365
15	17	10	7	15.7233512	1.27664877	15.8134766	1.18652344	16.0426662	0.95733381
16	17	17	0	16.5744504	0.42554959	16.4067383	0.59326172	16.4255997	0.57440029
17	18	17	1	17.5248168	0.4751832	17.2033691	0.79663086	17.0553598	0.94464017
18	18	18	0	17.8416056	0.1583944	17.6016846	0.39831543	17.4332159	0.5667841
19	14	18	4	15.2805352	1.2805352	15.8008423	1.80084229	16.0599295	2.05992954
20	14	14	0	14.4268451	0.42684507	14.9004211	0.90042114	15.2359577	1.23595772
21	14	14	0	14.1422817	0.14228169	14.4502106	0.45021057	14.7415746	0.74157463
22	19	14	5	17.3807606	1.61923944	16.7251053	2.27489471	16.4449448	2.55505522
23	12	19	7	13.7935869	1.79358685	14.3625526	2.36255264	14.6669669	2.66696687
24	13	12	1	13.264529	0.26452895	13.6812763	0.68127632	14.0001801	1.00018012
25	14	13	1	13.754843	0.24515702	13.8406382	0.15936184	14.0001081	0.00010807
26			2.91304348		0.84055449		1.23867161		1.48027795

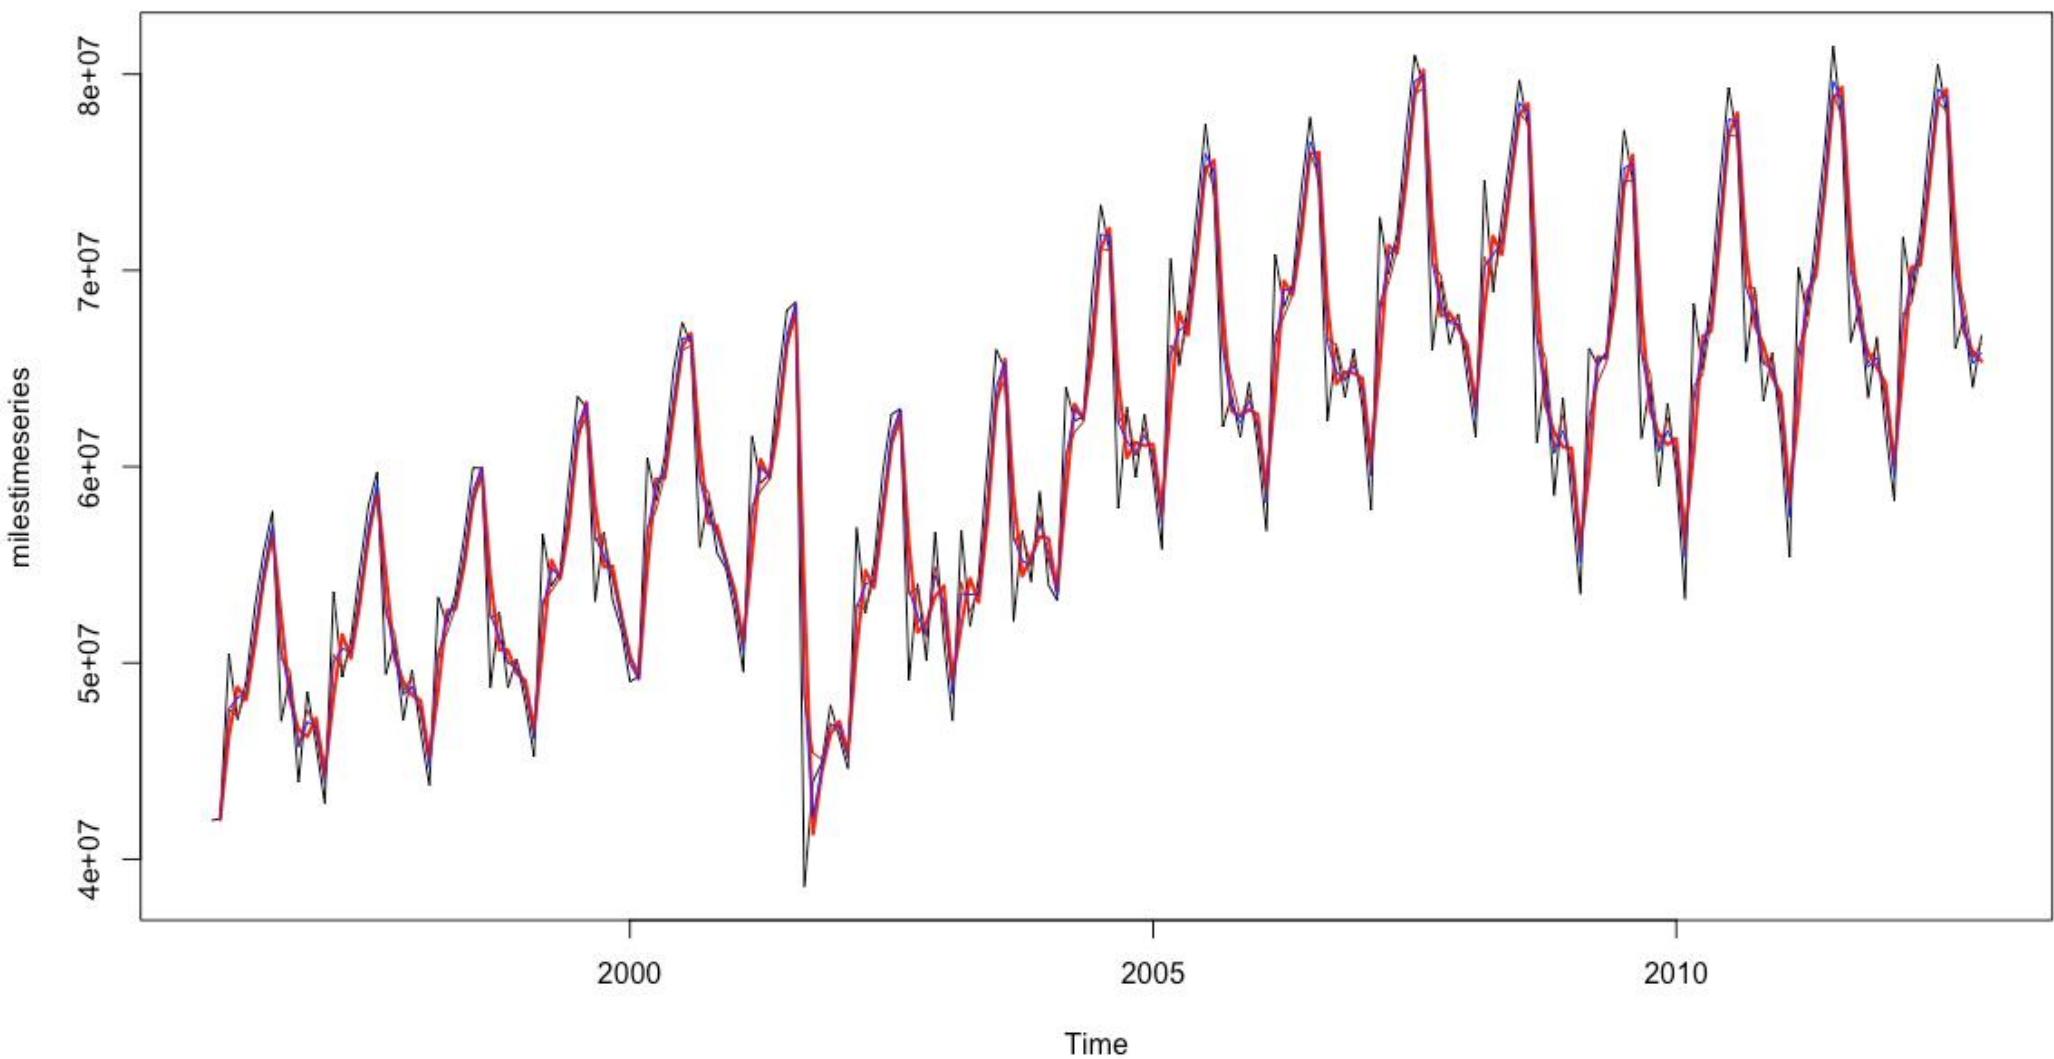




CSE 7202C



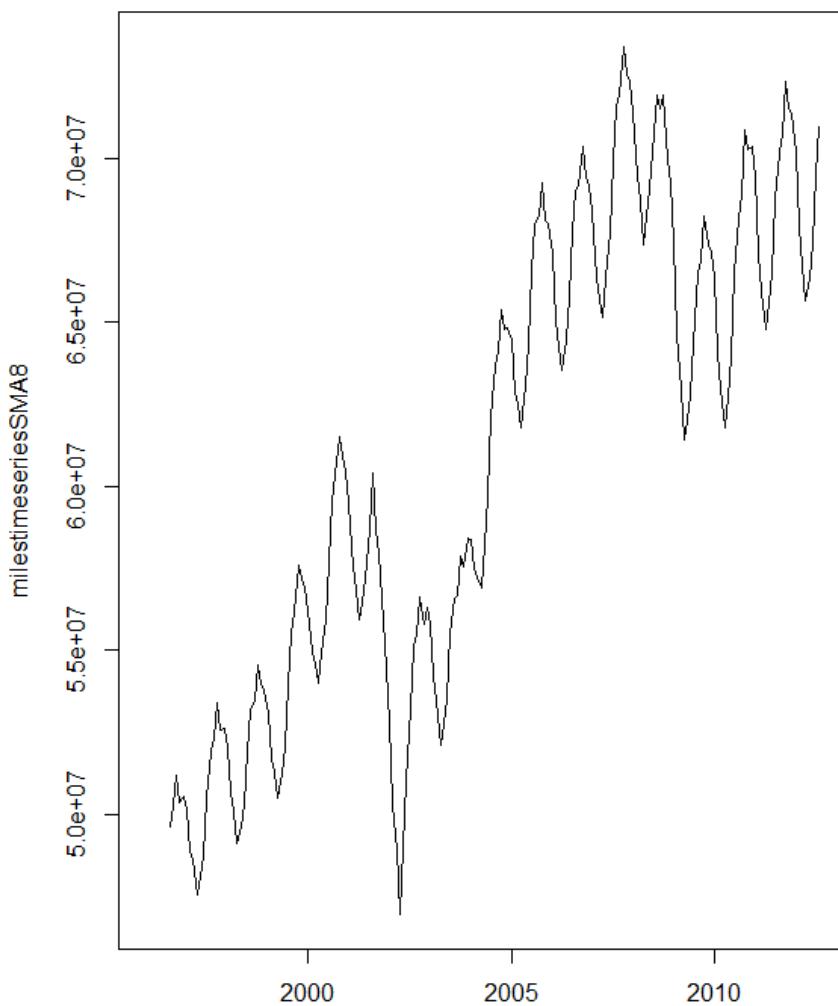
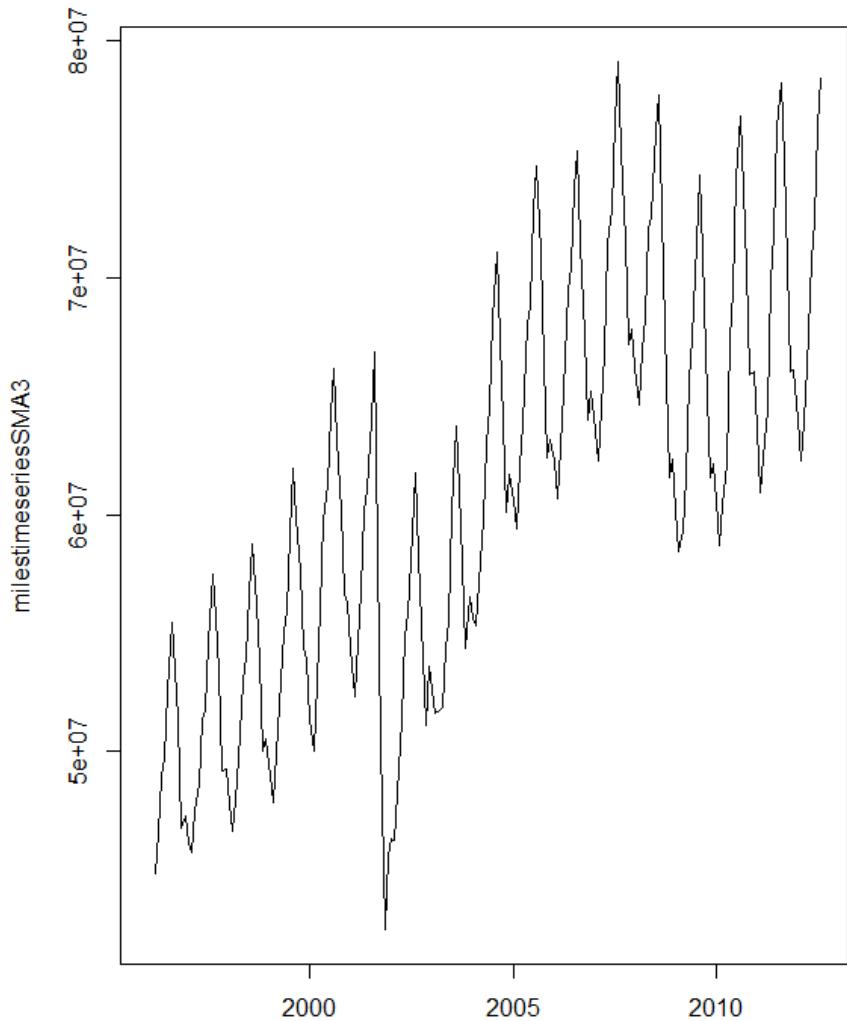
SMA, WMA and Exponential Smoothing – RPM



> MAPE-SMA 4.087519 > MAPE-WMA 2.725013 > MAPE-EMA 2.542173



Effect of k – Revenue Passenger Miles



SMA

K=3 vs K=8

Advanced Time Series Methods

HOLT-WINTERS

CSE 7202c



Holt-Winters Method

- 3 components – Trend, Seasonality and Random/Level/Irregular.
- 3 weights – smoothing parameters – are used to update components at each period.

Holt-Winters Method

Additive Seasonality

$$\hat{y}_t = l_{t-1} + b_{t-1} + s_{t-m}$$

$$\hat{y}_{t+h|t} = l_t + hb_t + s_{t+h-m}$$

Multiplicative Seasonality

$$\hat{y}_t = (l_{t-1} + b_{t-1})s_{t-m}$$

$$\hat{y}_{t+h|t} = (l_t + hb_t)s_{t+h-m}$$

The 3 smoothing equations are:

$$l_t = \alpha(y_t - s_{t-m}) + (1 - \alpha)(l_{t-1} + b_{t-1})$$

$$b_t = \beta(l_t - l_{t-1}) + (1 - \beta)b_{t-1}$$

$$s_t = \gamma(y_t - (l_{t-1} + b_{t-1})) + (1 - \gamma)s_{t-m}$$

$$l_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(l_{t-1} + b_{t-1})$$

$$b_t = \beta(l_t - l_{t-1}) + (1 - \beta)b_{t-1}$$

$$s_t = \gamma \frac{y_t}{l_{t-1} + b_{t-1}} + (1 - \gamma)s_{t-m}$$

Recall Exponential Smoothing form: $\hat{Y}_t = \alpha Y_{t-1} + (1 - \alpha)\hat{Y}_{t-1}$

Holt-Winters Method

- Initial values for level and trend components are obtained using:
 - Average of the first few observations. Need to decide the number of observations.
 - The first observation. Works well for long series.
 - Average of the first year (period) of data for *level* and average of the slopes for each period in the first 2 years for *slope*. First forecast for time period $m+1$ and not for period 1.
 - $$l_m = (y_1 + \dots + y_m)/m$$
 - $$b_m = \frac{\frac{(y_{m+1}-y_1)}{m} + \frac{(y_{m+2}-y_2)}{m} + \dots + \frac{(y_{m+m}-y_m)}{m}}{m}$$
 - $$= [(y_{m+1} + y_{m+2} + \dots + y_{m+m}) - ((y_1 + y_2 + \dots + y_m))] / m^2$$
 - $s_i = y_i - l_m$ for additive seasonality
 - $s_i = \frac{y_i}{l_m}$ for multiplicative seasonality

Holt-Winters Method

- **Initial values** for level and trend components are obtained using:
 - Linear regression on time using first few years (periods) of data.
 - Initial *level* l_0 is set to the intercept.
 - Initial *slope* b_0 is set to the regression slope.
 - Decomposition method (*implemented in R's HoltWinters() function and described next*).

Holt-Winters Method

- **Initial values** for seasonal component are obtained from:
 - A dummy variable regression using de-trended data.
 - Decomposition method (*implemented in R's HoltWinters() function and described next*).
- In the level equation, the series is seasonally adjusted by subtracting the seasonal component.

Holt-Winters Implementation in R's HoltWinters() function

STEP 1: Fit a $2 \times m$ moving average smoother to the first 2 years (periods) of data. Moving averages remove some randomness and leave a smooth trend component at that point.

Note: For shorter time series, a simple linear model with time trend is used in place of the MA smoothing. First order Fourier approximation is used for the seasonal component. The rest of the steps remain same as follows.

U.S. Air Carrier Traffic Statistics	12-MA	2X12-MA Trend
41972194		
42054796		
50443045		
47112397		
49118248		
52880510	48812034	48973636
55664750	49135237	49167910
57723208	49200583	49332997
47035464	49465412	49555846
49263120	49646280	49732680
43937074	49819080	49895193
48539606	49971306	50068400
45850623	50165495	50248505
42838949	50331514	50430794
53620994	50530075	50604898
49282817	50679721	

Holt-Winters Implementation in R's HoltWinters() function

STEP 2: Get de-trended data by subtracting (for additive Holt-Winters) or dividing (for multiplicative Holt-Winters) the smooth trend from the original data. The initial seasonal values are then obtained by averaging the de-trended values from the same season.

U.S. Air Carrier Traffic Statistics	12-MA	2X12-MA Trend	Detrended Series	Seasonal Component
41972194				
42054796				
50443045				
47112397				
49118248				
52880510	48812034			
55664750	49135237	48973636	6691114	10086225
57723208	49200583	49167910	8555298	8613219
47035464	49465412	49332997	-2297533	-4652131
49263120	49646280	49555846	-292726	-1338310
43937074	49819080	49732680	-5795606	-5113908
48539606	49971306	49895193	-1355587	-2406656
45850623	50165495	50068400	-4217777	-5951588
42838949	50331514	50248505	-7409556	-9546987
53620994	50530075	50430794	3190200	2875281
49282817	50679721	50604898	-1322081	-364949

Average of all July de-trended values

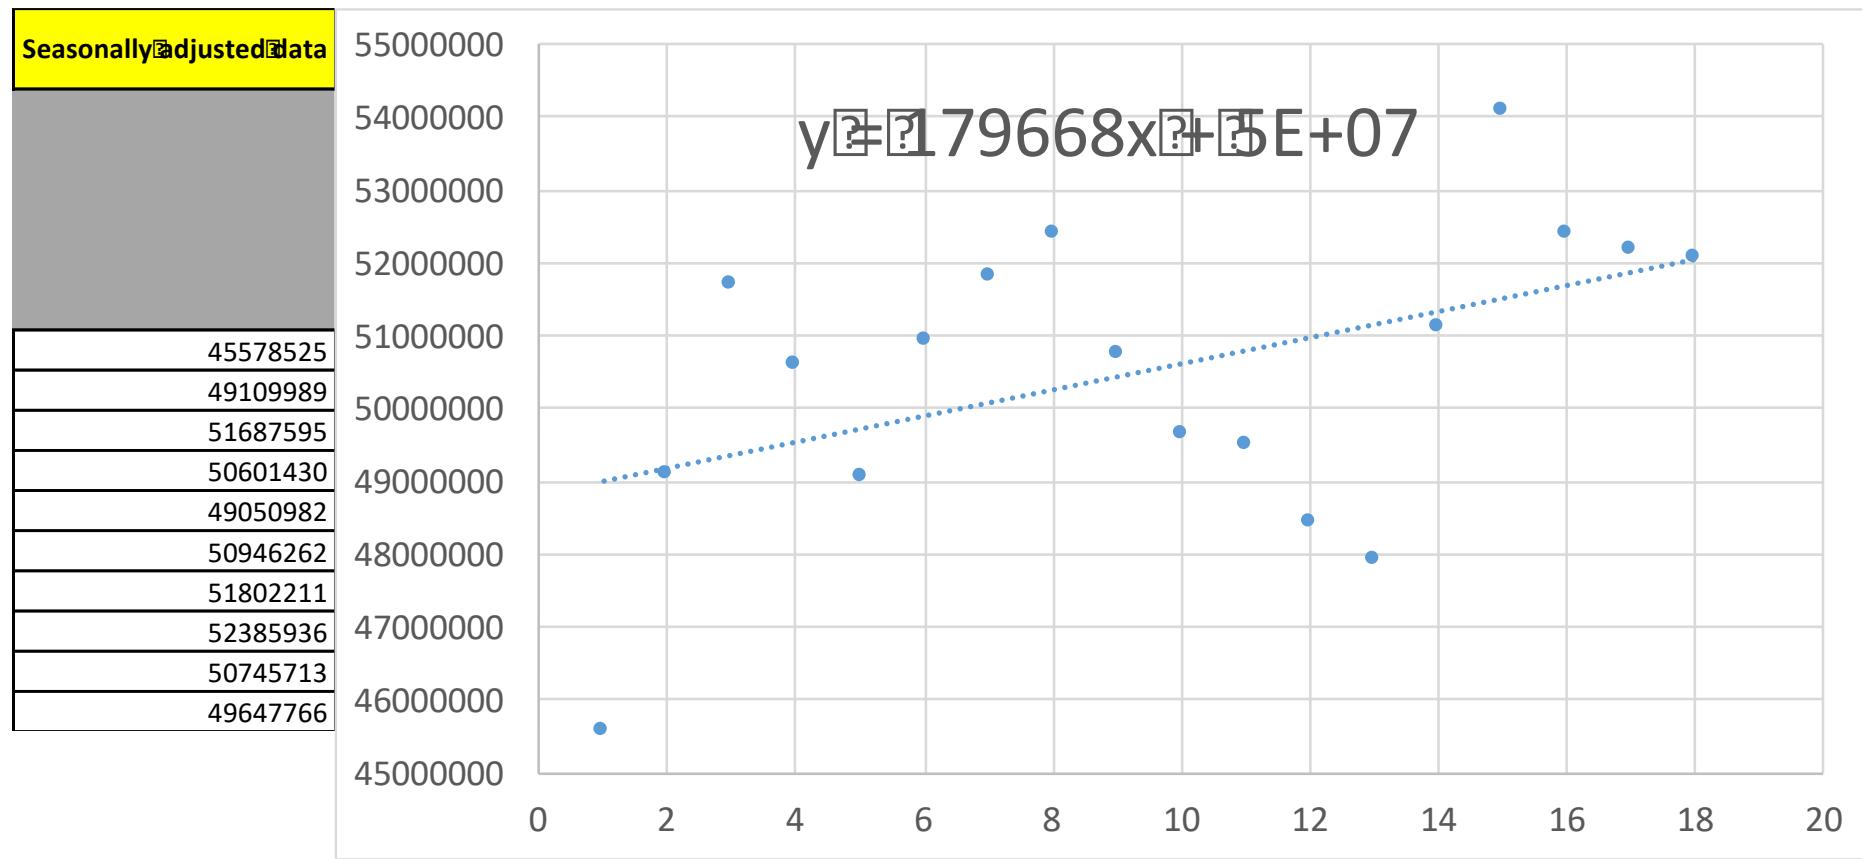
Holt-Winters Implementation in R's HoltWinters() function

STEP 3: Get seasonally adjusted data by subtracting (for additive Holt-Winters) or dividing (for multiplicative Holt-Winters) the seasonal component from the original data.

U.S. Air Carrier Traffic Statistics	12-MA	2X12-MA Trend	Detrended Series	Seasonal Component	Seasonally Adjusted Data
41972194					
42054796					
50443045					
47112397					
49118248					
52880510	48812034				
55664750	49135237	48973636	6691114	10086225	45578525
57723208	49200583	49167910	8555298	8613219	49109989
47035464	49465412	49332997	-2297533	-4652131	51687595
49263120	49646280	49555846	-292726	-1338310	50601430
43937074	49819080	49732680	-5795606	-5113908	49050982
48539606	49971306	49895193	-1355587	-2406656	50946262
45850623	50165495	50068400	-4217777	-5951588	51802211
42838949	50331514	50248505	-7409556	-9546987	52385936
53620994	50530075	50430794	3190200	2875281	50745713
49282817	50679721	50604898	-1322081	-364949	49647766

Holt-Winters Implementation in R's HoltWinters() function

STEP 4: Fit a linear trend to the seasonally adjusted data to get initial values for the level l_0 (the intercept) and the initial slope b_0 .



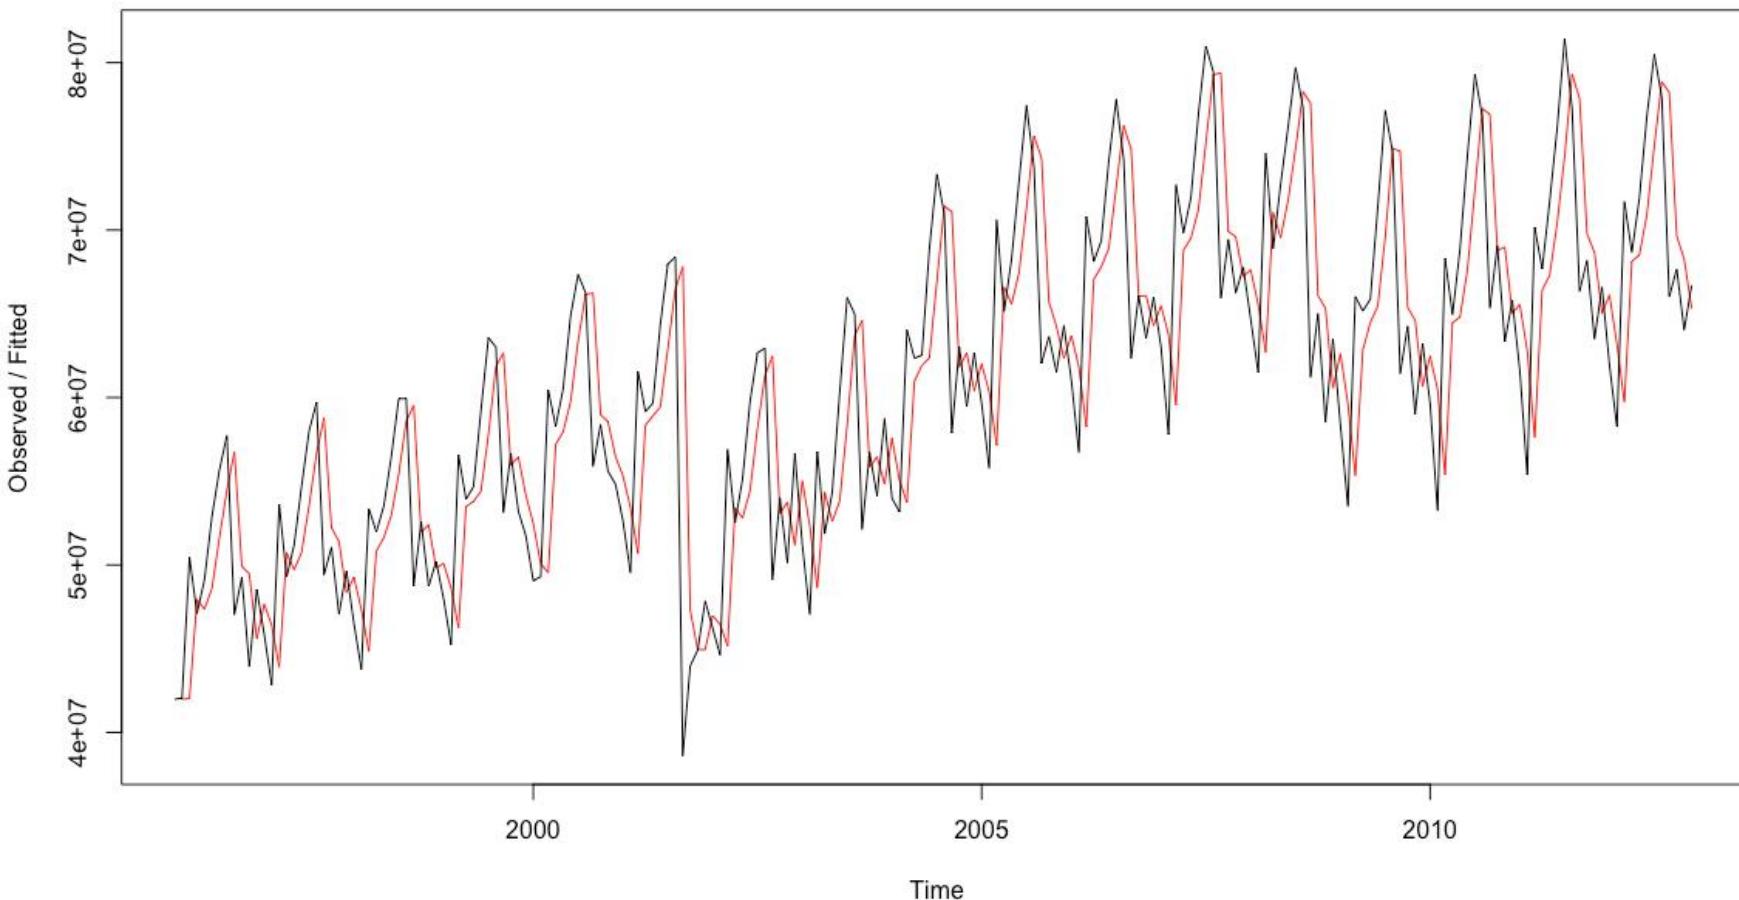


CSE 7202c



Holt-Winters Method: Only Randomness

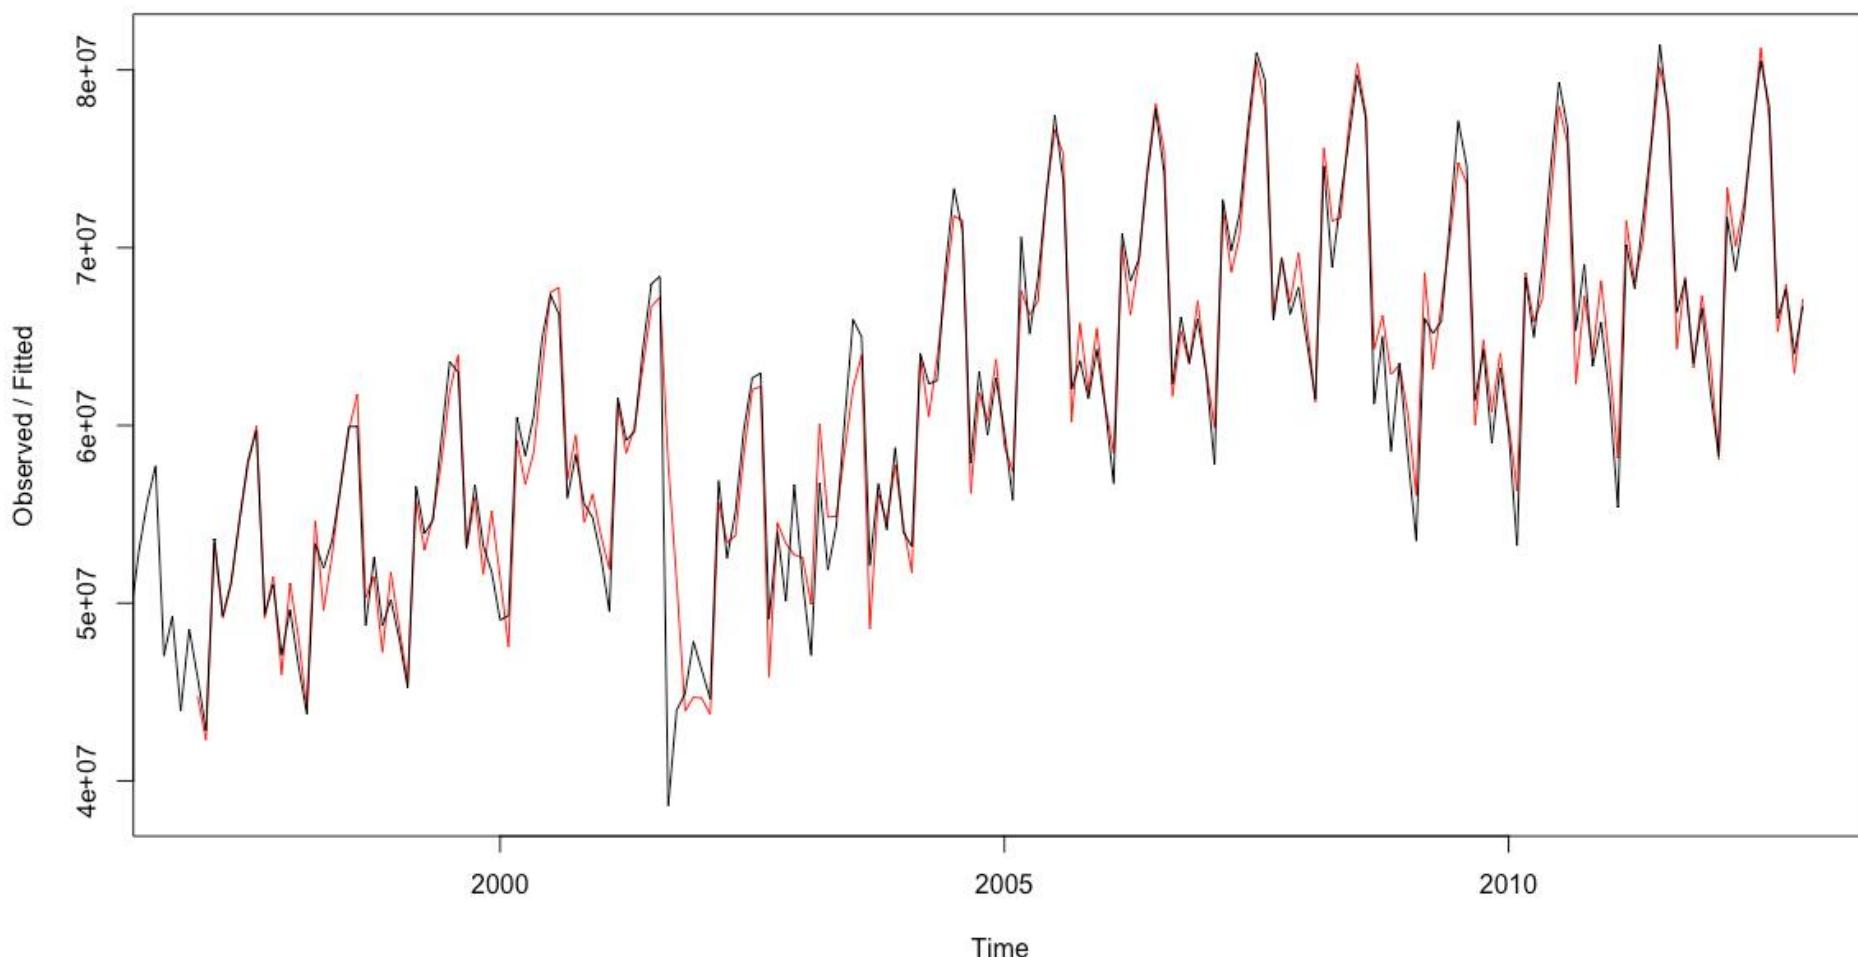
Holt-Winters filtering



```
milesforecast$SSE  
[1] 7.485495e+15
```

Holt-Winters Method: All Components

Holt-Winters filtering



milesforecast\$SSE
[1] 8.5203e+14

CSE 72026



Holt-Winters Method: All Components

```
> milesforecast
```

Holt-Winters exponential smoothing with trend and additive seasonal component.

Call:

```
HoltWinters(x = milestimeseries)
```

Smoothing parameters:

```
alpha: 0.5024519  
beta : 0  
gamma: 0.6698853
```

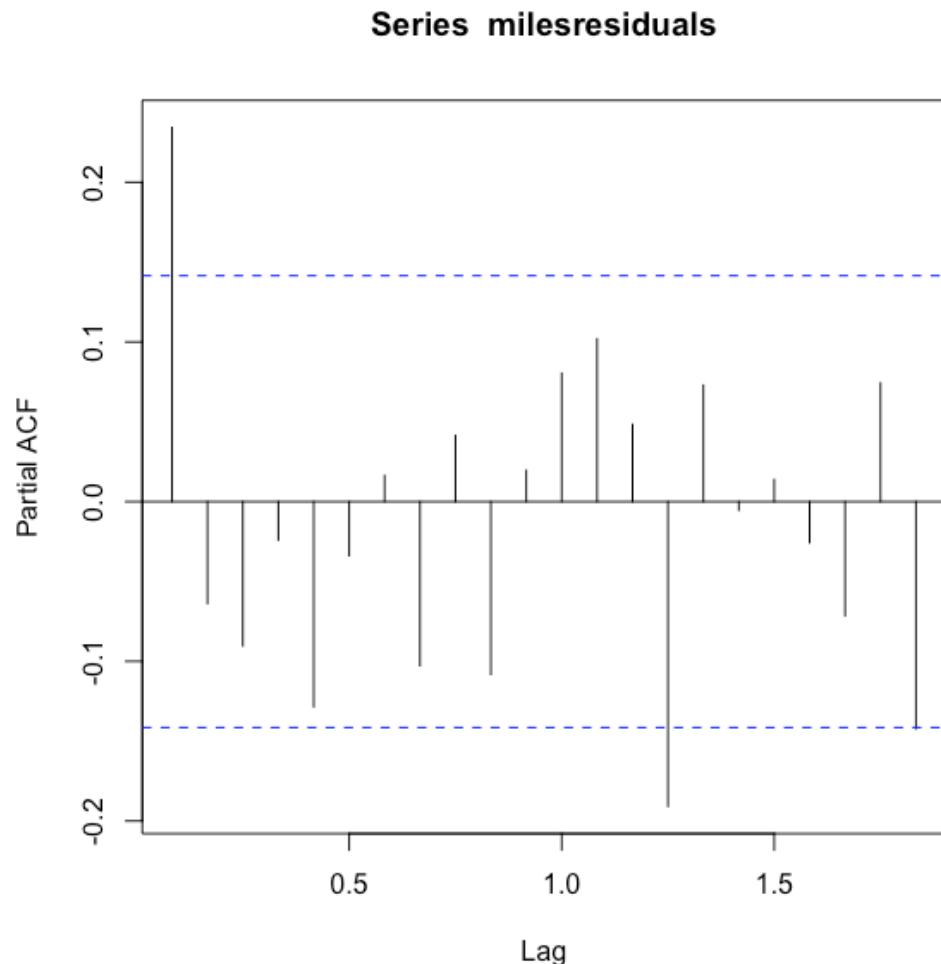
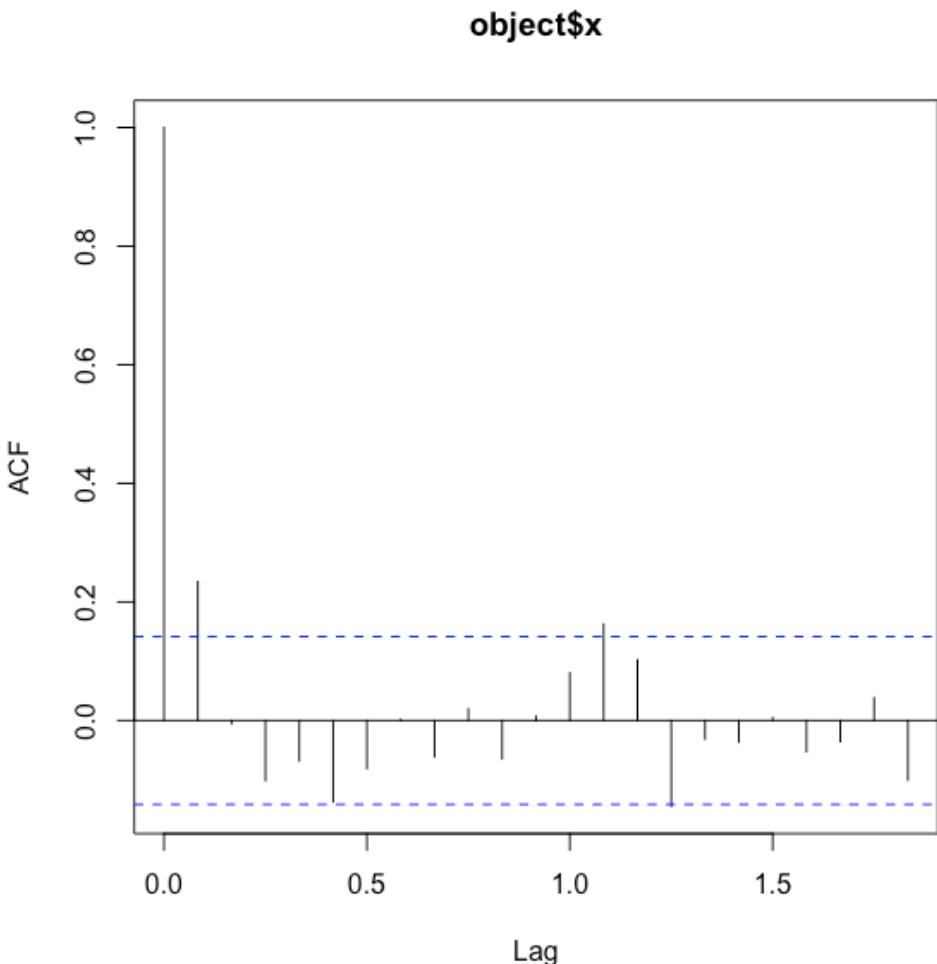
Coefficients:

```
 [,1]  
a    71153143.7  
b     181508.7  
s1   -8052713.5  
s2   -12314754.8  
s3    2106026.2  
s4   -447843.5  
s5    2758176.7  
s6    7232862.9  
s7   11309888.8  
s8    8022549.6  
s9   -4490660.8  
s10  -2715650.2  
s11  -7315682.9  
s12  -4384446.7
```

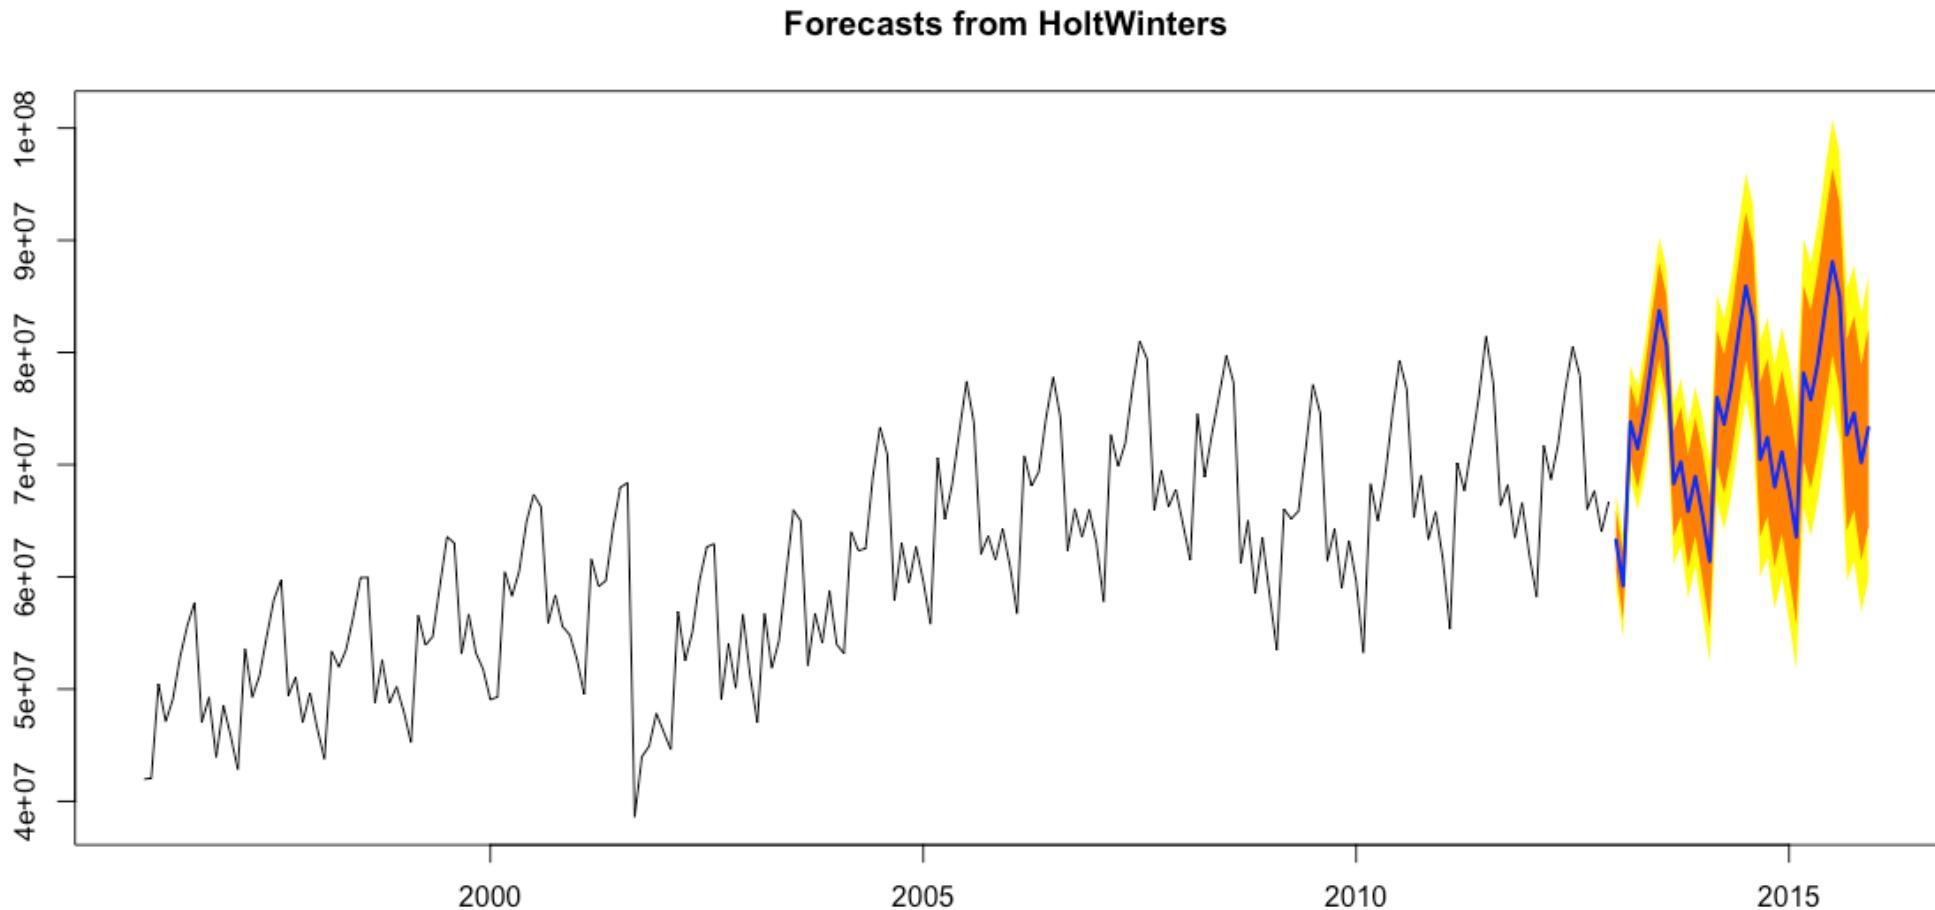
```
> milesforecast$fitted
```

	xhat	level	trend	season
Jan 1997	44780388	48804078	181508.7	-4205198.60
Feb 1997	42307860	49523328	181508.7	-7396976.76

Holt-Winters Method: Residuals



Holt-Winters Method: Forecasting



CSE 7202c



Advanced Time Series Methods

ARIMA



AutoRegressive (AR) Models

- \hat{y}_t depends only on its own past values $y_{t-1}, y_{t-2}, y_{t-3}$, etc. Thus,

$$\hat{y}_t = f(y_{t-1}, y_{t-2}, y_{t-3}, \dots, \varepsilon_t)$$

- Autoregressive model where it depends on **p** of its **past values**

(lagged values having significant relationship with the most recent value) is an **AR(p)** model and is represented as:

$$\hat{y}_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + e_t$$

Moving Average (MA) Models

- A moving average model is one where \hat{y}_t **depends** only on the **random error term**, i.e.,

$$\hat{y}_t = f(\varepsilon_t, \varepsilon_{t-1}, \varepsilon_{t-2}, \varepsilon_{t-3}, \dots)$$

- Moving average model where it depends on **q** of its **past values (residuals or lagged errors from earlier estimates)** is an **MA(q)** model and is represented as

$$\hat{y}_t = c + e_t + \theta_1 e_{t-1} + \theta_2 e_{t-2} + \dots + \theta_q e_{t-q}$$

AutoRegressive Moving Average (ARMA) Models

- ARMA(p,q)
 - Mix of both AR and MA
- General form of such a time-series model

$$\hat{y}_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + e_t + \theta_1 e_{t-1} + \theta_2 e_{t-2} + \dots + \theta_q e_{t-q}$$

Integration of order d

- Non-stationary series can be made stationary series after differencing.
- Integration of order d
 - A series that is differentiated d times
 - Denoted as $I(d)$
 - If the series is differentiated once then it is $I(1)$.
 - $I(0)$ means series is stationary without differencing

Box-Jenkins Methodology

- Model identification and model selection
 - Make sure variables are stationary. Difference as necessary to get a constant mean and transformations to get constant variance.
 - Check for seasonality: Decays and spikes at regular intervals in ACF plots.
- Parameter estimation
 - Compute coefficients that best fit the selected model.
- Model checking
 - Check if residuals are independent of each other and constant in mean and variance over time (white noise).

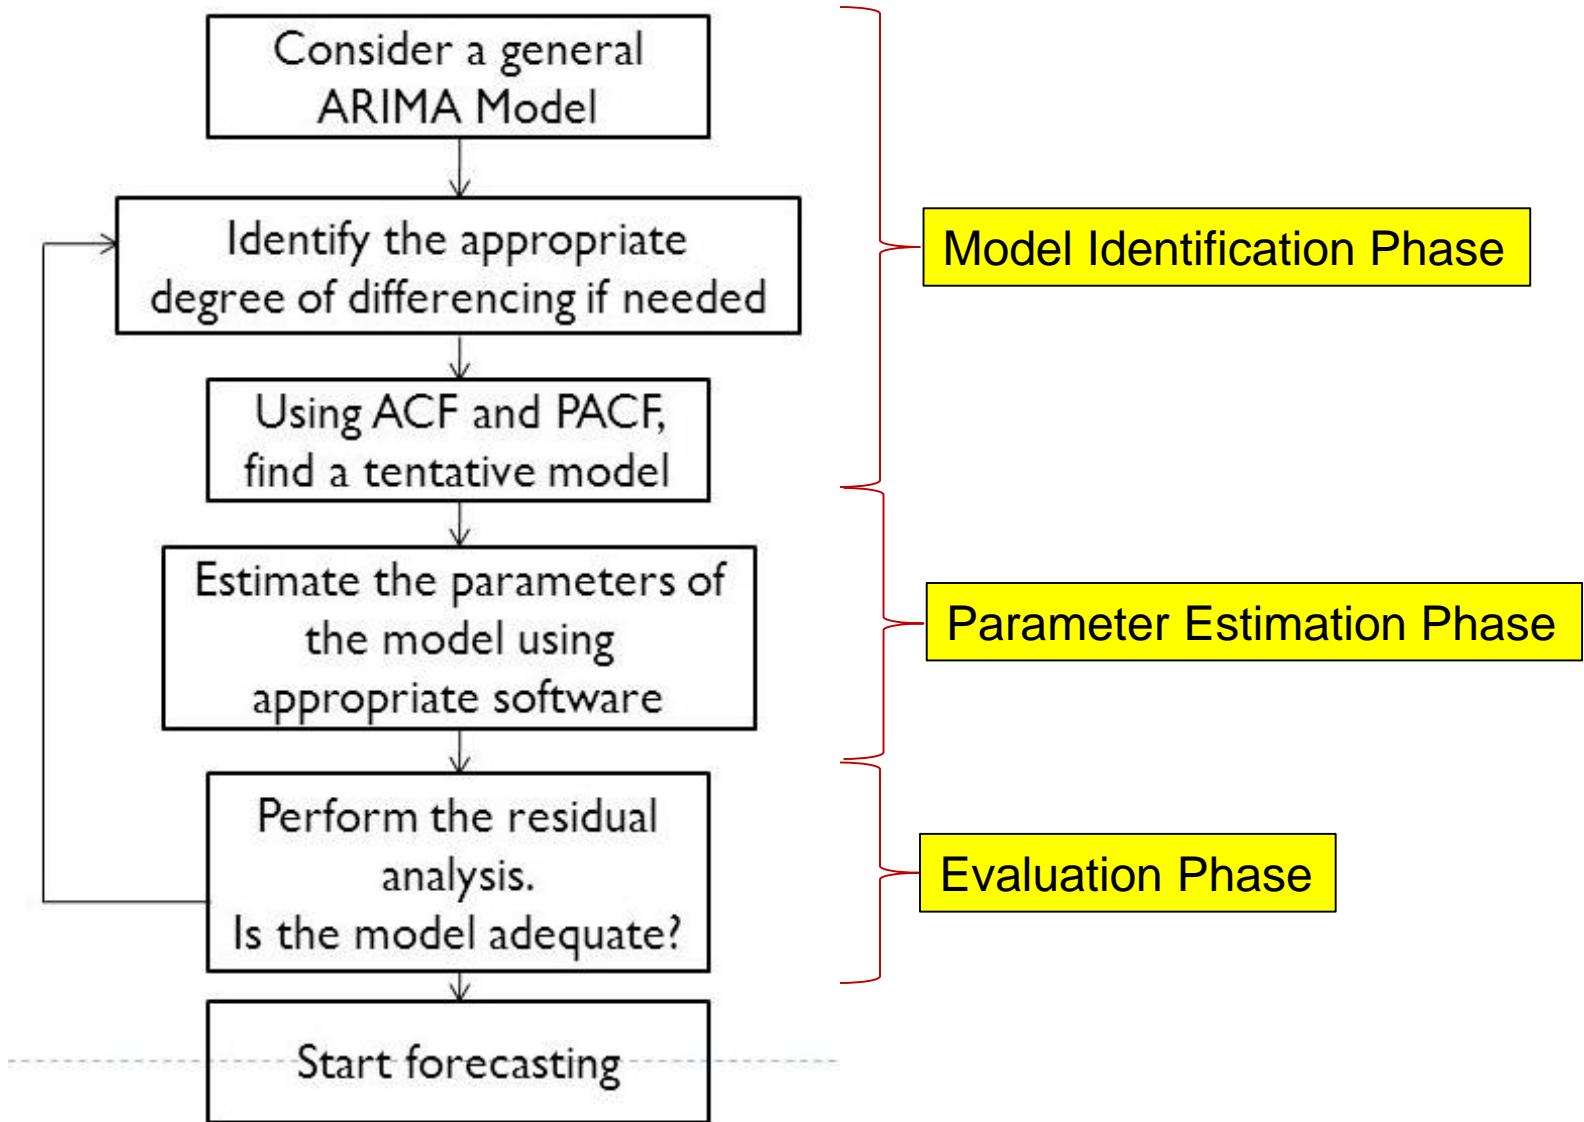
http://www.ncss.com/wp-content/themes/ncss/pdf/Procedures/NCSS/The_Box-Jenkins_Method.pdf

ARIMA(p,d,q) Model

- p is the number of autoregressive [AR(p)] terms (a linear regression of the current value of the series against one or more prior values of the series)
 - Maximum lag beyond which PACF is 0
- d is the number of non-seasonal differences (order of the differencing) used to make the time series stationary [I(d)]
- q is the order of the moving average [MA(q)] model
 - Maximum lag beyond which the ACF is 0

- Non-seasonal ARIMA models are denoted ARIMA(p,d,q)
- Seasonal ARIMA (SARIMA) models are denoted ARIMA(p,d,q)(P,D,Q)_m, where m refers to the number of periods in each season and (P,D,Q) refer to the autoregressive, differencing and moving average terms of the seasonal part of the ARIMA model.

Time Series Model Building Using ARIMA





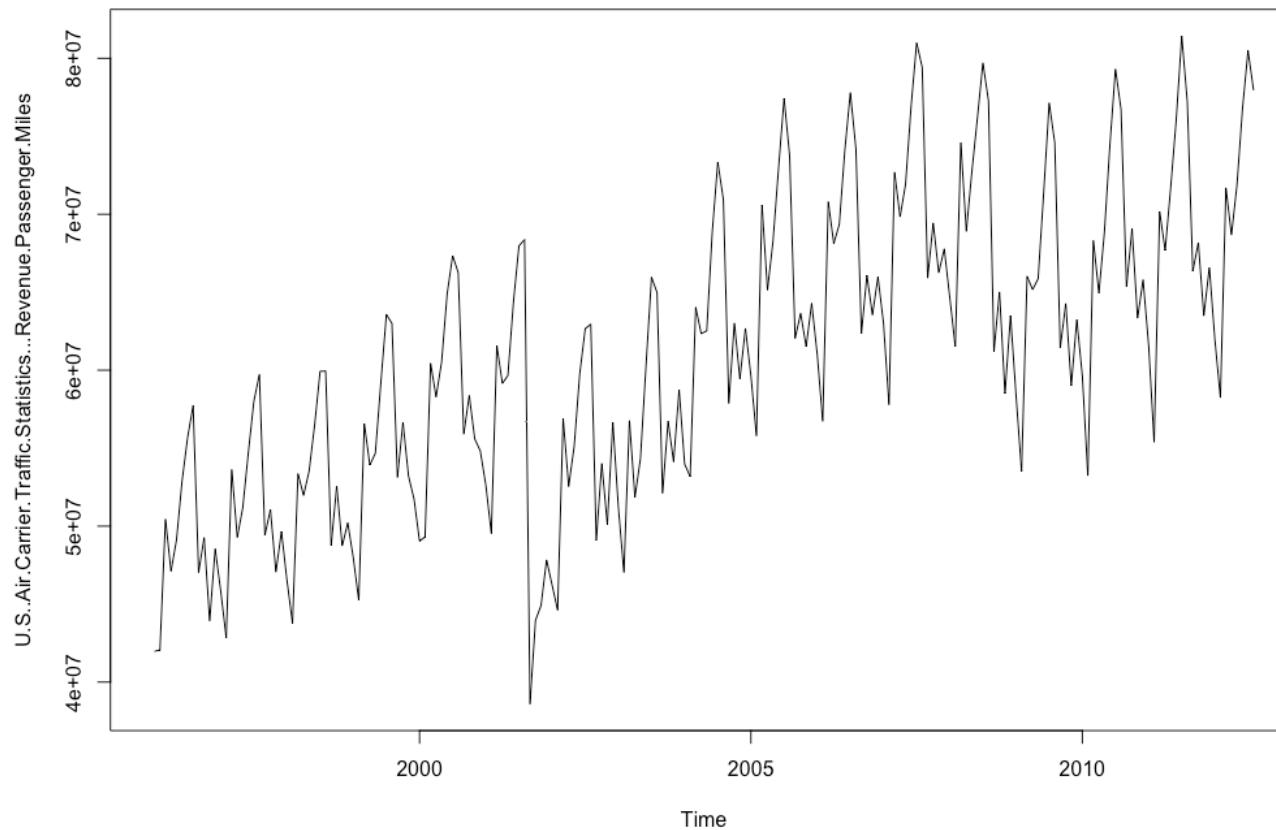
CSE 7202C



Time Series Model Building Using ARIMA

Identification Phase

Step 1: Plot the data (transform data to stabilize variance, if required)



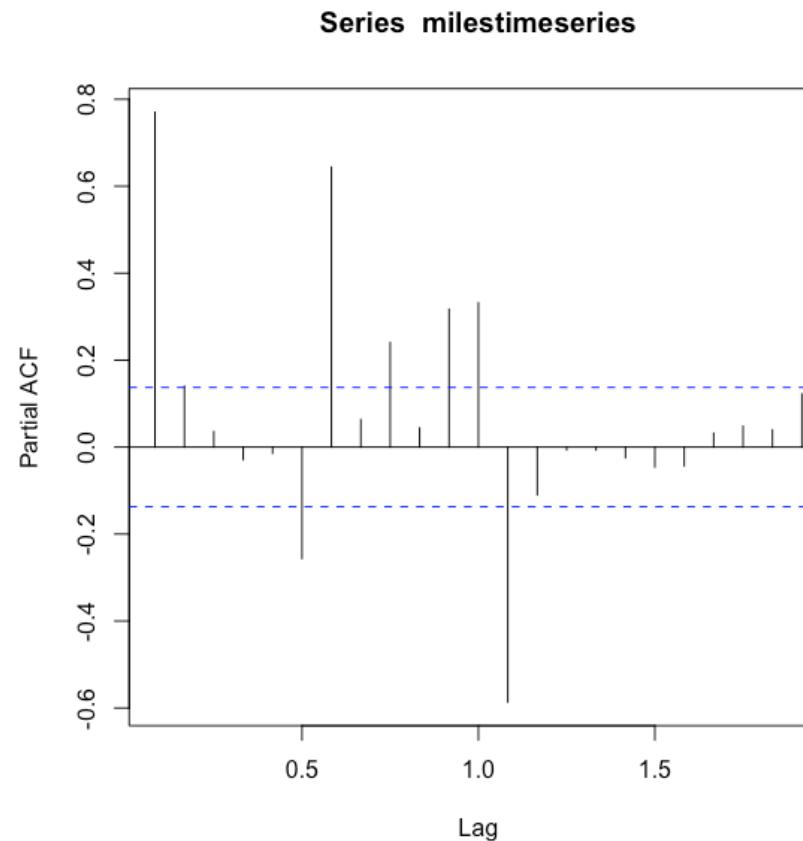
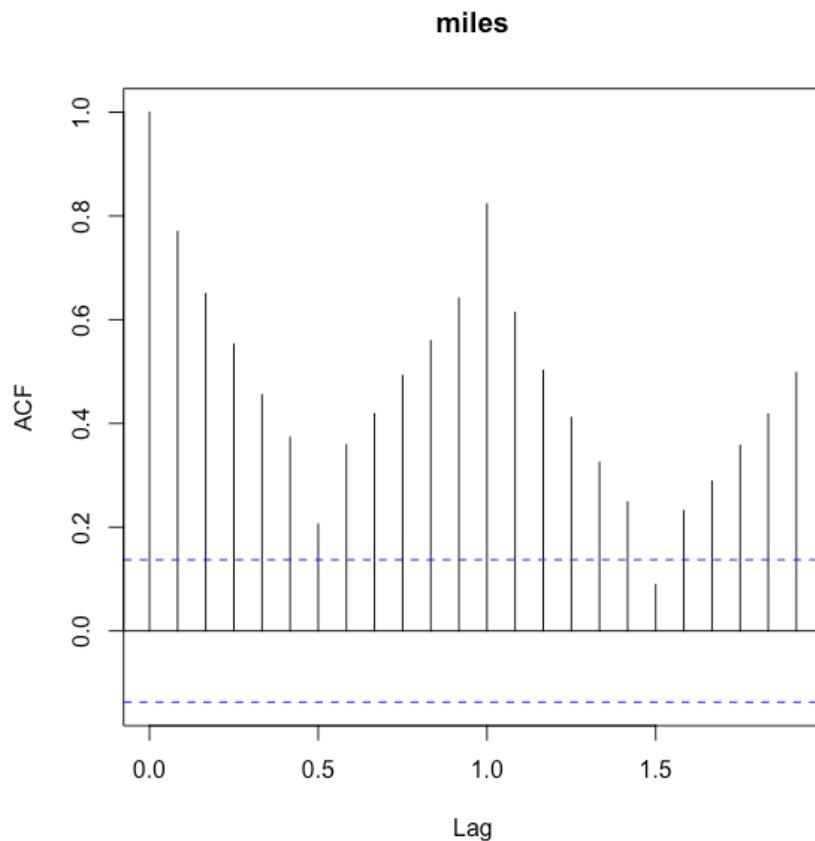
CSE 72026



Time Series Model Building Using ARIMA

Identification Phase

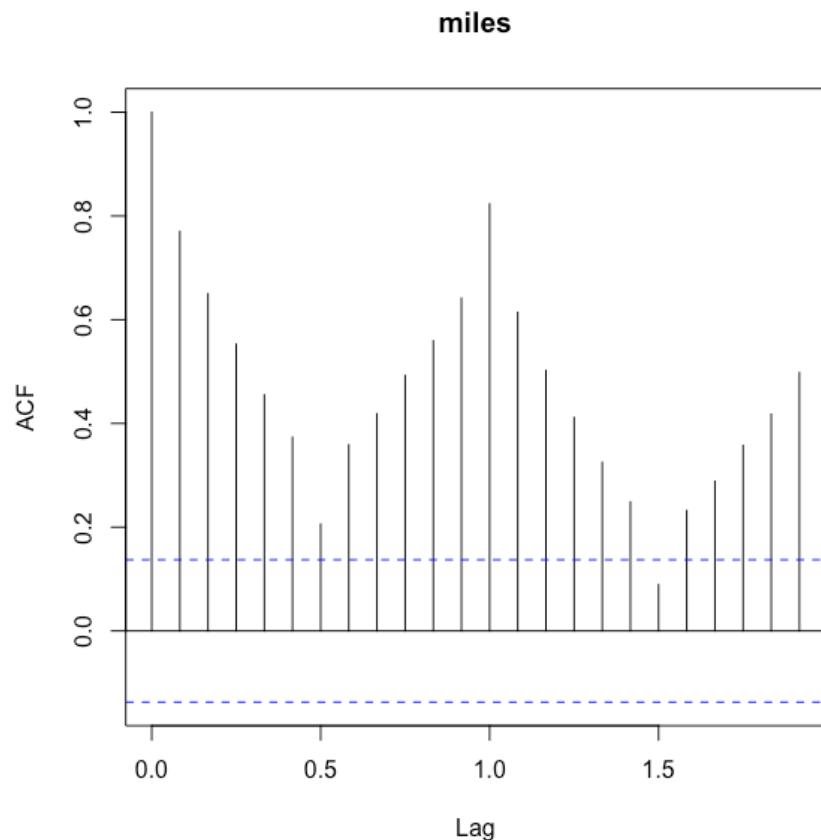
Step 2: Plot ACF and PACF to get preliminary understanding of the processes involved.



Time Series Model Building Using ARIMA

Identification Phase

Step 2: Plot ACF and PACF to get preliminary understanding of the processes involved.

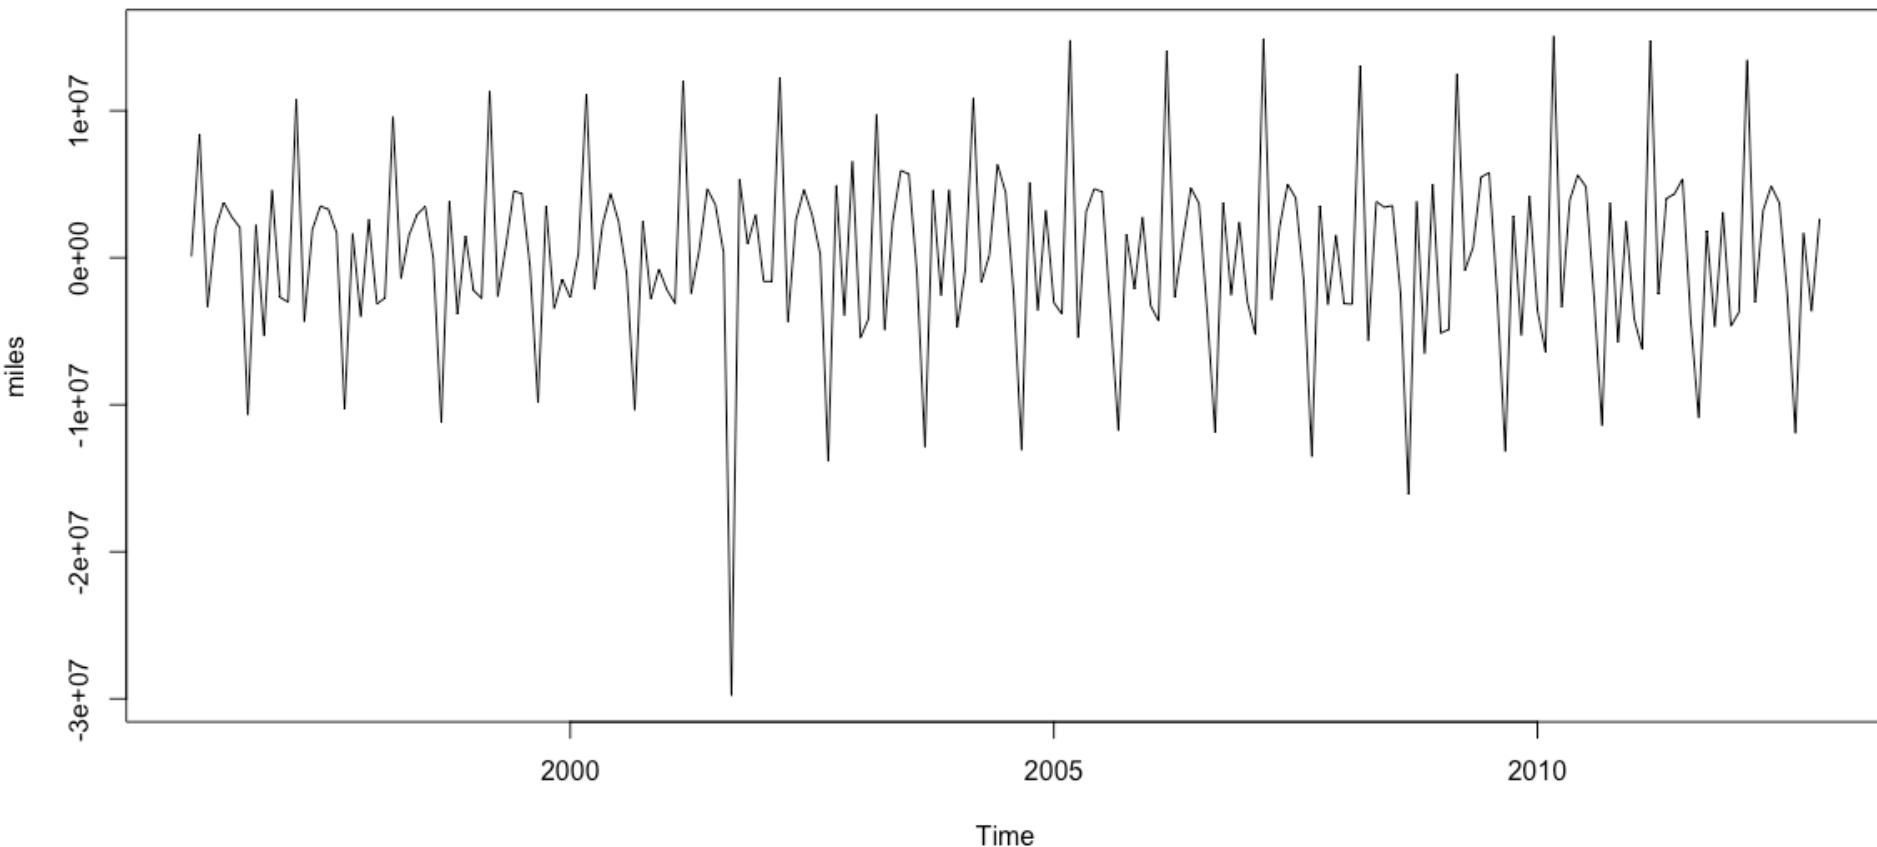


The suspension bridge pattern in ACF (also, positive and negative spikes in PACF) suggests non-stationarity and strong seasonality.

Time Series Model Building Using ARIMA

Identification Phase

Step 3: Perform a non-seasonal difference. It is the same as an ARIMA(0,1,0) model.



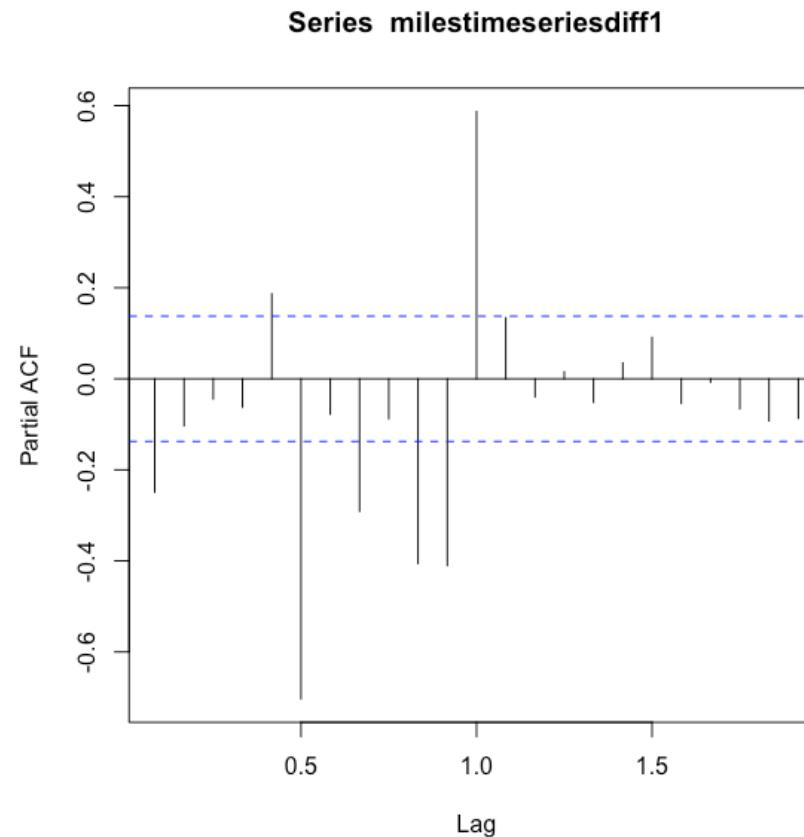
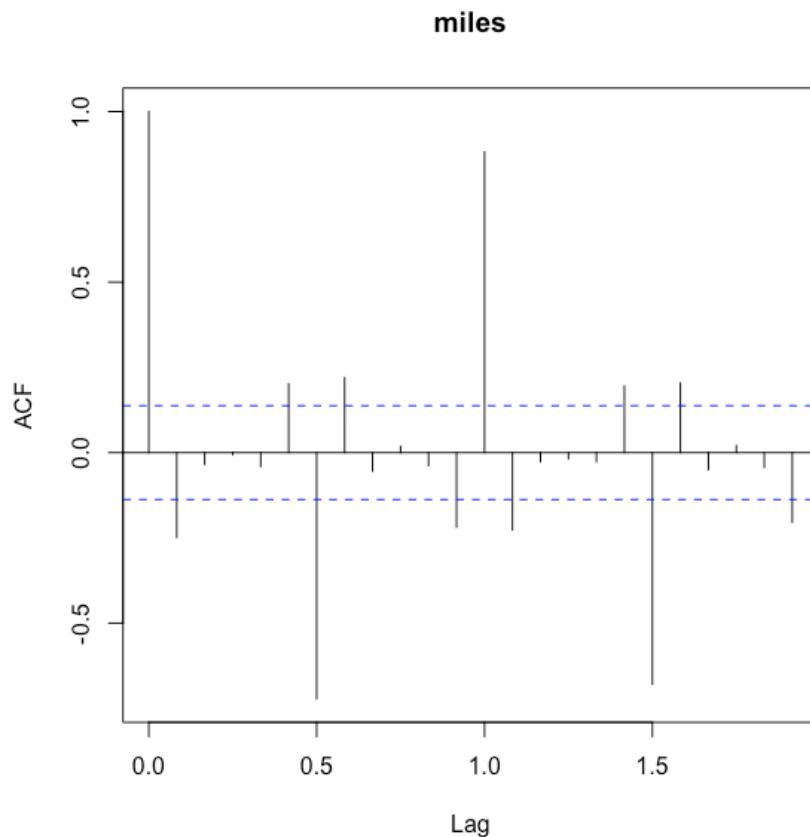
CSE 72026



Time Series Model Building Using ARIMA

Identification Phase

Step 4: Check ACF and PACF of differenced data to explore remaining dependencies.

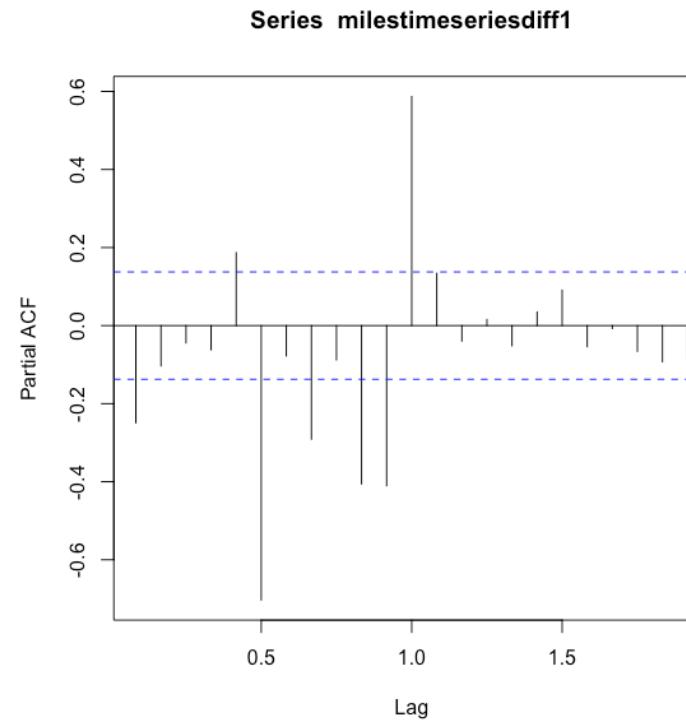
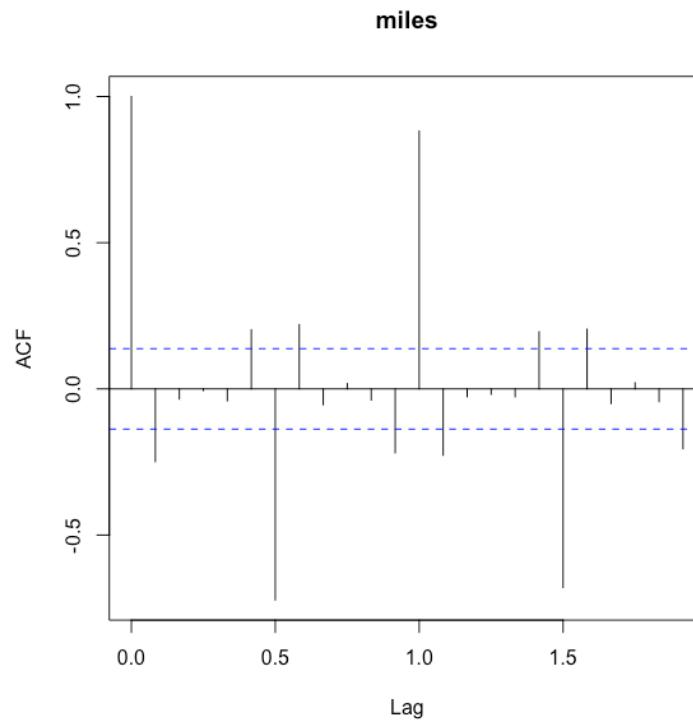


Time Series Model Building Using ARIMA

Identification Phase

Step 4: Check ACF and PACF of differenced data to explore remaining dependencies.

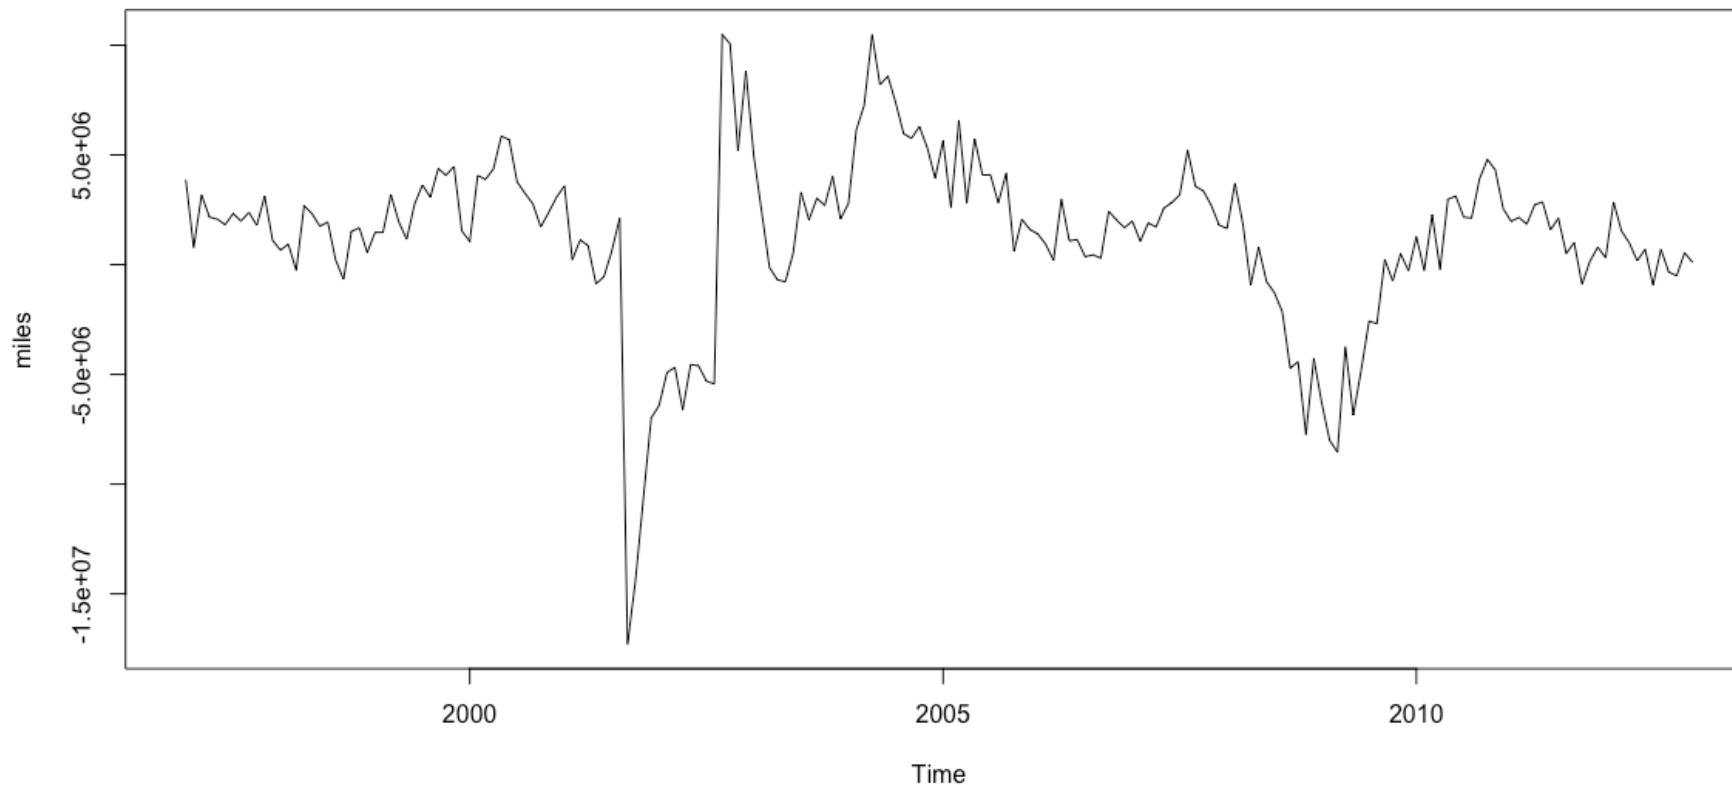
The differenced series looks somewhat stationary but has strong seasonal lags.



Time Series Model Building Using ARIMA

Identification Phase

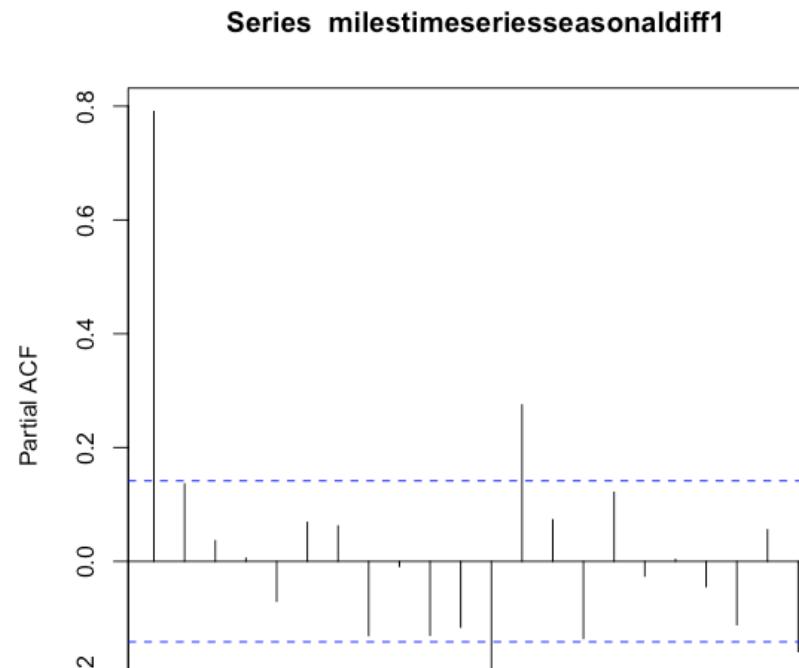
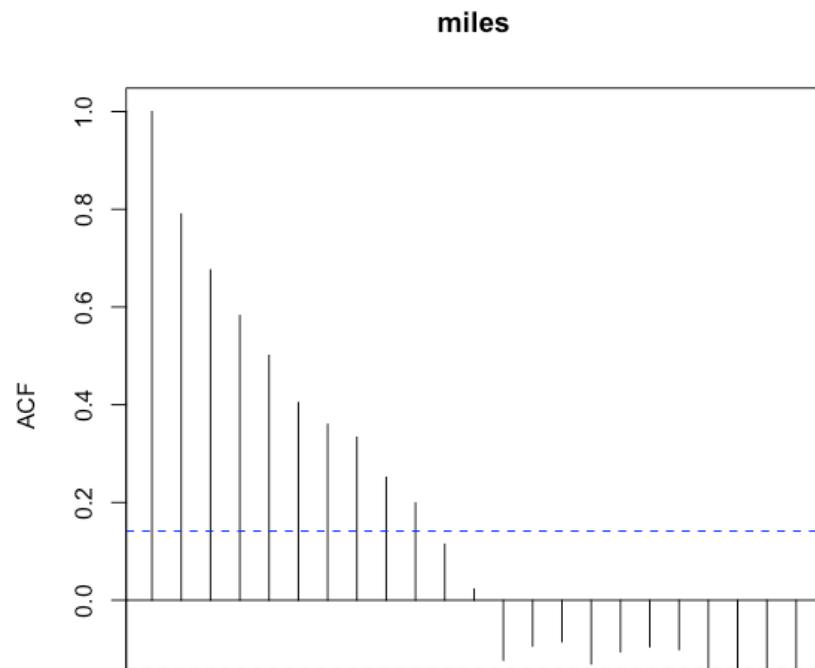
Step 5: Perform seasonal differencing (t_0-t_{12} , t_1-t_{13} , etc.) on the original time series to get seasonal stationarity. This is the same as an ARIMA(0,0,0)(0,1,0)₁₂ model.



Time Series Model Building Using ARIMA

Identification Phase

Step 6: Check ACF and PACF of seasonally differenced data to explore remaining dependencies and identify model(s).

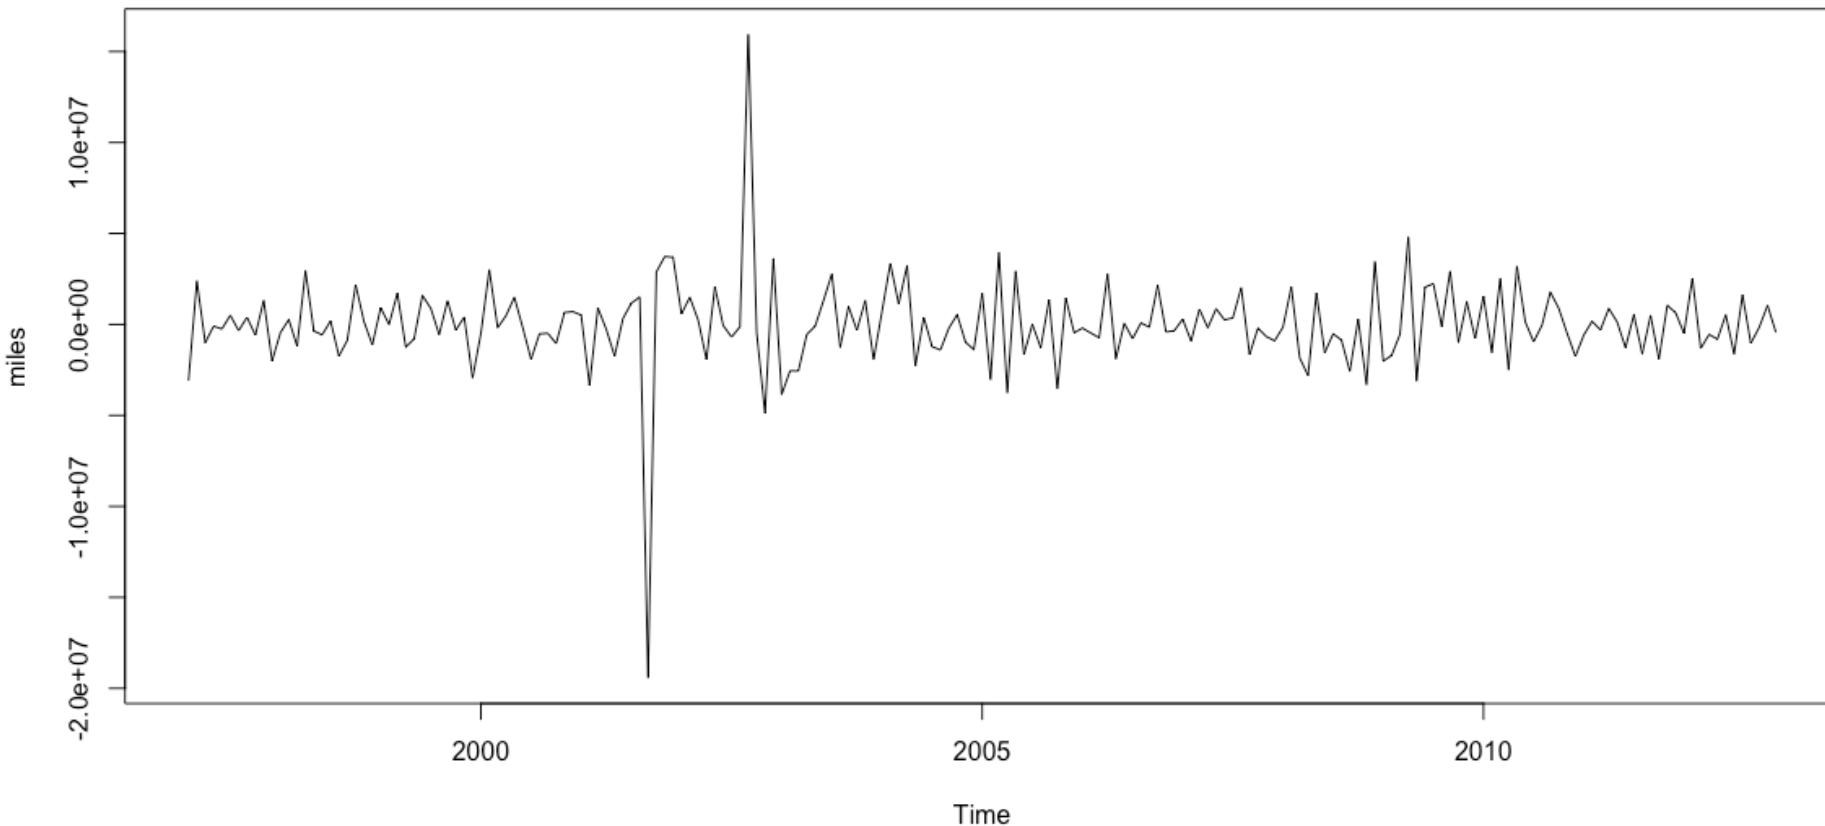


Strong positive autocorrelation indicates need for either an AR term or a non-seasonal differencing.

Time Series Model Building Using ARIMA

Identification Phase

Step 7: Perform a non-seasonal differencing on seasonally differenced data. This is like an $\text{ARIMA}(0,1,0)(0,1,0)_{12}$ model.



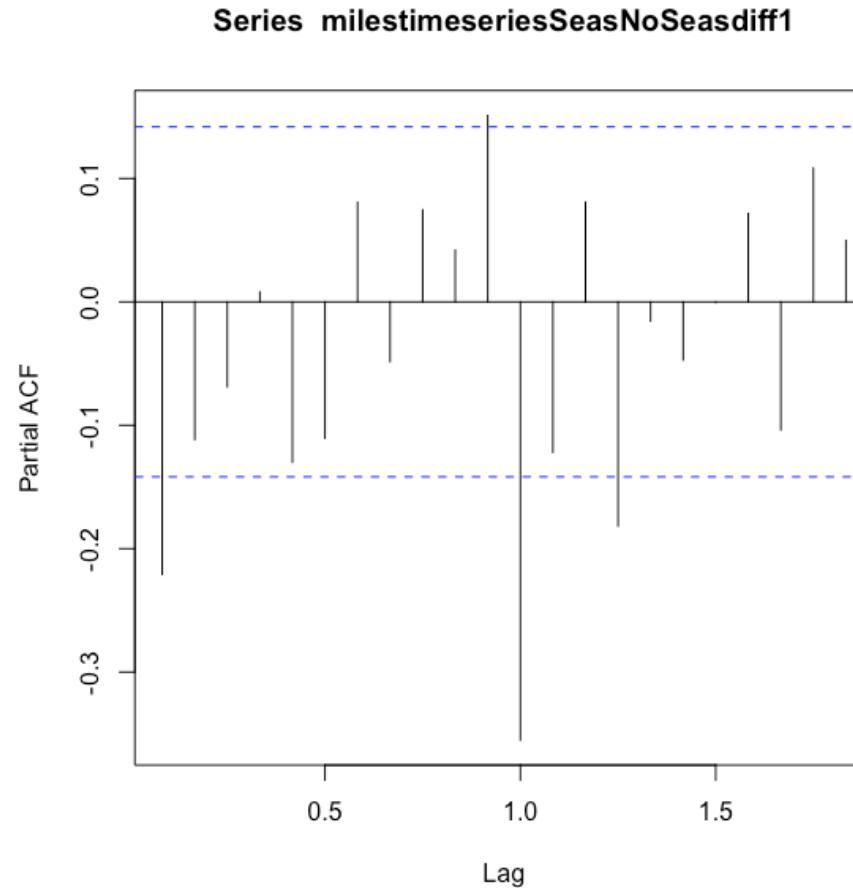
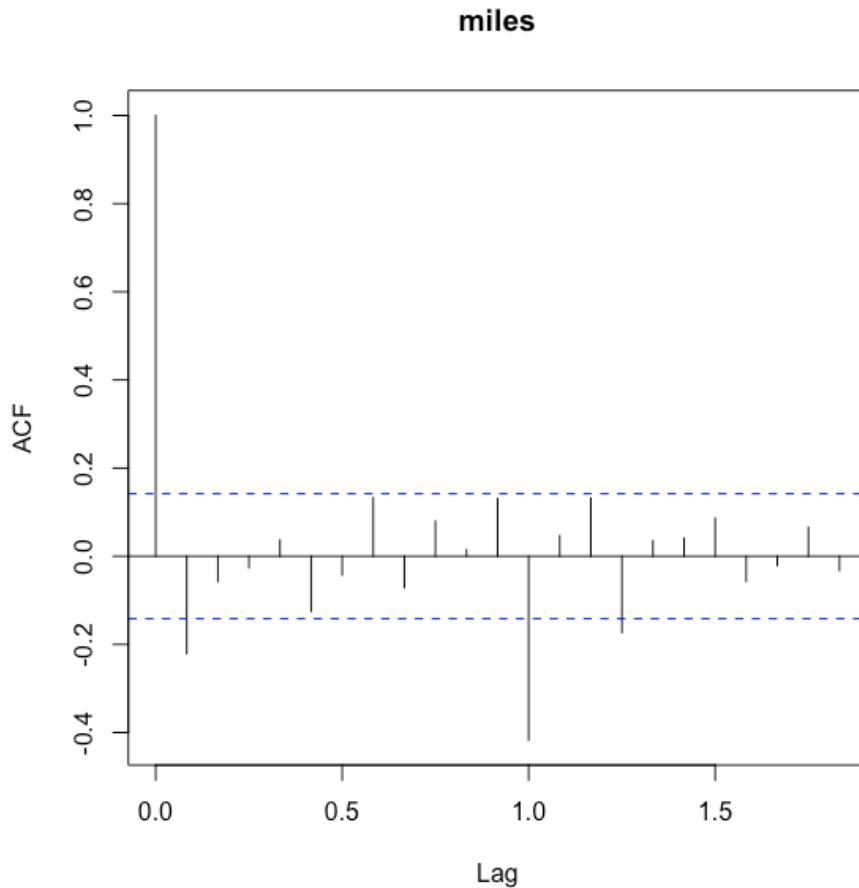
CSE 7202c



Time Series Model Building Using ARIMA

Identification Phase

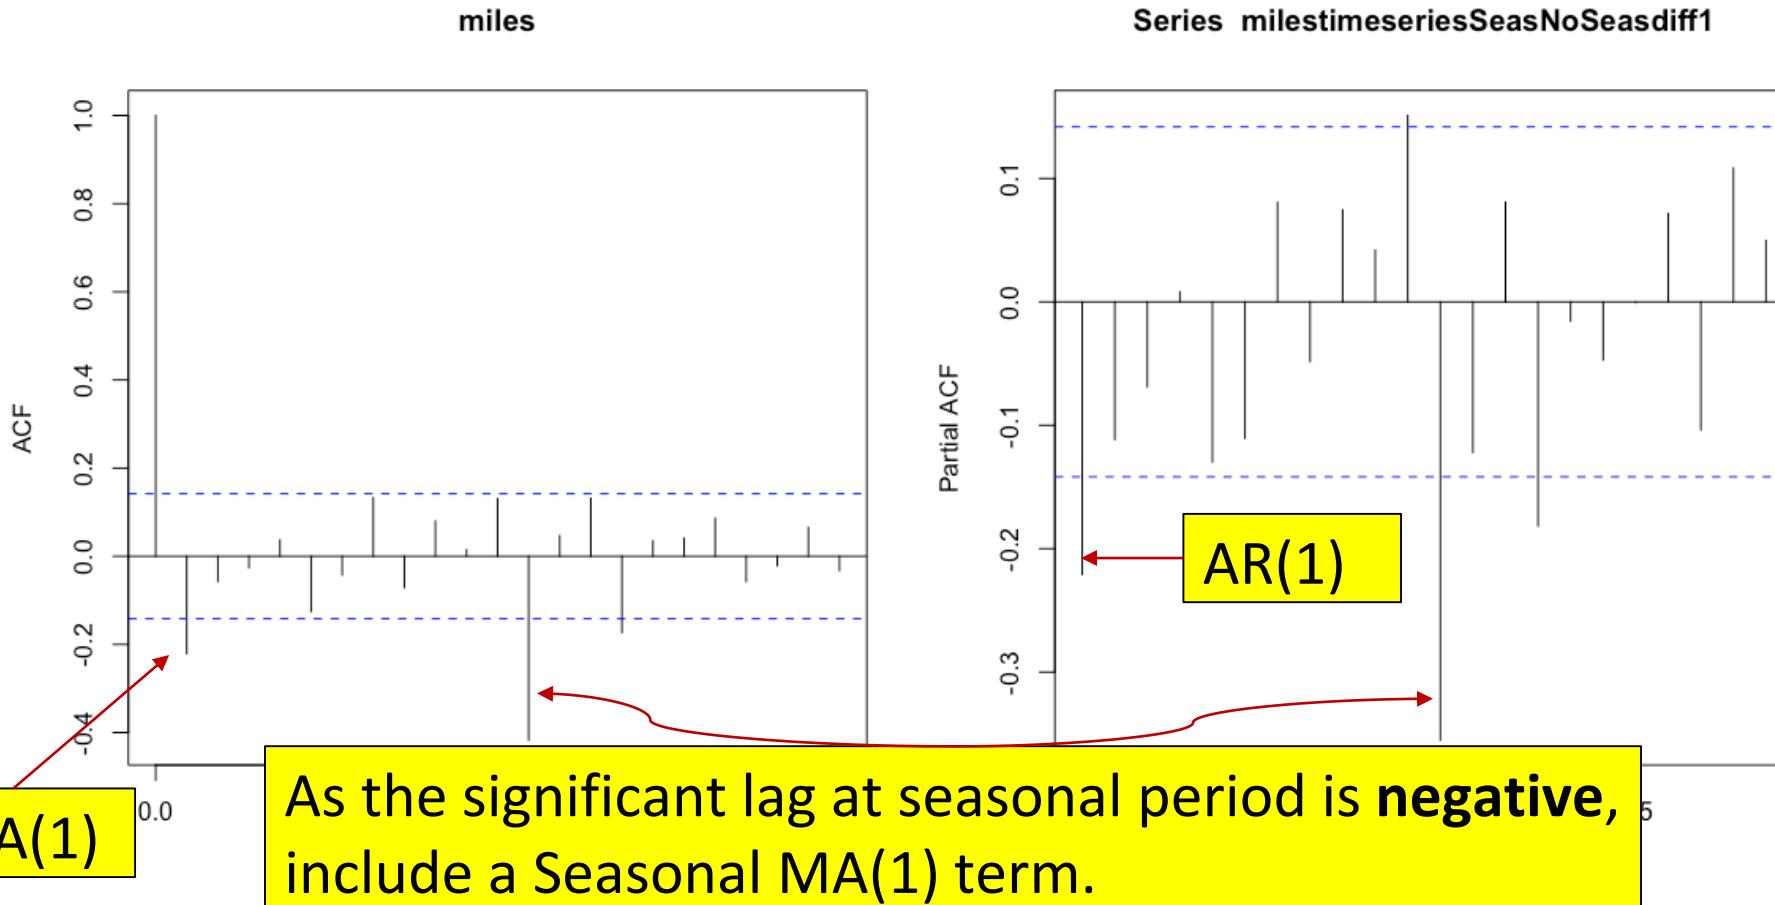
Step 8: Check ACF and PACF to explore remaining dependencies.



Time Series Model Building Using ARIMA

Identification Phase

Step 8: This indicates an ARIMA(1,1,1)(0,1,1)₁₂ model.



Time Series Model Building Using ARIMA

Parameter Estimation Phase

Step 9: Calculate parameters using the identified model(s). Use AIC to pick the best model.

```
Series: milestoneseries
ARIMA(1,1,1)(0,1,1)[12]
```

Coefficients:

	ar1	ma1	sma1
	0.4501	-0.7035	-0.7393
s.e.	0.1755	0.1407	0.0641

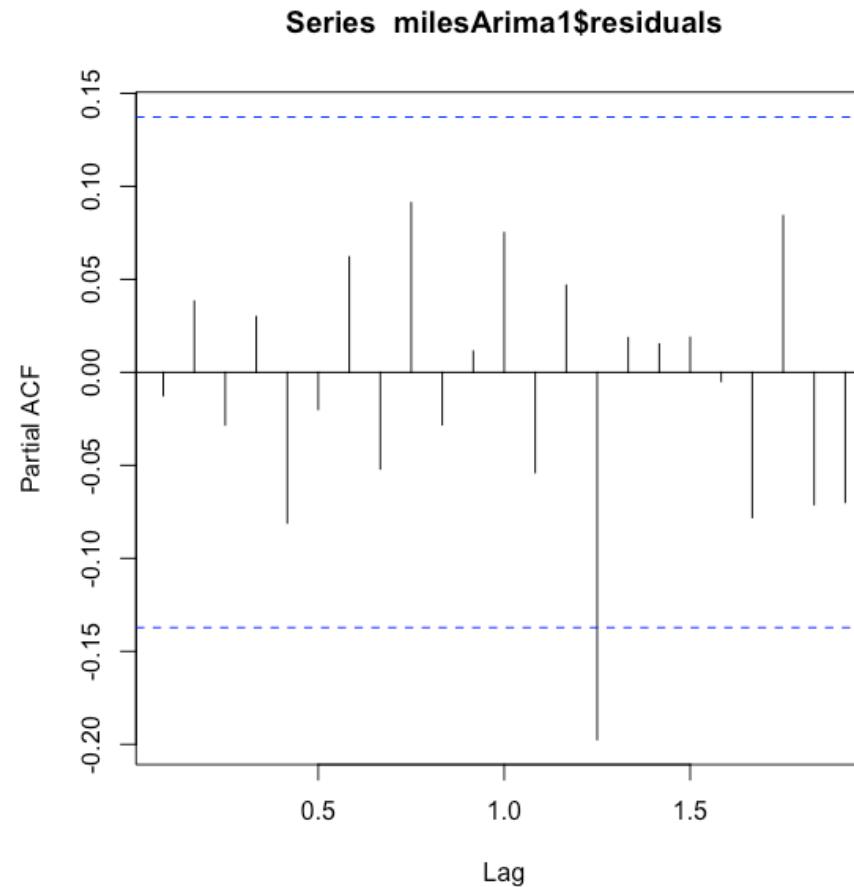
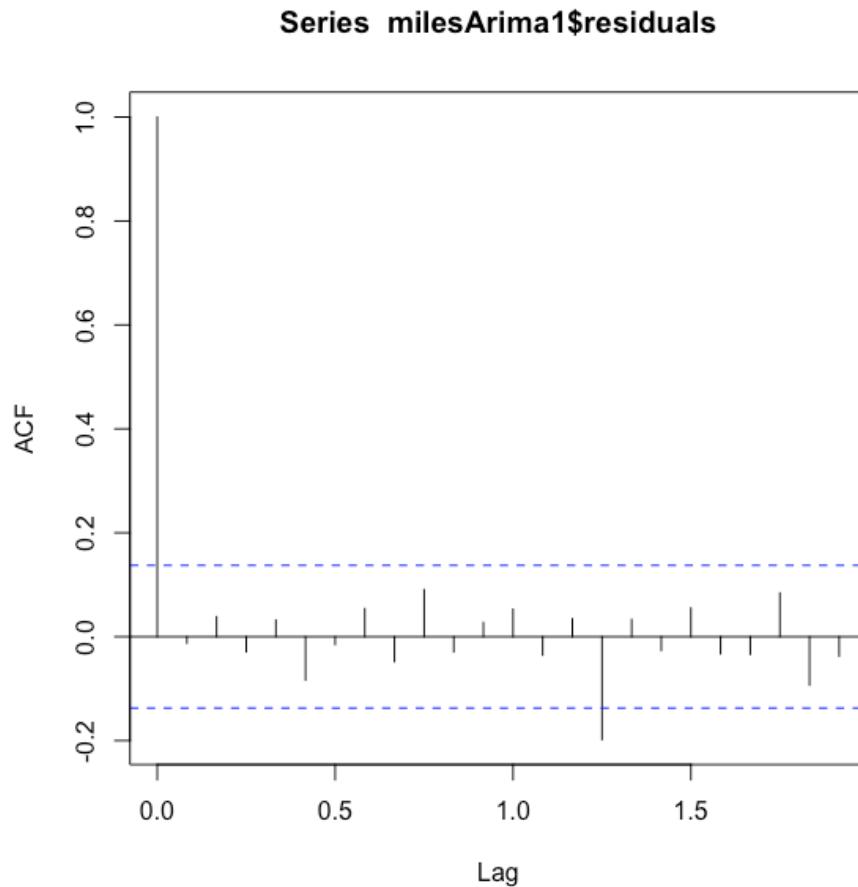
σ^2 estimated as 3.917e+12: log likelihood=-3043.49

AIC=6094.99 AICc=6095.2 BIC=6107.99

Time Series Model Building Using ARIMA

Evaluation Phase

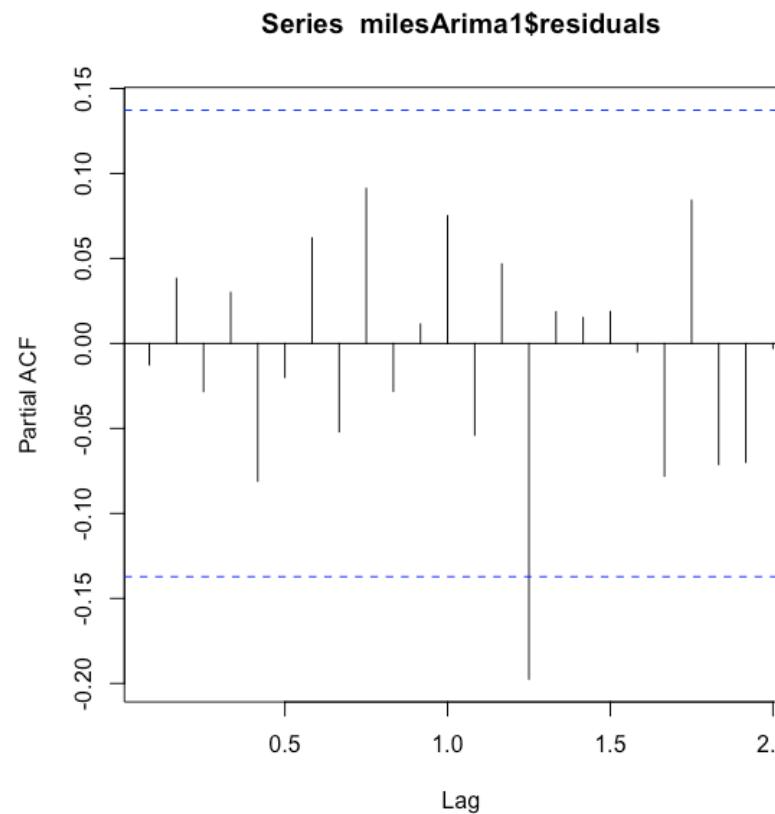
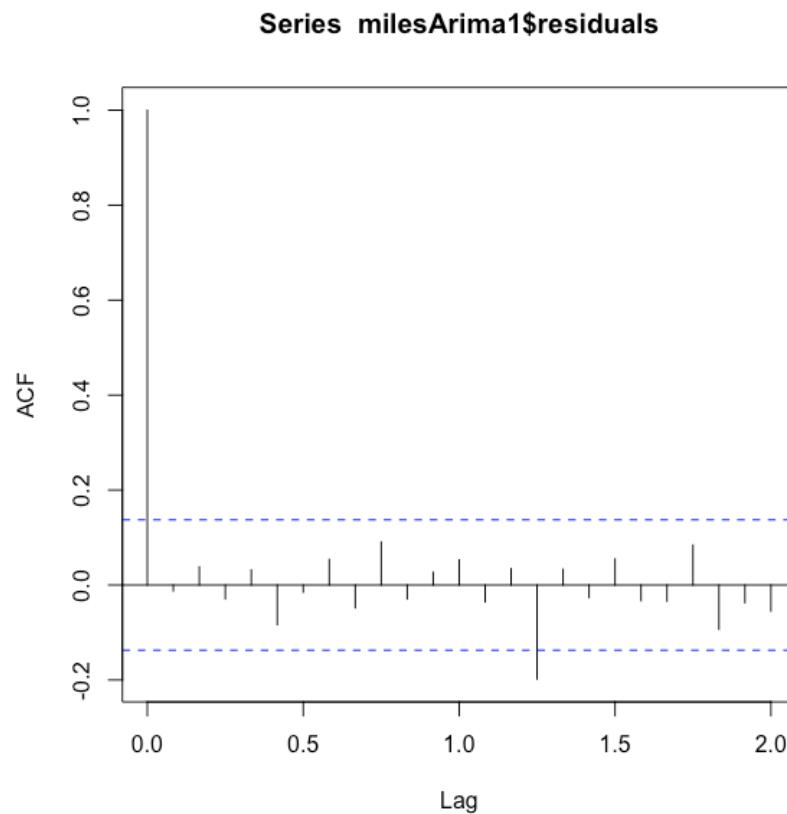
Step 10: Check ACF and PACF of the residuals to evaluate model.



Time Series Model Building Using ARIMA

Evaluation Phase

Step 10: The residuals indicate white noise. Indicates a good model that can be used for forecasting.



Time Series Model Building Using ARIMA

Evaluation Phase

Step 10: The residuals indicate white noise. Can be checked using Ljung-Box test.

$$Q^* = n(n + 2) \sum_{k=1}^h \frac{r_k^2}{n - k}$$

h is the maximum lag being considered
 n is the # of observations (length of the time series)
 r_k is the autocorrelation

* For non-seasonal time series, use $h = \min(10, n/5)$
For seasonal time series, use $h = \min(2m, n/5)$, where m is the seasonal period

Box-Ljung test

```
data: milesArima1$residuals
X-squared = 21.65, df = 24, p-value = 0.6002
```

If residuals are white noise (purely random), then Q has a χ^2 distribution with $h-p$ degrees of freedom, where p is the number of parameters estimated in the model

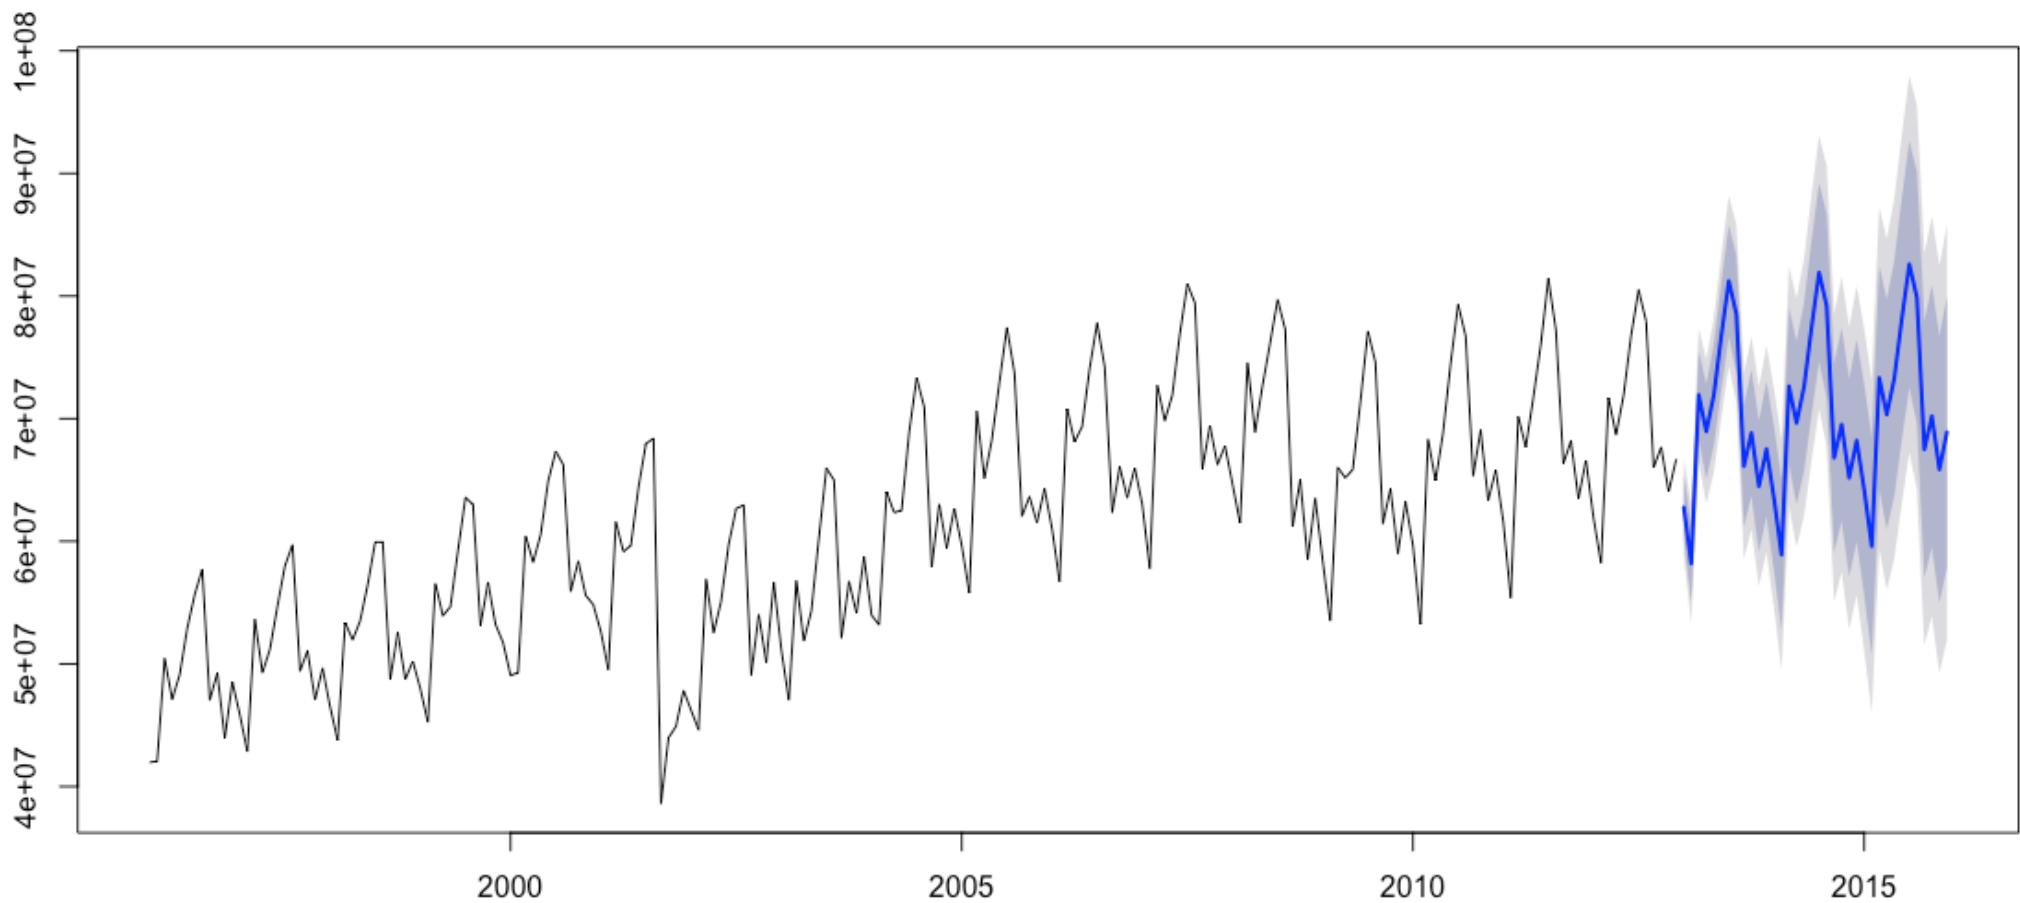
* <http://robjhyndman.com/hyndtsight/ljung-box-test/>

Time Series Model Building Using ARIMA

Forecasting Phase

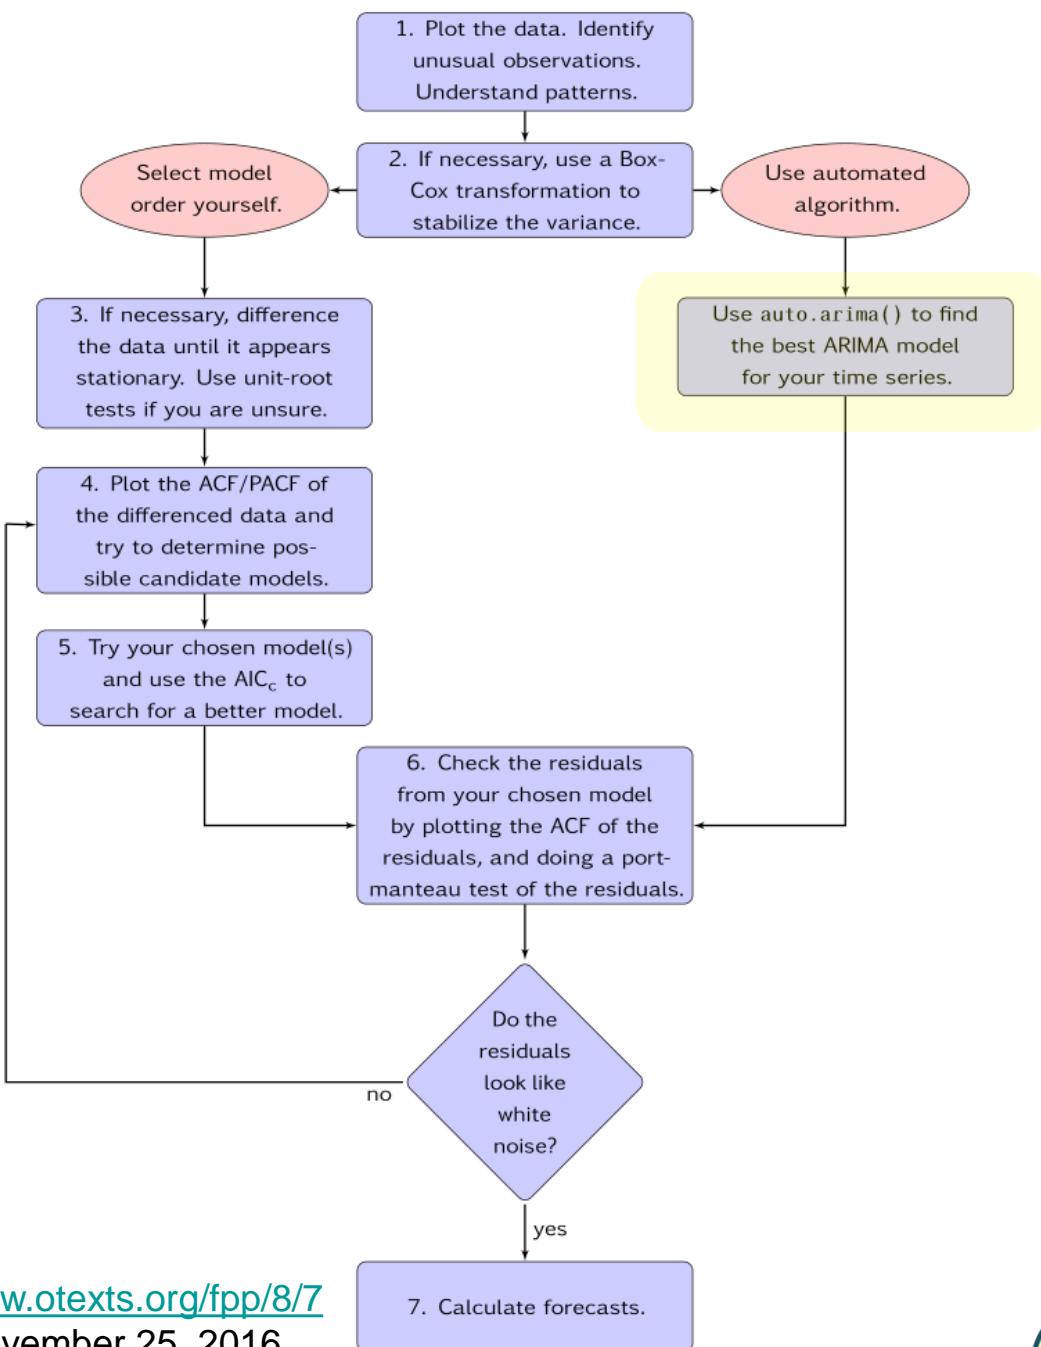
Step 11: Start forecasting.

Forecasts from ARIMA(1,1,1)(0,1,1)[12]



Model Selection in Practice

There are techniques that automate model selection



Source: <https://www.otexts.org/fpp/8/7>
Last accessed: November 25, 2016

Model Selection in Practice

`auto.arima()` in R using Hyndman-Khandakar algorithm.

1. Find d for stationarity.
2. p and q are selected using a stepwise search to minimize AICc on the differenced data.
 - a. The best model among the following is selected as the current model: ARIMA(2, d ,2), ARIMA(0, d ,0), ARIMA(1, d ,0), ARIMA(0, d ,1). If $d=0$, constant c is included; if $d\geq 1$, c is set to 0.
 - b. p and q are varied by ± 1 .
 - c. c is included/excluded.
 - d. 2(b) and (c) are repeated till AICc doesn't reduce further.

Source: <https://www.otexts.org/fpp/8/7>
Last accessed: November 25, 2016



Forecast using Auto ARIMA

```
Series: milestoneseries  
ARIMA(1,0,1)(0,1,1)[12] with drift
```

Coefficients:

ar1	ma1	sma1	drift
0.9092	-0.2128	-0.7257	108456.69
s.e.	0.0358	0.0873	31546.77

```
sigma^2 estimated as 3.834e+12: log likelihood=-3056.9  
AIC=6123.81 AICc=6124.13 BIC=6140.09
```

Auto ARIMA

```
Series: milestoneseries  
ARIMA(1,1,1)(0,1,1)[12]
```

Coefficients:

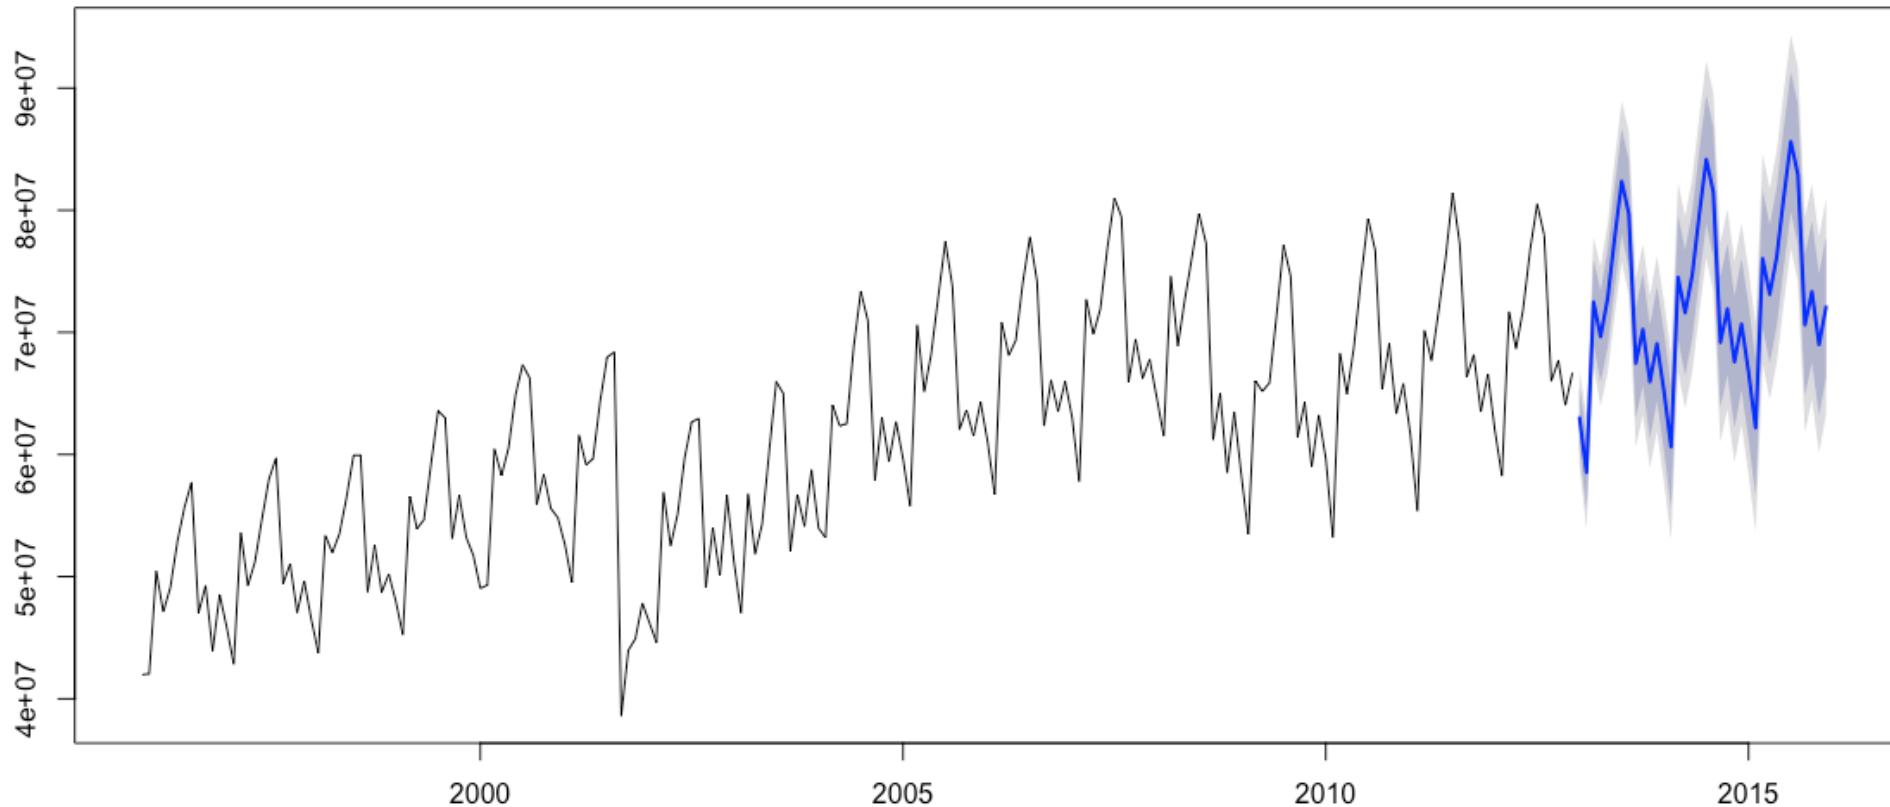
ar1	ma1	sma1
0.4501	-0.7035	-0.7393
s.e.	0.1755	0.1407

```
sigma^2 estimated as 3.917e+12: log likelihood=-3043.49  
AIC=6094.99 AICc=6095.2 BIC=6107.99
```

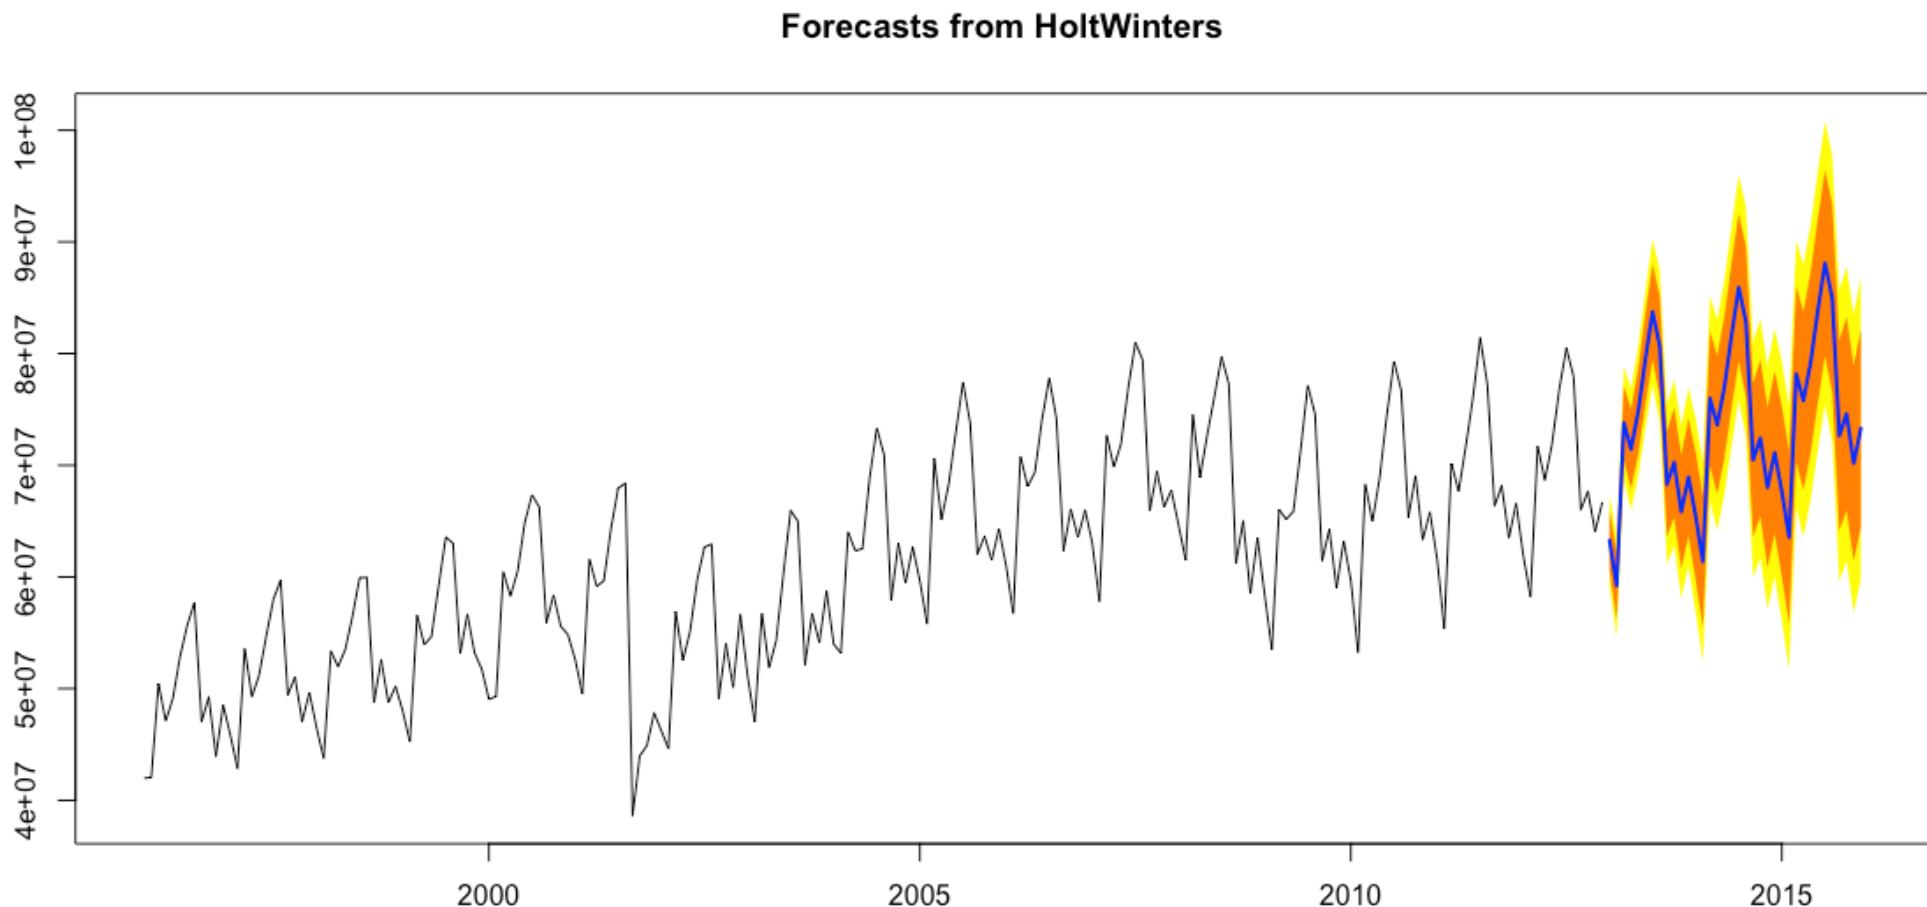
Manual ARIMA

Forecast using Auto ARIMA - RPM

Forecasts from ARIMA(1,0,1)(0,1,1)[12] with drift



Forecast using Holt-Winters - RPM



Forecast window	Holt-Winters	Manual ARIMA	Auto ARIMA
3-month	1.18%	0.62%	0.45%
6-month	1.81%	1.17%	0.85%
12-month	1.88%	1.58%	1.12%
18-month	2.10%	1.85%	1.14%
24-month	2.03%	2.19%	1.11%
36-month	1.98%	3.44%	1.61%

Comparing Models - MAPE

MAPE

0.04

0.03

0.02

0.01

0

Manual ARIMA

Holt-Winters

Auto ARIMA

3 6 9 12 15 18 21 24 27 30 33 36

Forecast Window (Months)

CSE 7202C



Time Series Model Building Using ARIMA

A nice summary of rules for identifying ARIMA models

<http://people.duke.edu/~rnau/arimrule.htm>

CSE 7202c





Manufacturing Case Study

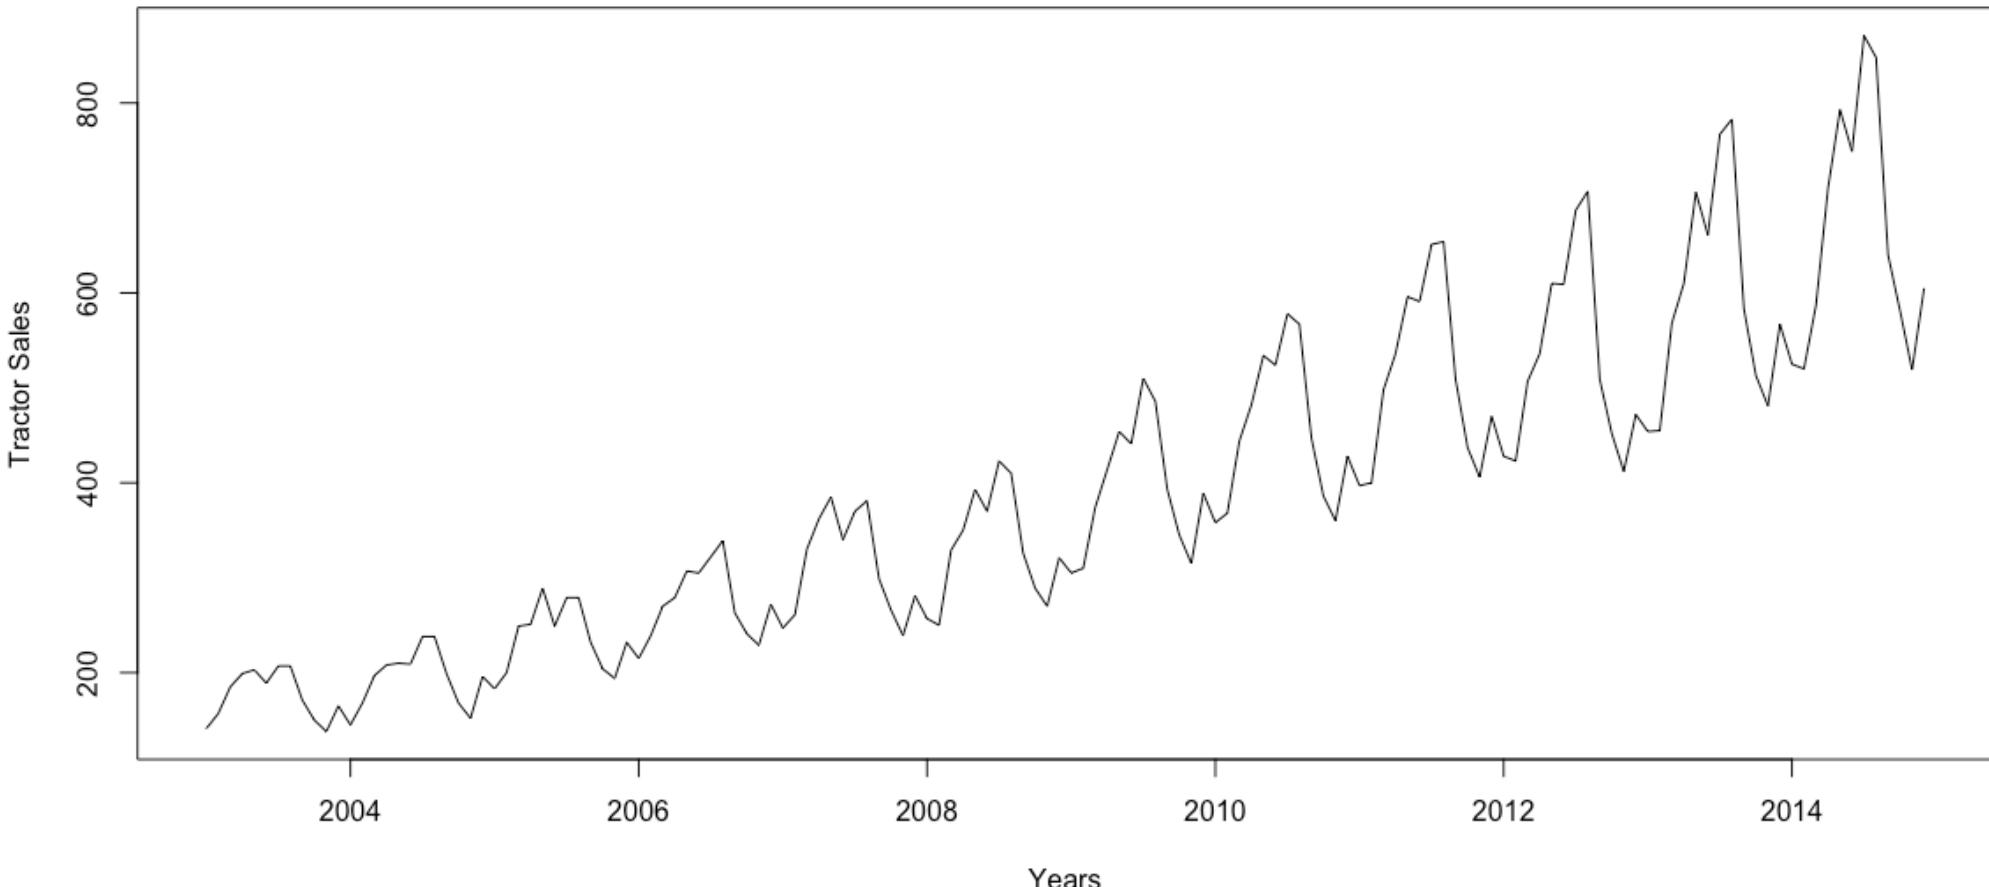
FORECASTING TRACTOR SALES

CSE 7202C



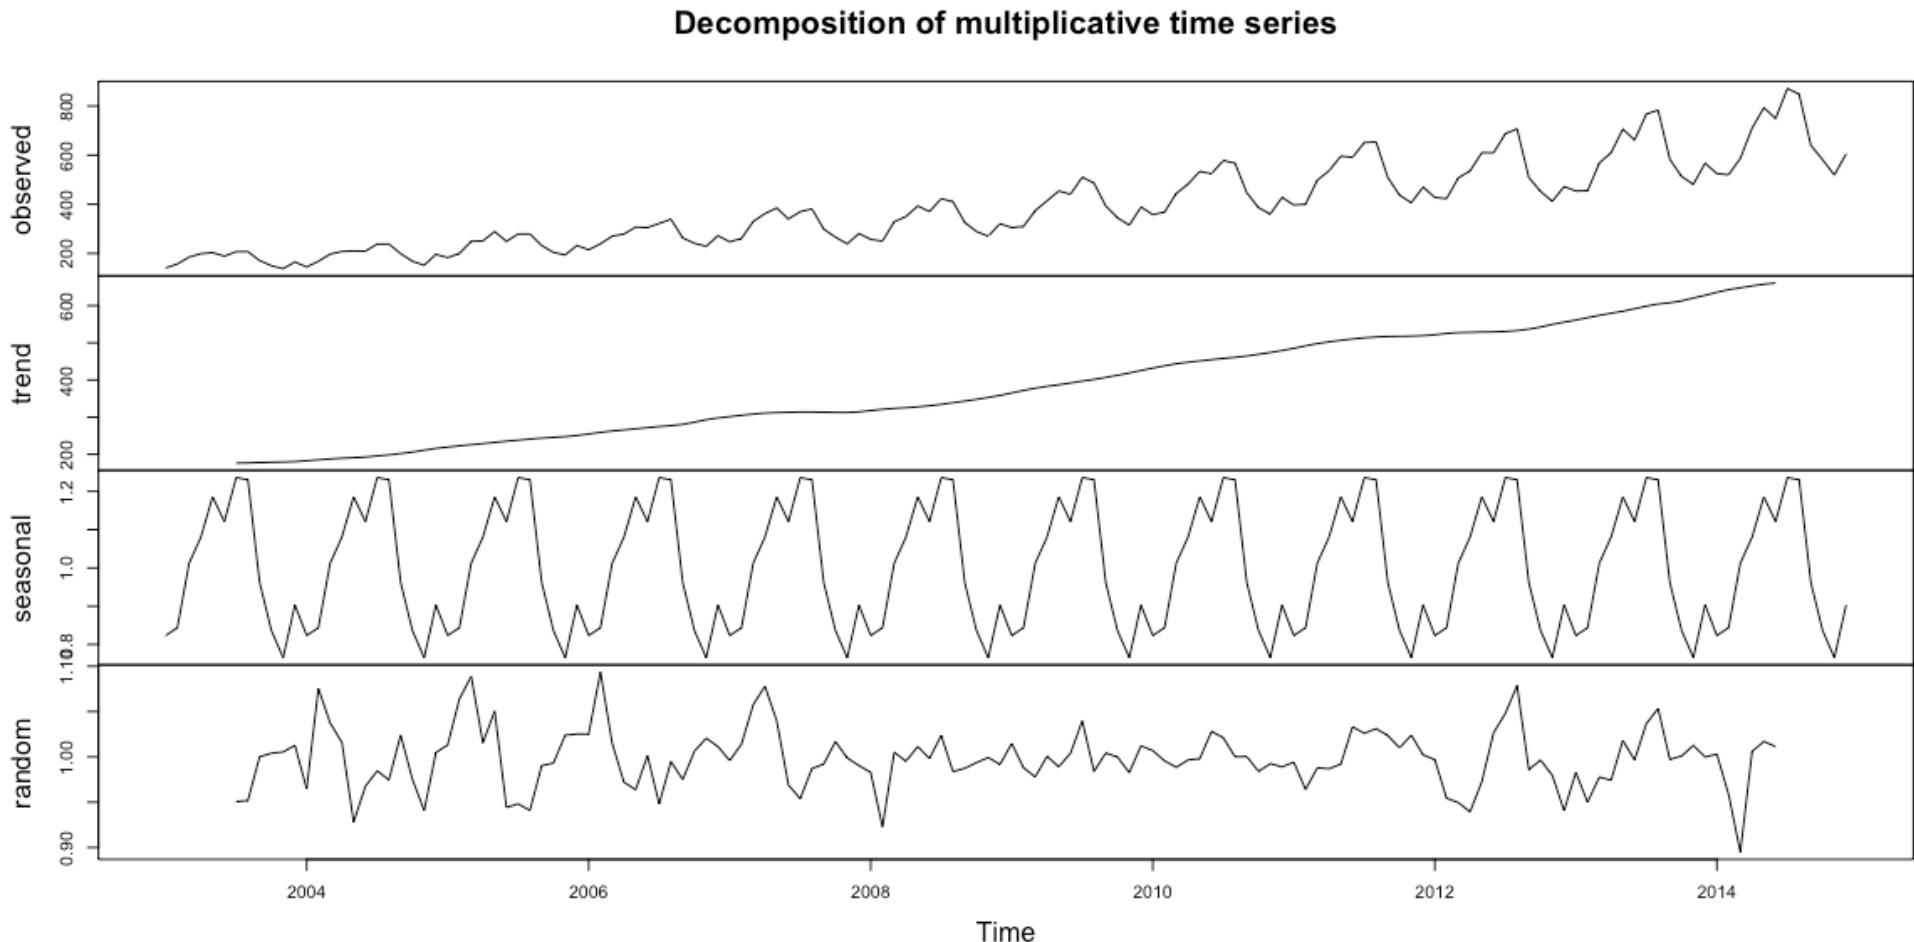
Tractor Sales Case Study

Step 1: Convert data into time series, and plot and decompose into components to understand it.



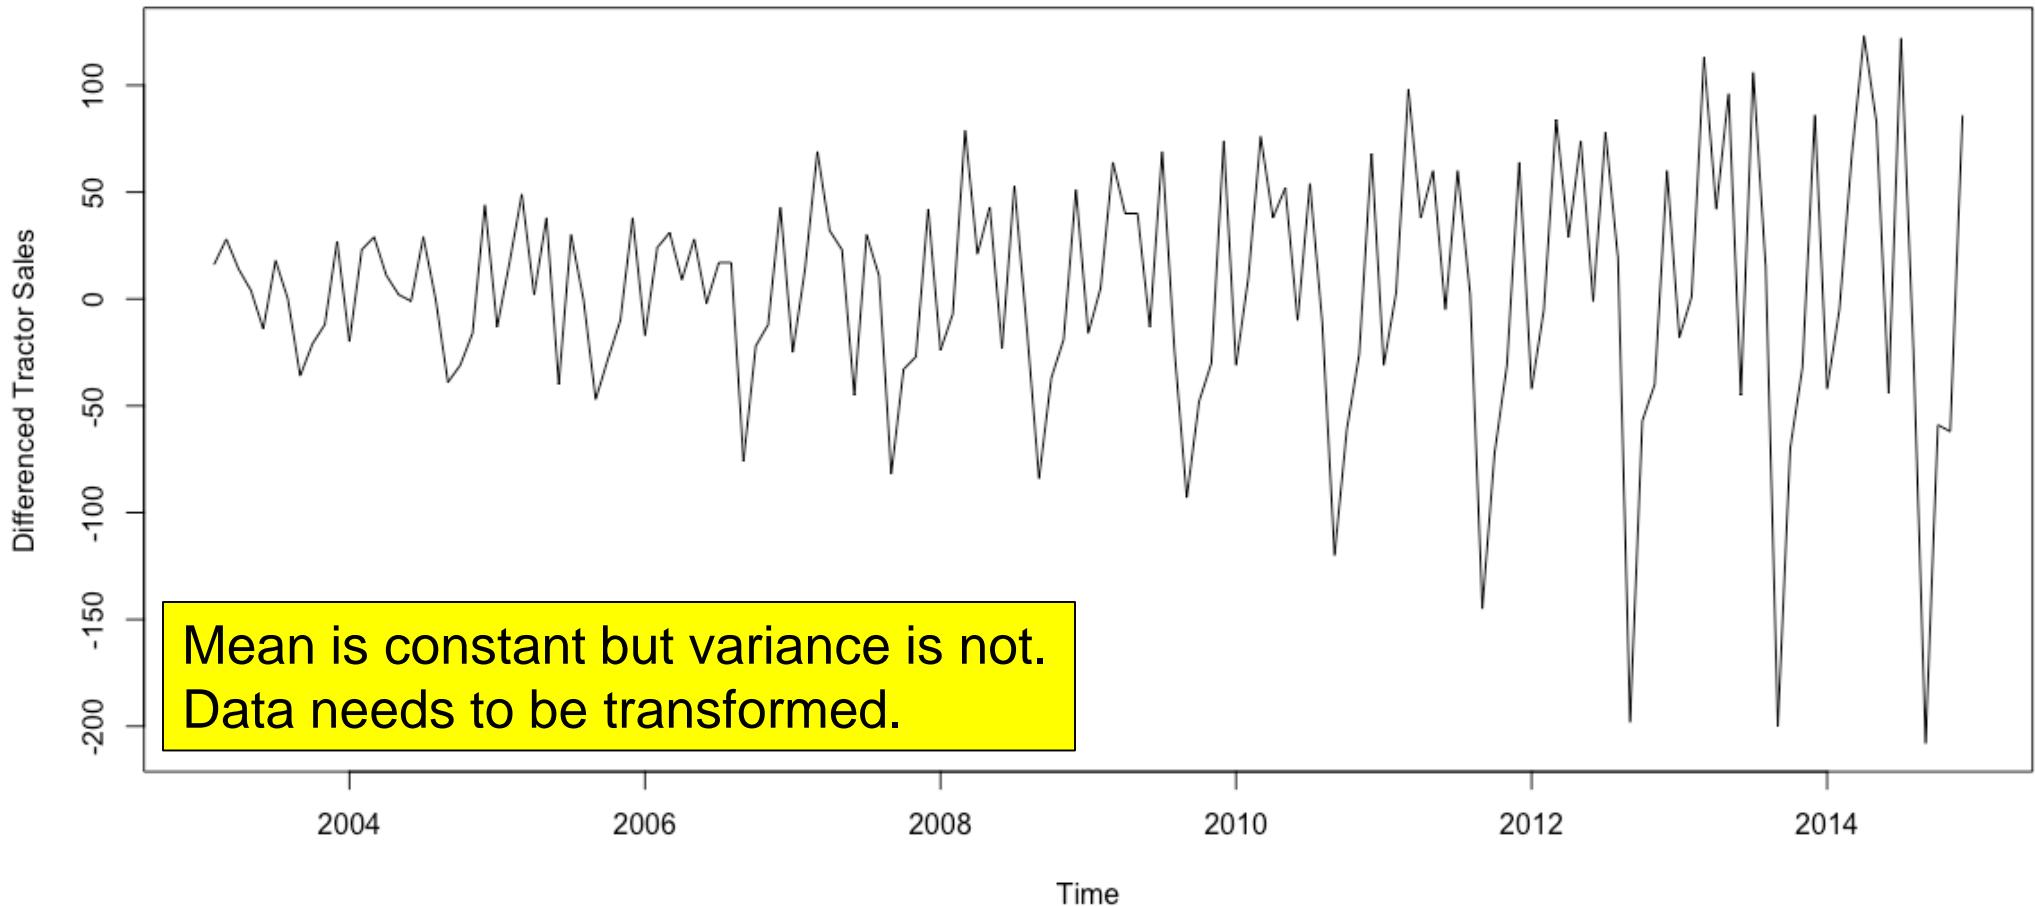
Tractor Sales Case Study

Step 1: Convert data into time series, and plot and decompose into components to understand it.



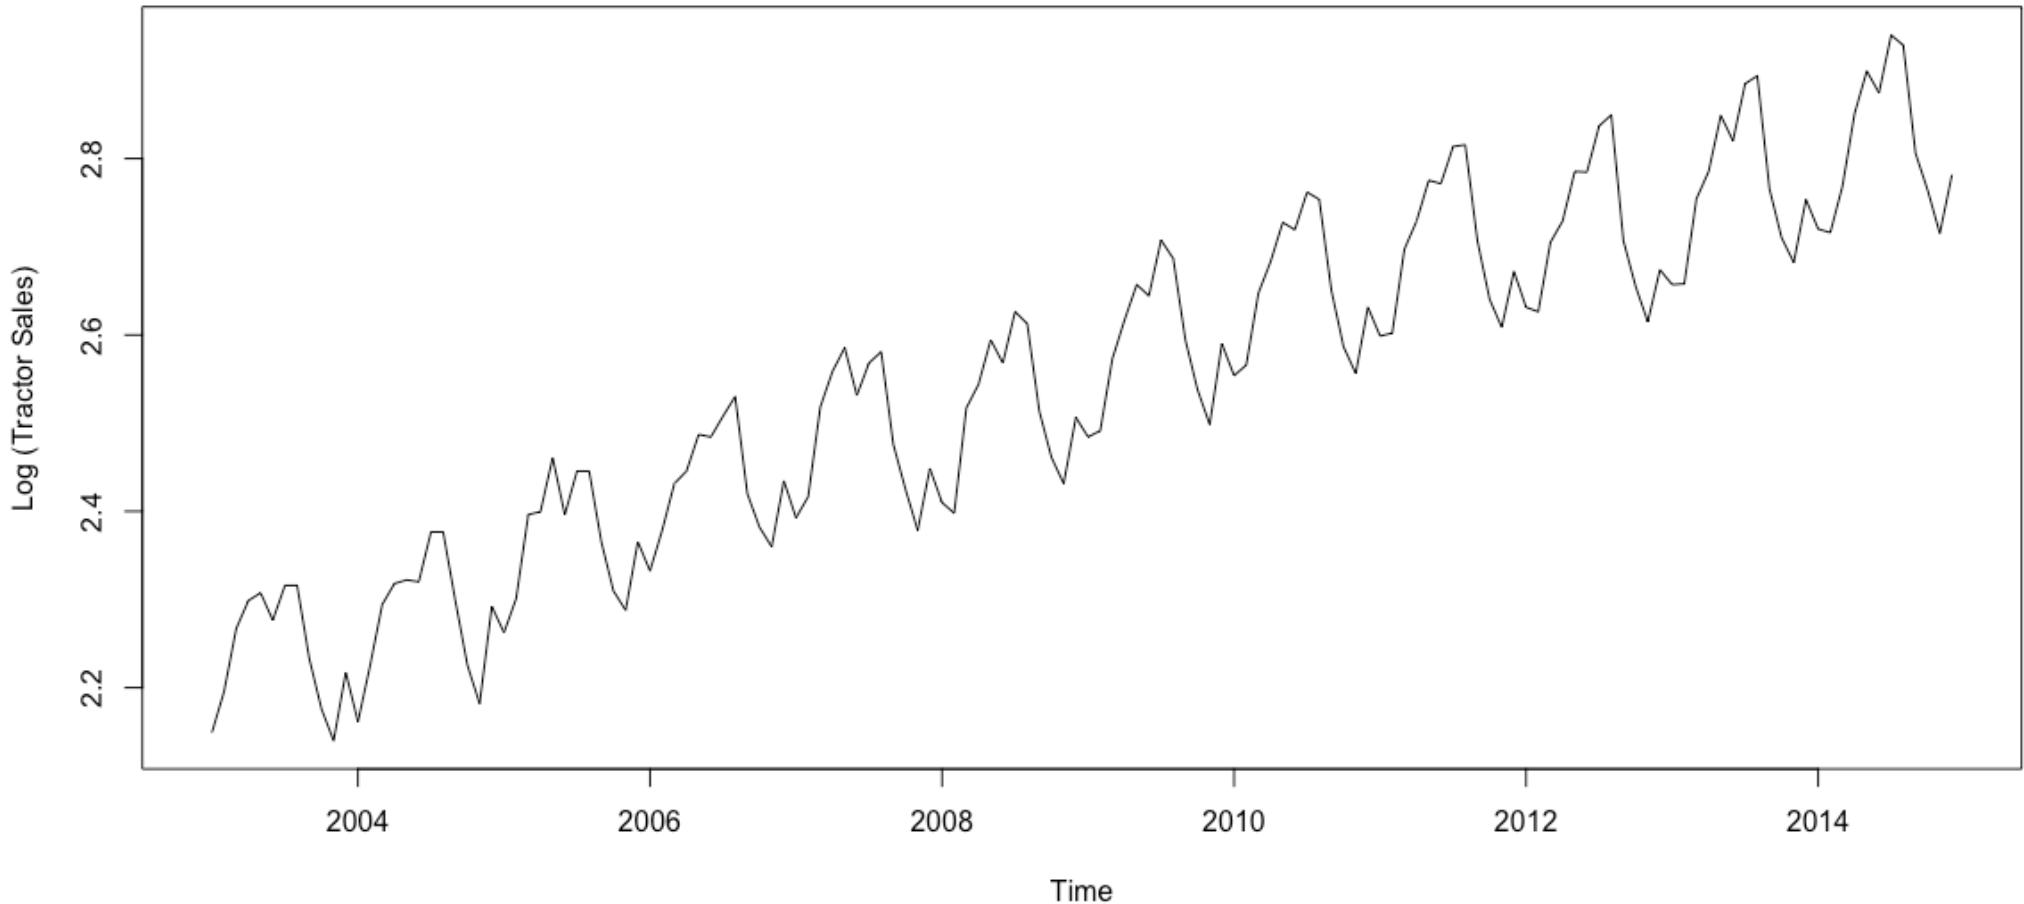
Tractor Sales Case Study

Step 2: Difference the data to make it stationary.



Tractor Sales Case Study

Step 3: Log transform the data.

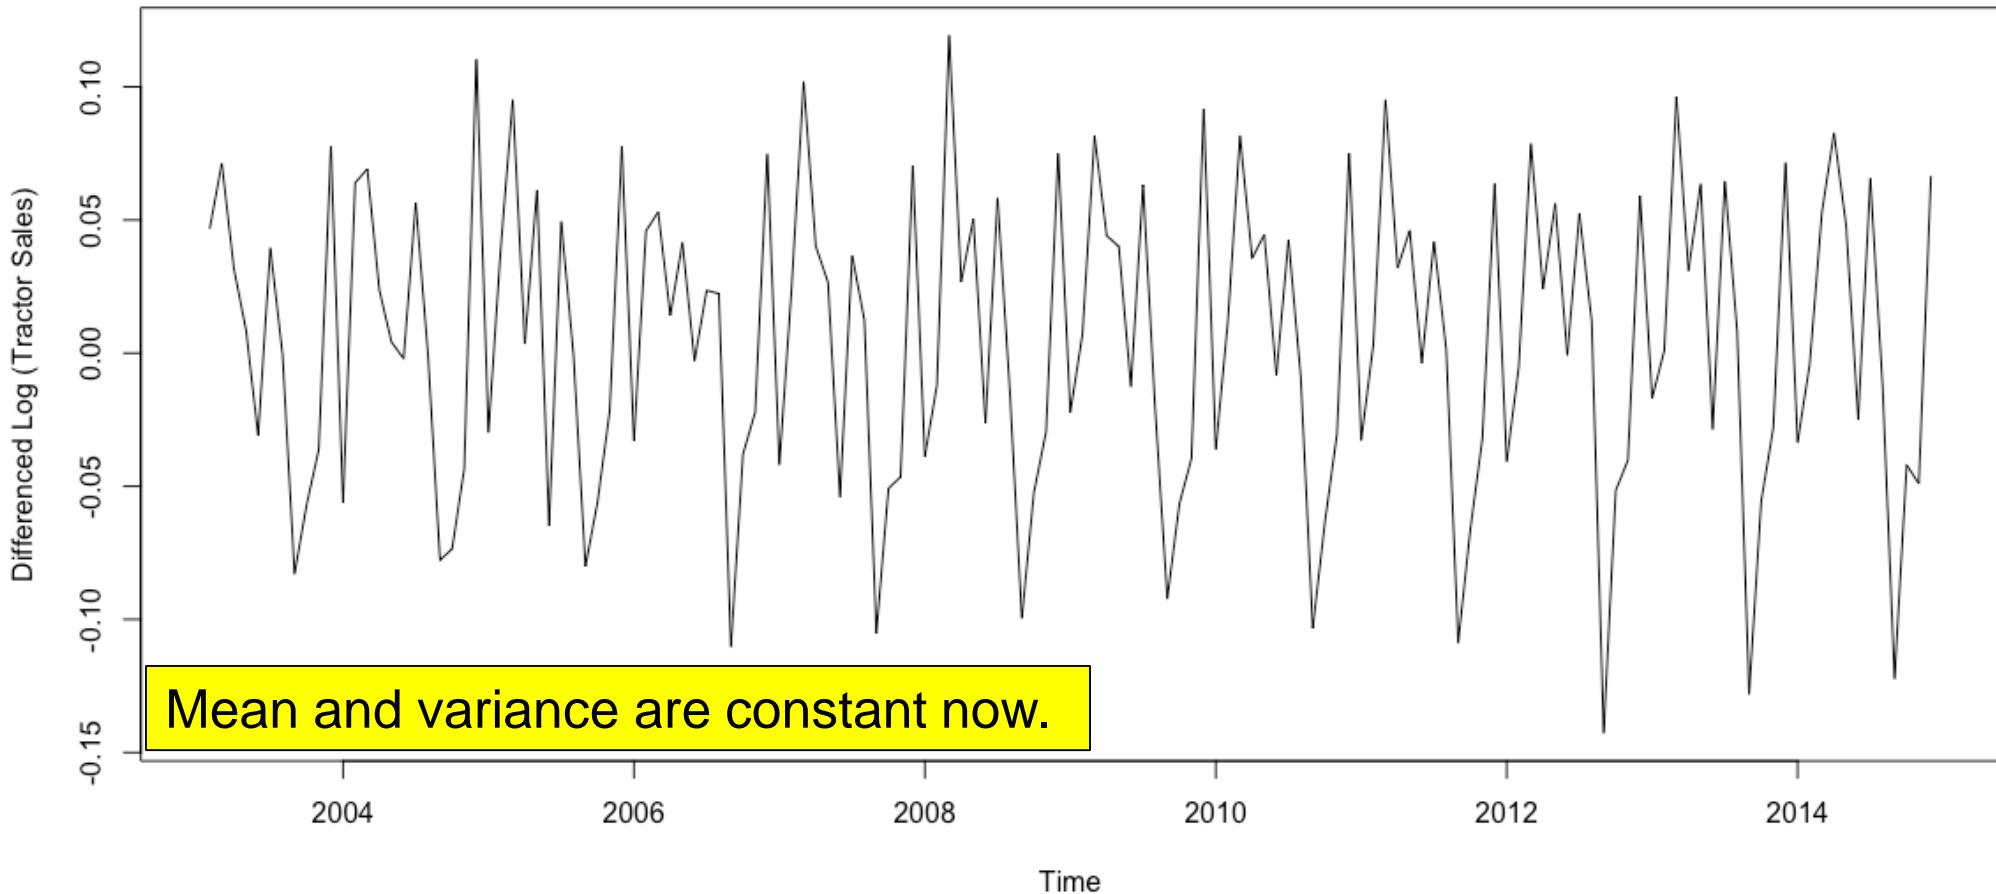


CSE 7202c



Tractor Sales Case Study

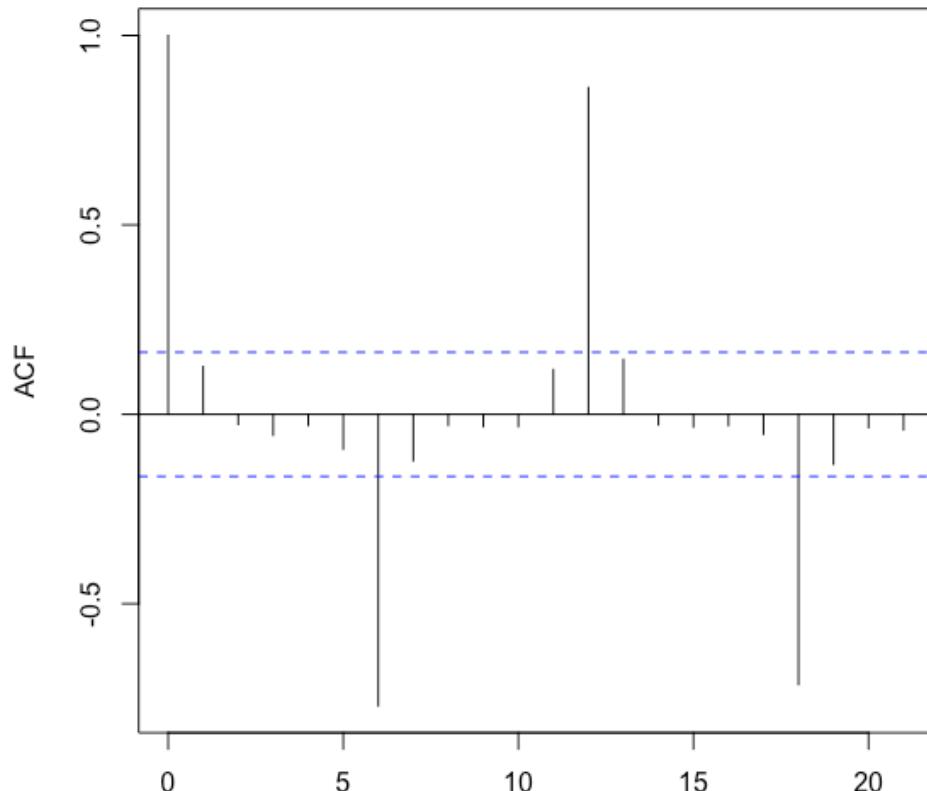
Step 4: Difference the log transformed data to check for stationarity.



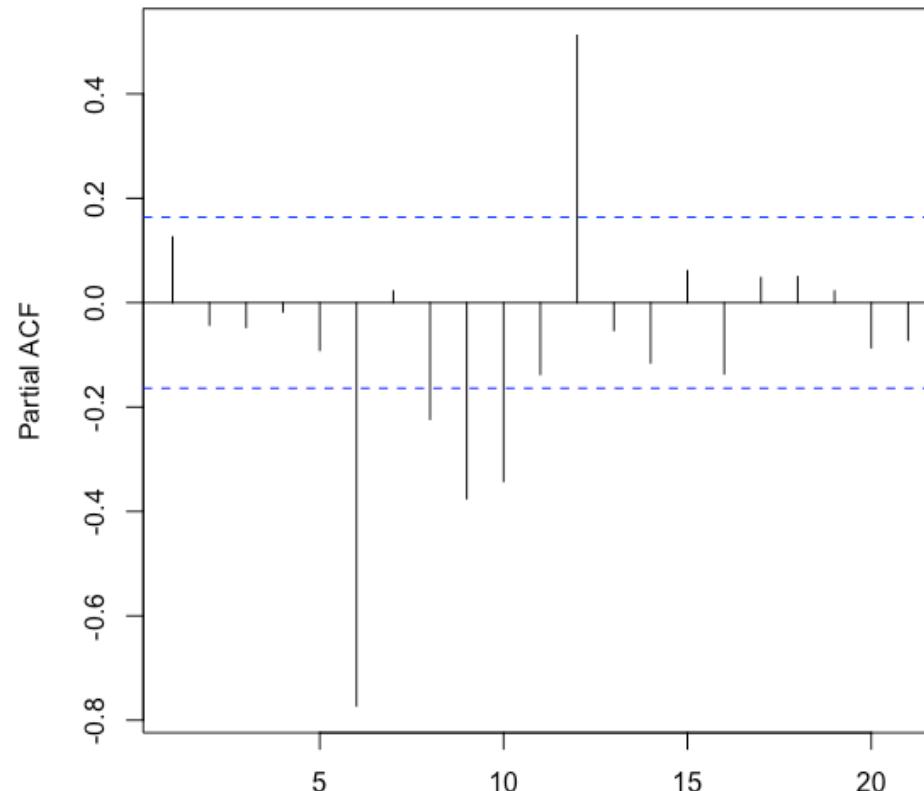
Tractor Sales Case Study

Step 5: Check ACF and PACF to explore remaining dependencies.

ACF Tractor Sales



PACF Tractor Sales



ACF has cutoff but PACF shows some decay. This indicates an MA process. Also, there is strong seasonality.

Tractor Sales Case Study

Step 6: Run Auto ARIMA.

```
Series: log10(TractorSalesTS)
ARIMA(0,1,1)(0,1,1)[12]
```

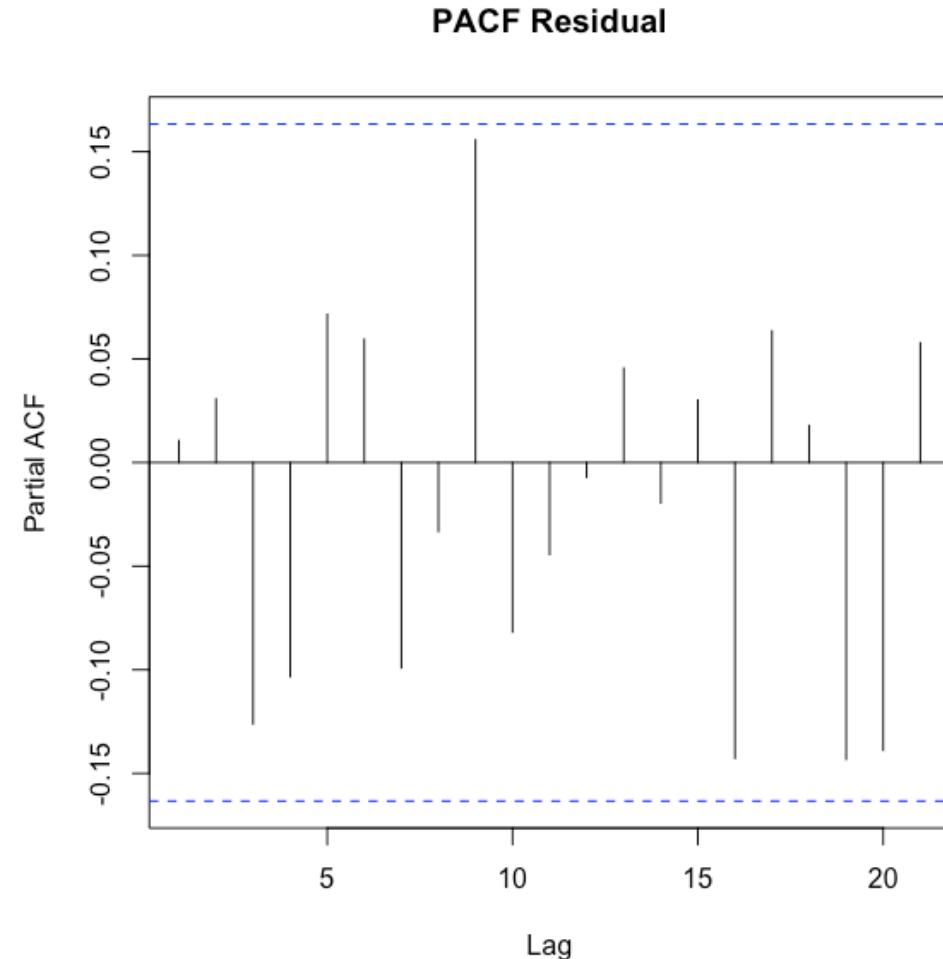
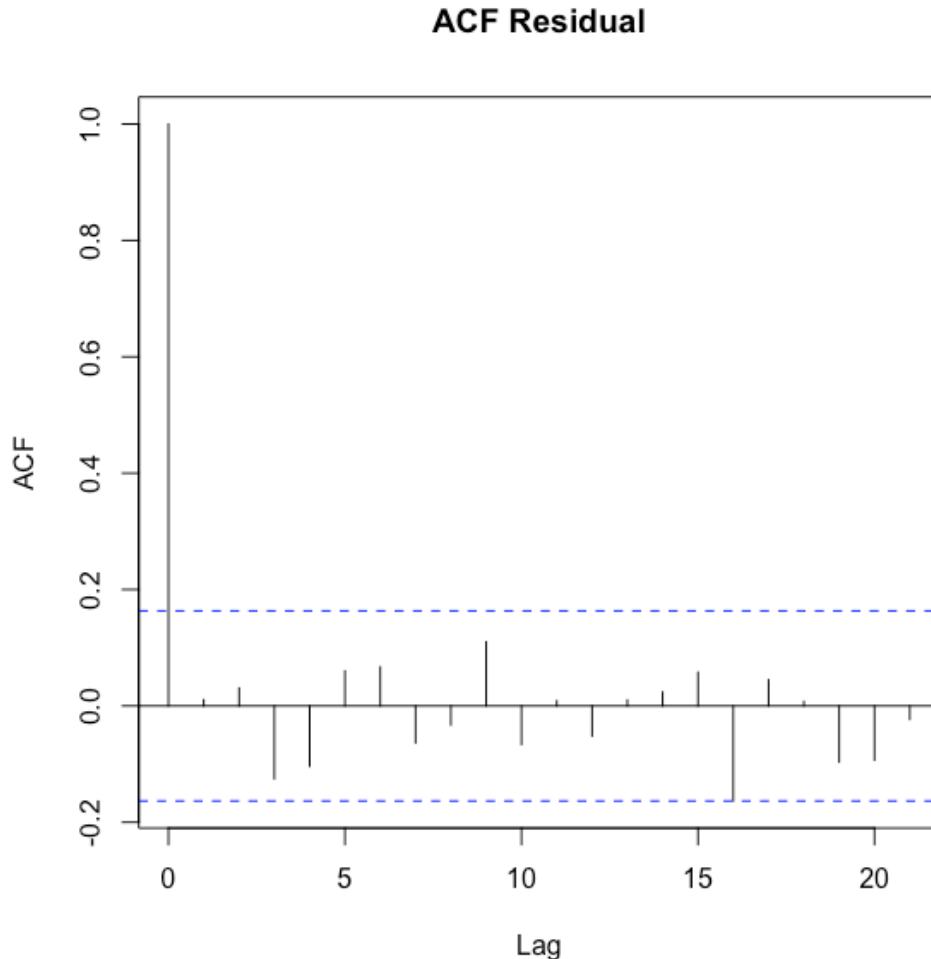
Coefficients:

	ma1	sma1
	-0.4047	-0.5529
s.e.	0.0885	0.0734

```
sigma^2 estimated as 0.0002571: log likelihood=354.4
AIC=-702.79    AICc=-702.6    BIC=-694.17
```

Tractor Sales Case Study

Step 7: Check ACF and PACF of residuals to ensure they are white noise.



Tractor Sales Case Study

Step 7: Use Box-Ljung test to verify residuals are white noise. *It is good enough to do the verification visually on ACF and PACF.*

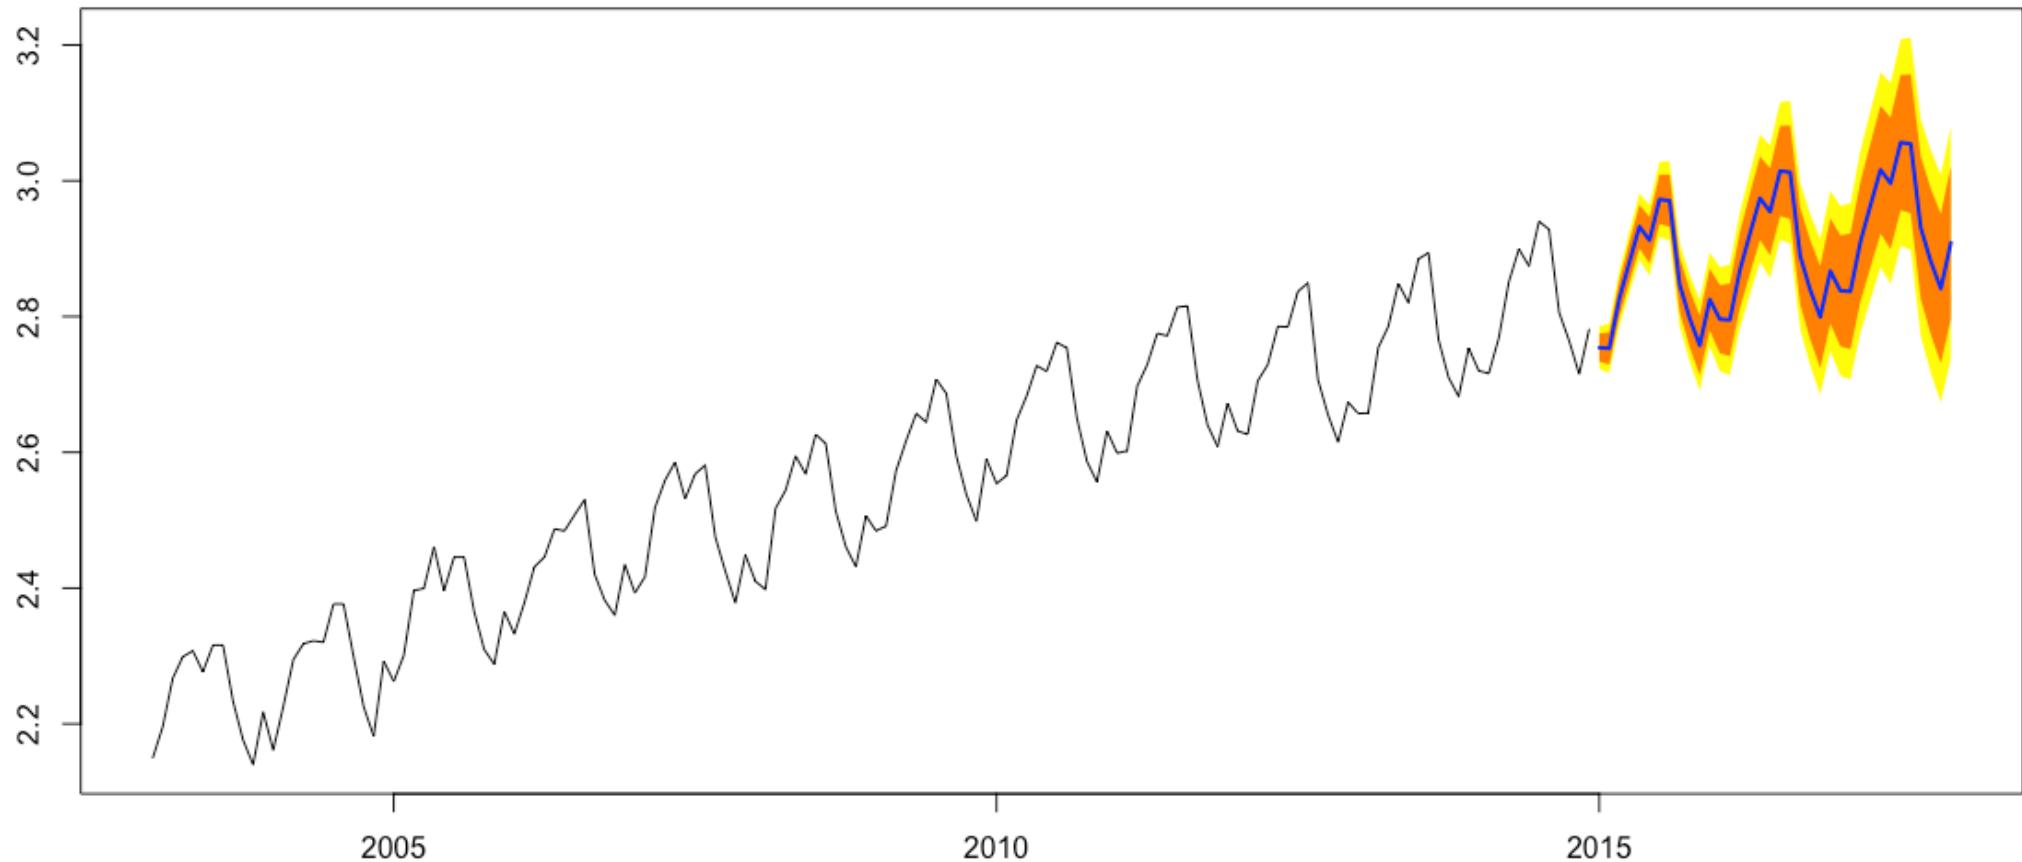
Box-Ljung test

```
data: LogTractorSalesARIMA$residuals  
X-squared = 26.219, df = 24, p-value = 0.3422
```

Tractor Sales Case Study

Step 8: Forecast.

Forecasts from ARIMA(0,1,1)(0,1,1)[12]

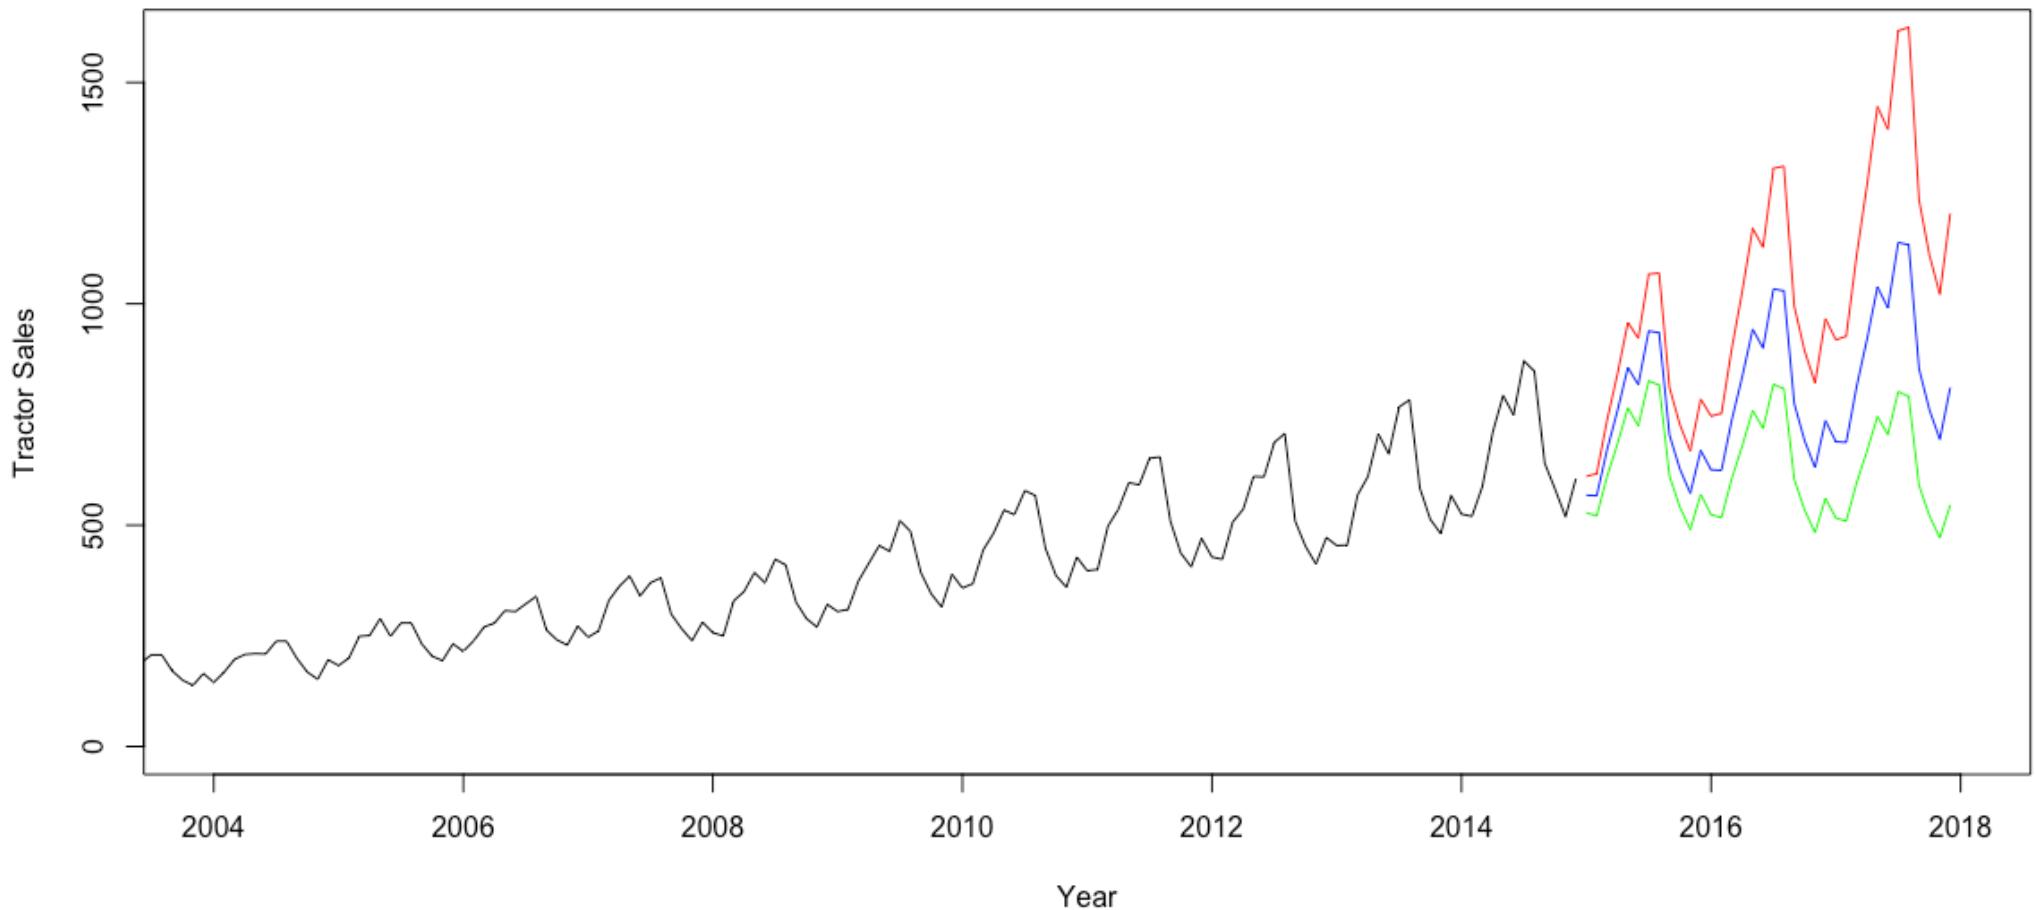


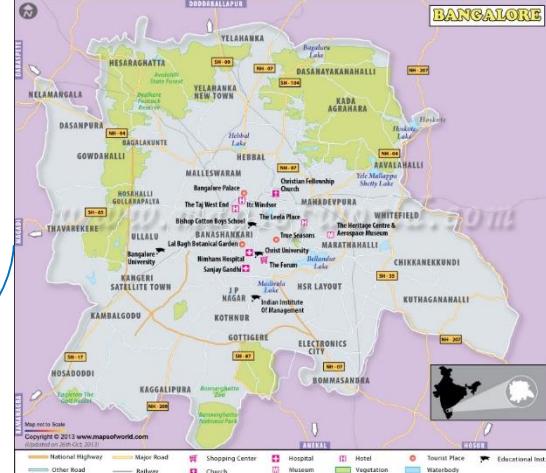
CEE 72026



Tractor Sales Case Study

Step 8: Forecast on the original scale and not the logged data.





HYDERABAD

Plot 63/A, Floors 1&2, Road # 13, Film Nagar,
Jubilee Hills, Hyderabad - 500 033
+91-9701685511 (Individuals)
+91-9618483483 (Corporates)

BENGALURU

Incubex, #728, Grace Platina, 4th Floor, CMH Road,
Indira Nagar, 1st Stage, Bengaluru – 560038
+91-9502334561 (Individuals)
+91-9502799088 (Corporates)

Social Media

- Web: <http://www.insofe.edu.in>
- Facebook: <https://www.facebook.com/insofe>
- Twitter: <https://twitter.com/Insofeedu>
- YouTube: <http://www.youtube.com/InsofeVideos>
- SlideShare: <http://www.slideshare.net/INSOFE>
- LinkedIn: <http://www.linkedin.com/company/international-school-of-engineering>

This presentation may contain references to findings of various reports available in the public domain. INSOFE makes no representation as to their accuracy or that the organization subscribes to those findings.